



Lecture: RANSAC and feature detectors

Juan Carlos Niebles and Ranjay Krishna
Stanford Vision and Learning Lab



CS 131 Roadmap

Pixels

Segments

Images

Videos

Web

Convolutions
Edges
Features

Resizing
Segmentation
Clustering

Recognition
Detection
Machine learning

Motion
Tracking

Neural networks
Convolutional
neural networks



What we will learn today?

- A model fitting method for line detection
 - RANSAC
- Local invariant features
 - Motivation
 - Requirements, invariances
- Keypoint localization
 - Harris corner detector



What we will learn today

- A model fitting method for line detection
 - RANSAC
- Local invariant features
 - Motivation
 - Requirements, invariances
- Keypoint localization
 - Harris corner detector



Fitting as Search in Parametric Space

- Choose a parametric model to represent a set of features
- Membership criterion is not local
 - Can't tell whether a point belongs to a given model just by looking at that point.
- Three main questions:
 - What model represents this set of features best?
 - Which of several model instances gets which feature?
 - How many model instances are there?
- Computational complexity is important
 - It is infeasible to examine every possible set of parameters and every possible combination of features



Example: Line Fitting

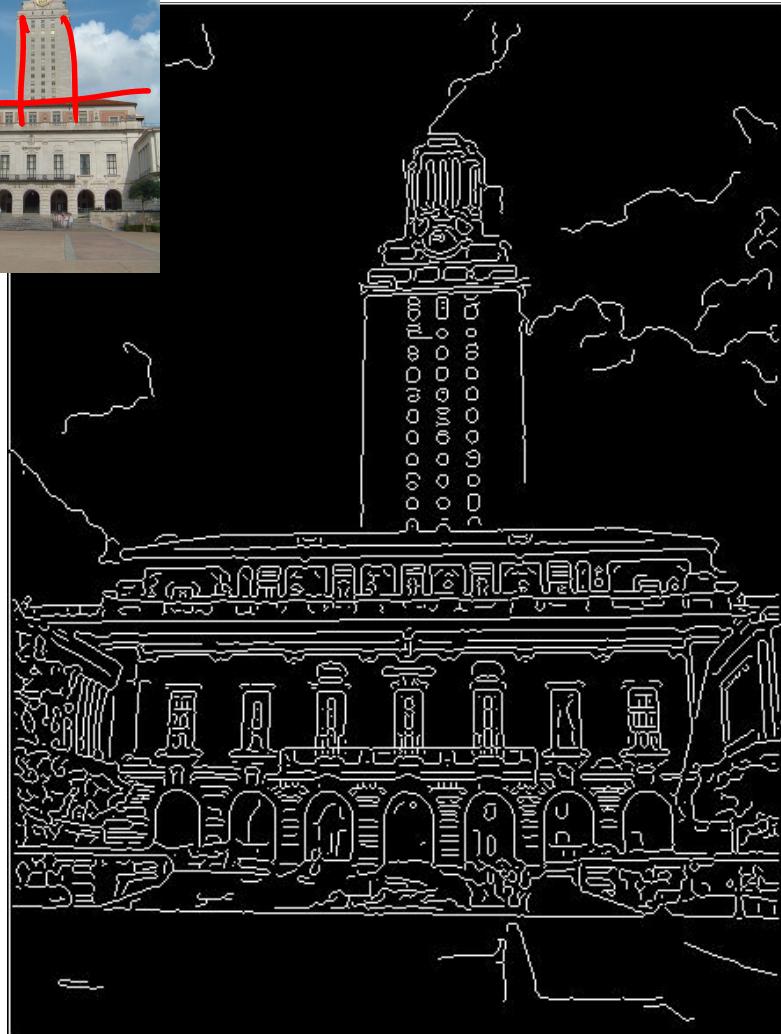
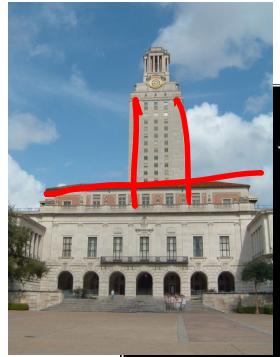
- Why fit lines?
Many objects characterized by presence of straight lines



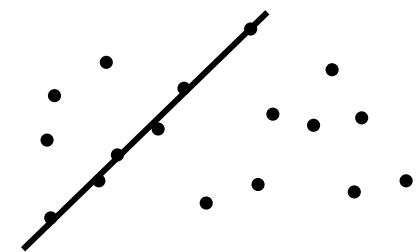
- Wait, why aren't we done just by running edge detection?



Difficulty of Line Fitting



- Extra edge points (clutter), multiple models:
 - Which points go with which line, if any?
- Only some parts of each line detected, and some parts are missing:
 - How to find a line that bridges missing evidence?
- Noise in measured edge points, orientations:
 - How to detect true underlying parameters?





Voting

- It's not feasible to check all combinations of features by fitting a model to each possible subset.
- Voting is a general technique where we let the features vote for all models that are compatible with it.
 - Cycle through features, cast votes for model parameters.
 - Look for model parameters that receive a lot of votes.
- Noise & clutter features will cast votes too, *but* typically their votes should be inconsistent with the majority of “good” features.
- Ok if some features not observed, as model can span multiple fragments.



RANSAC [Fischler & Bolles 1981]

- RANdom SAmple Consensus
- Approach: we want to avoid the impact of outliers, so let's look for “inliers”, and use only those.
- Intuition: if an outlier is chosen to compute the current fit, then the resulting line won't have much support from rest of the points.



RANSAC [Fischler & Bolles 1981]

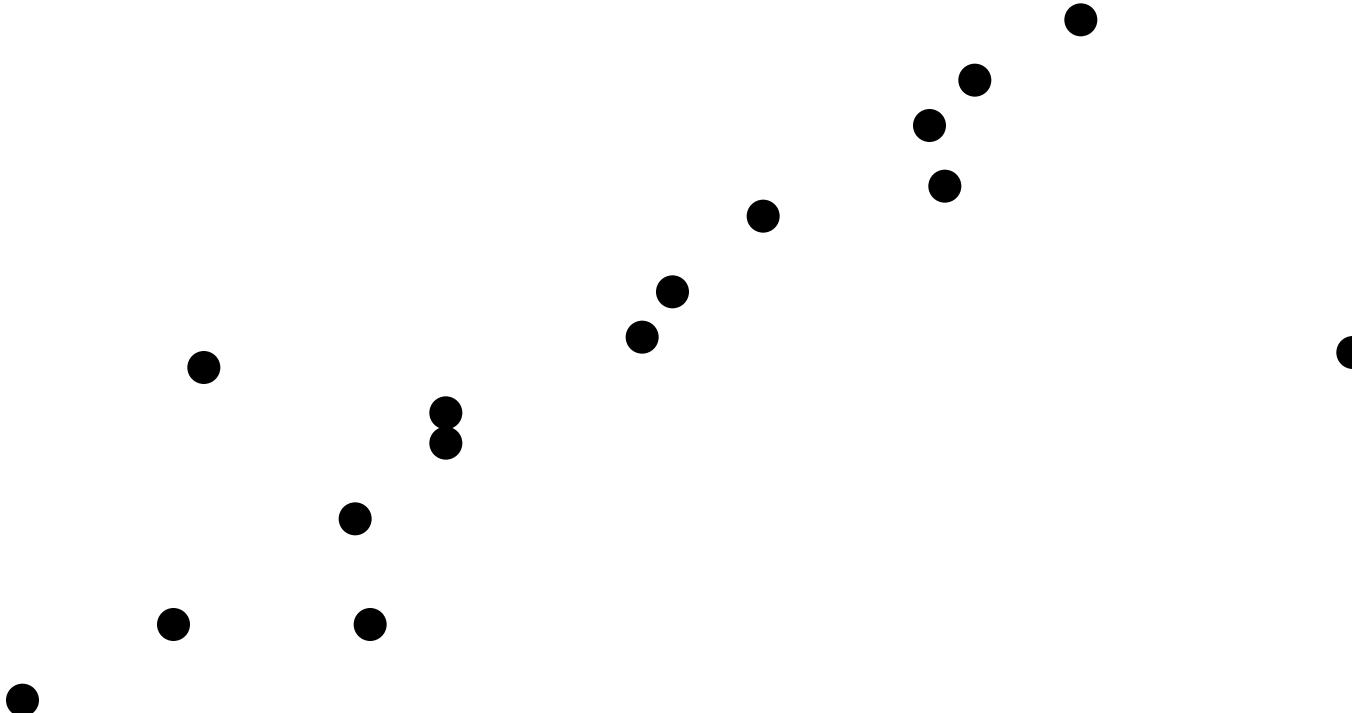
RANSAC loop:

1. Randomly select a *seed group* of points on which to base transformation estimate (e.g., a group of matches)
 2. Compute transformation from seed group
 3. Find *inliers* to this transformation
 4. If the number of inliers is sufficiently large, re-compute least-squares estimate of transformation on all of the inliers
- Keep the transformation with the largest number of inliers



RANSAC Line Fitting Example

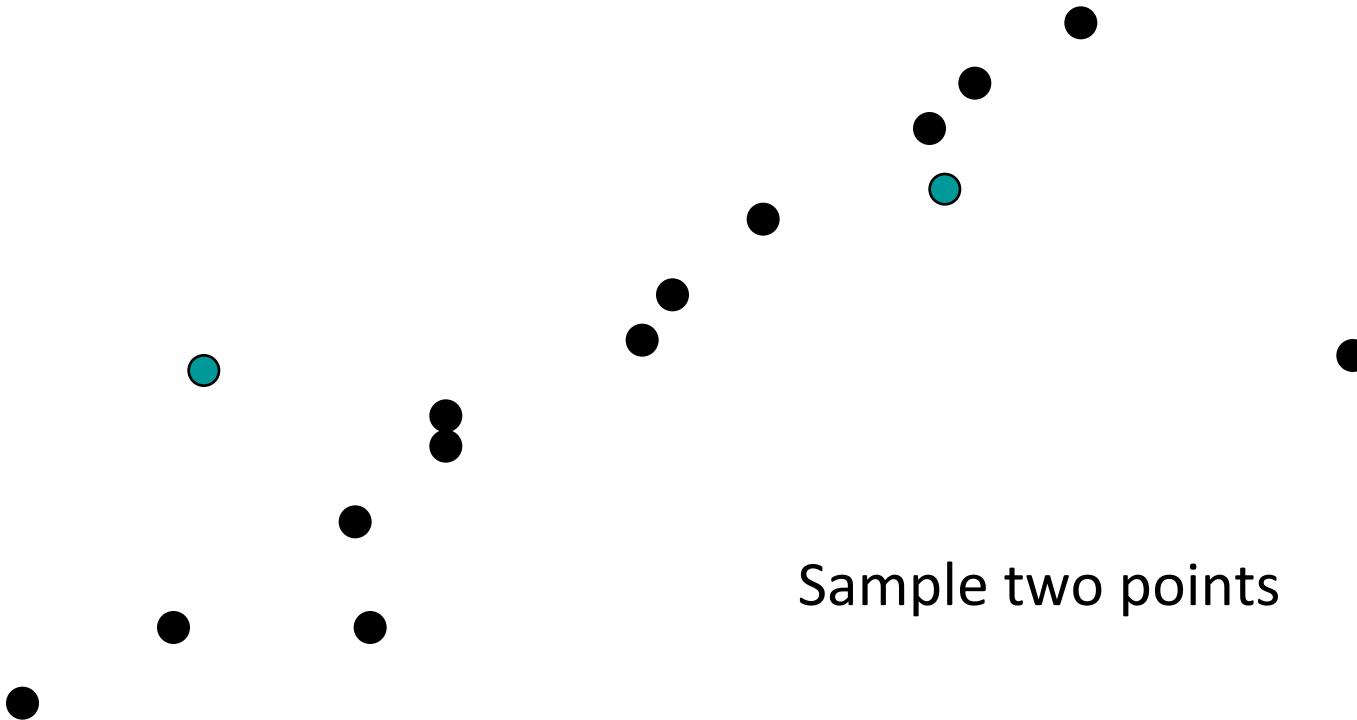
- Task: Estimate the best line
 - *How many points do we need to estimate the line?*





RANSAC Line Fitting Example

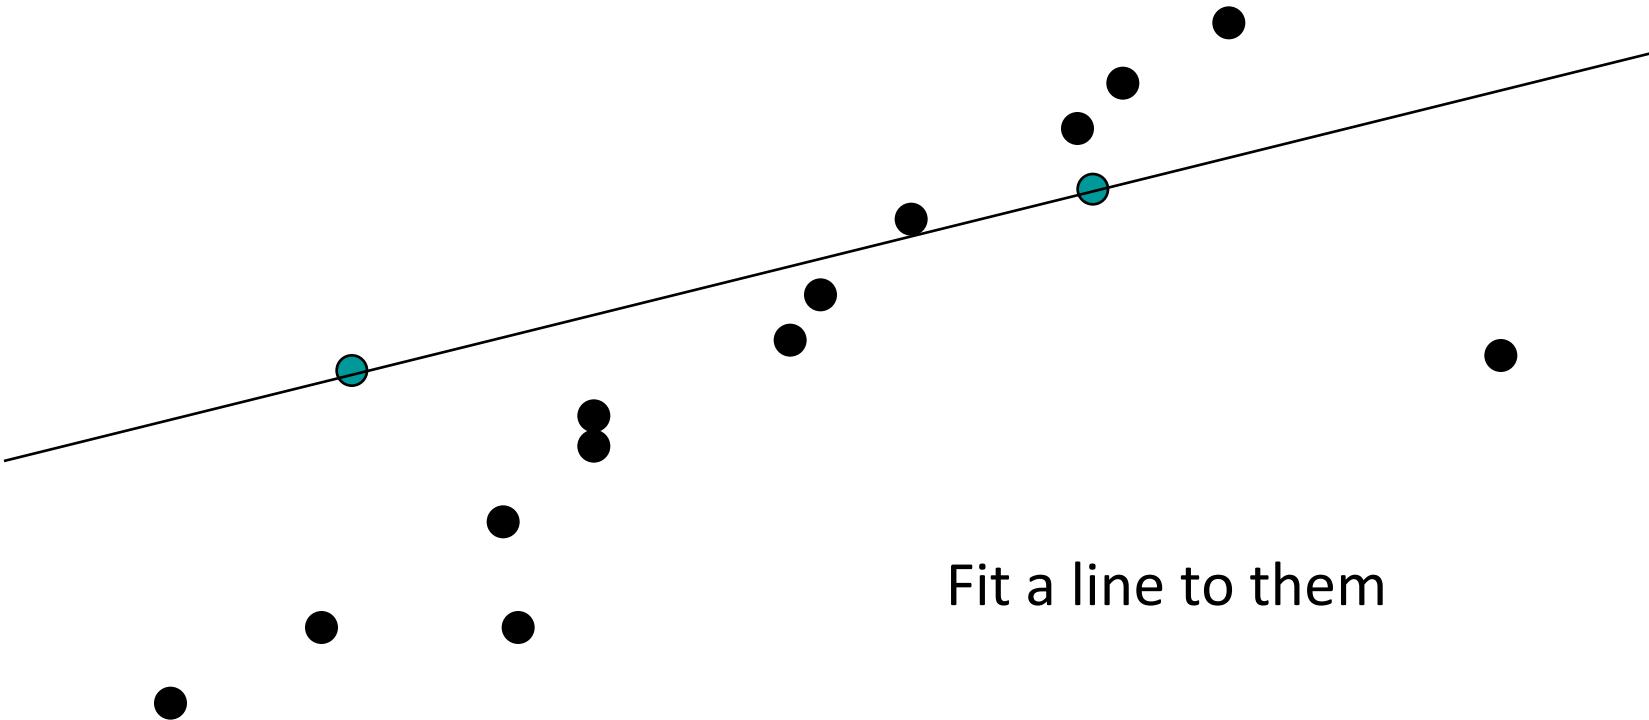
- Task: Estimate the best line





RANSAC Line Fitting Example

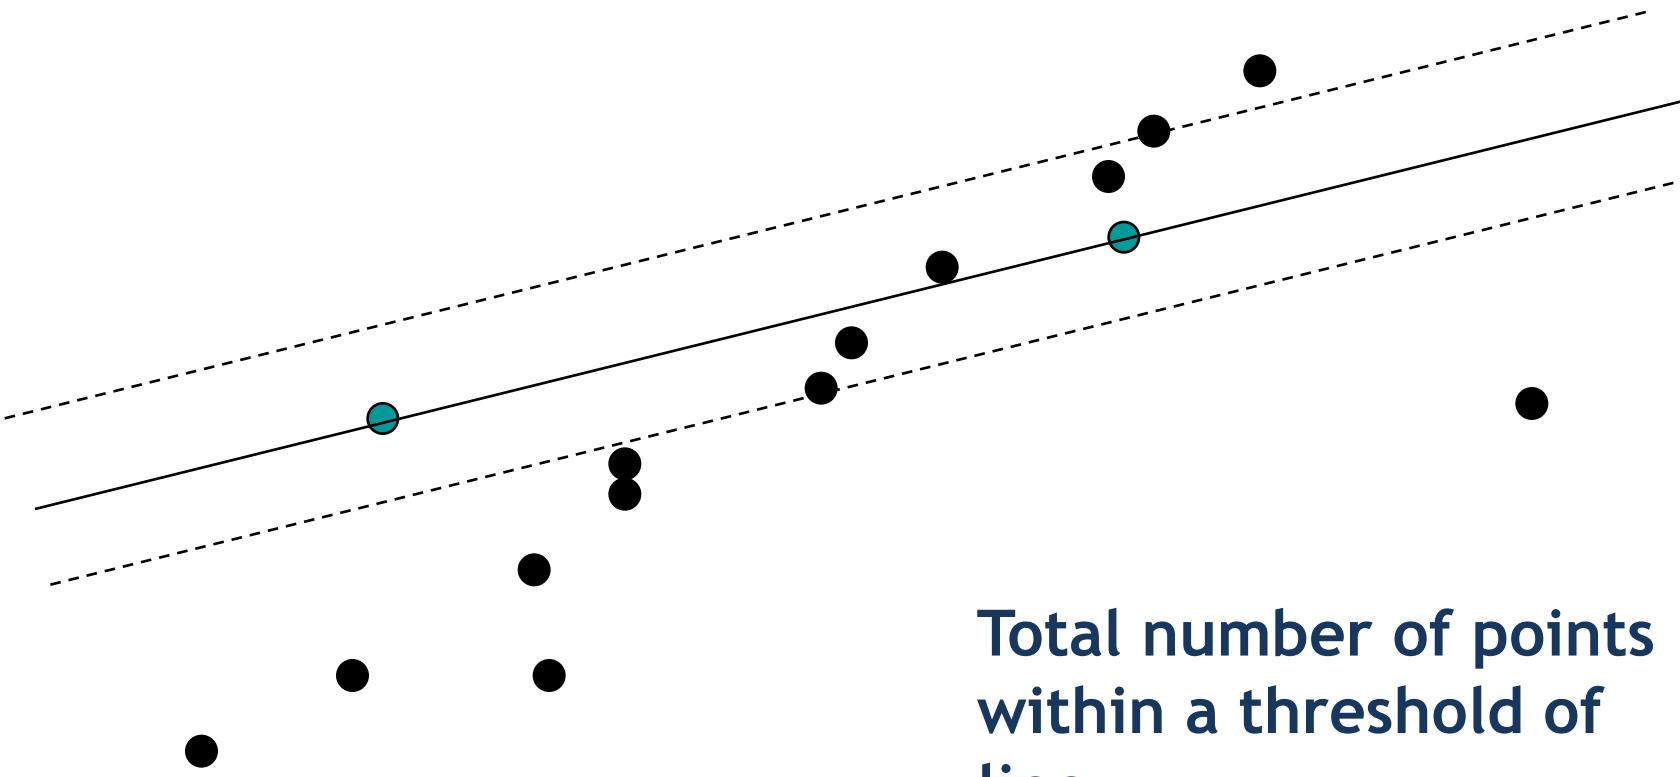
- Task: Estimate the best line





RANSAC Line Fitting Example

- Task: Estimate the best line

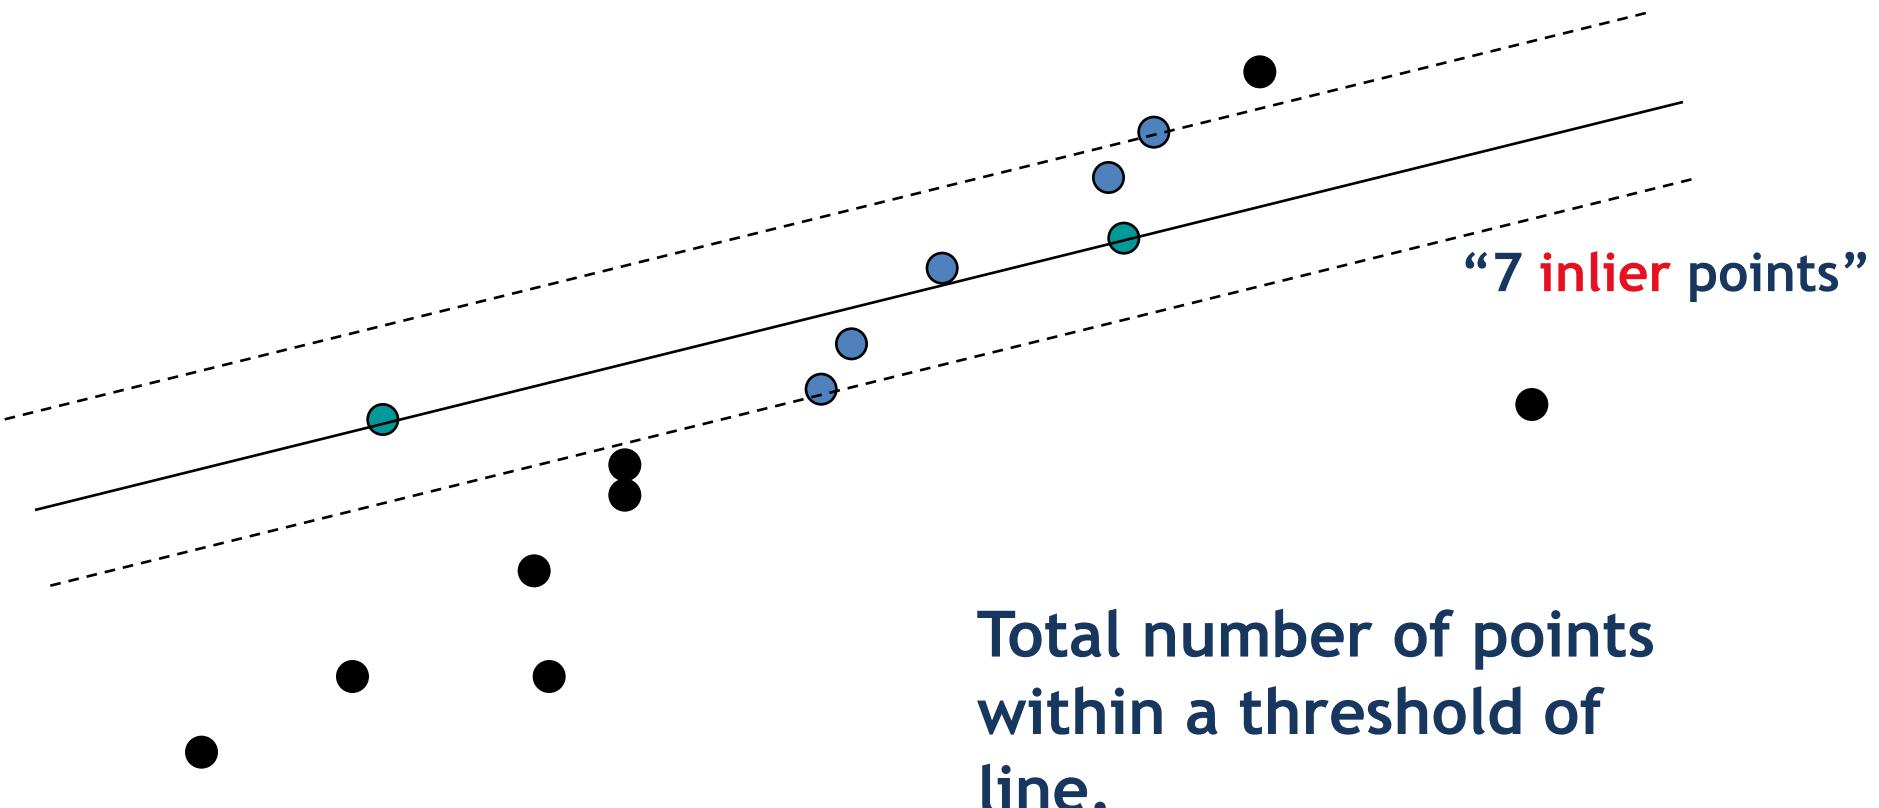


Total number of points
within a threshold of
line.



RANSAC Line Fitting Example

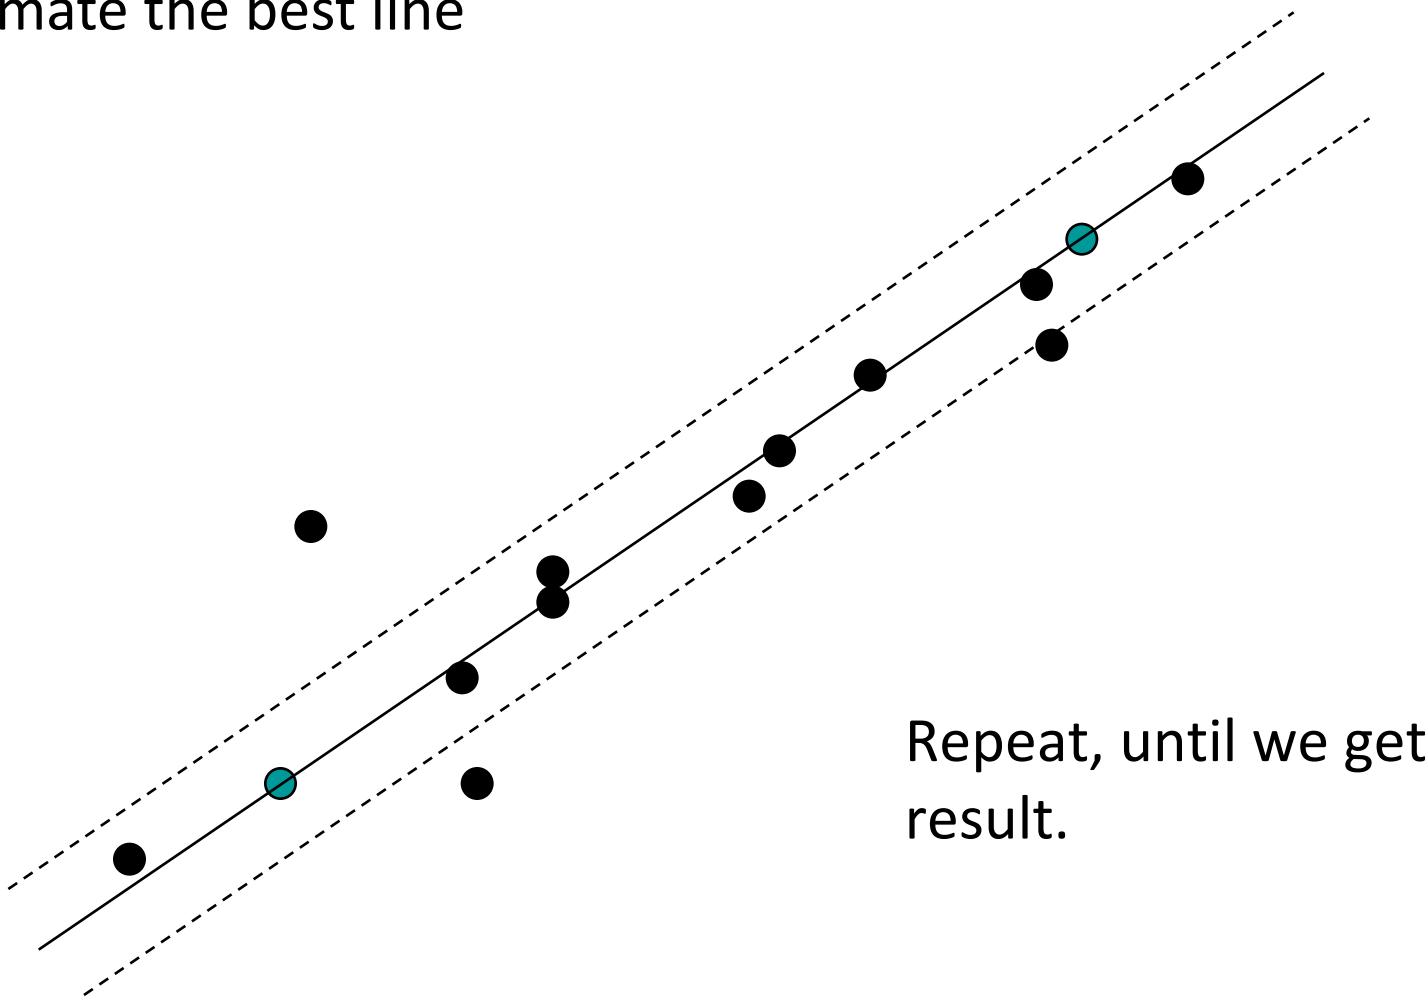
- Task: Estimate the best line





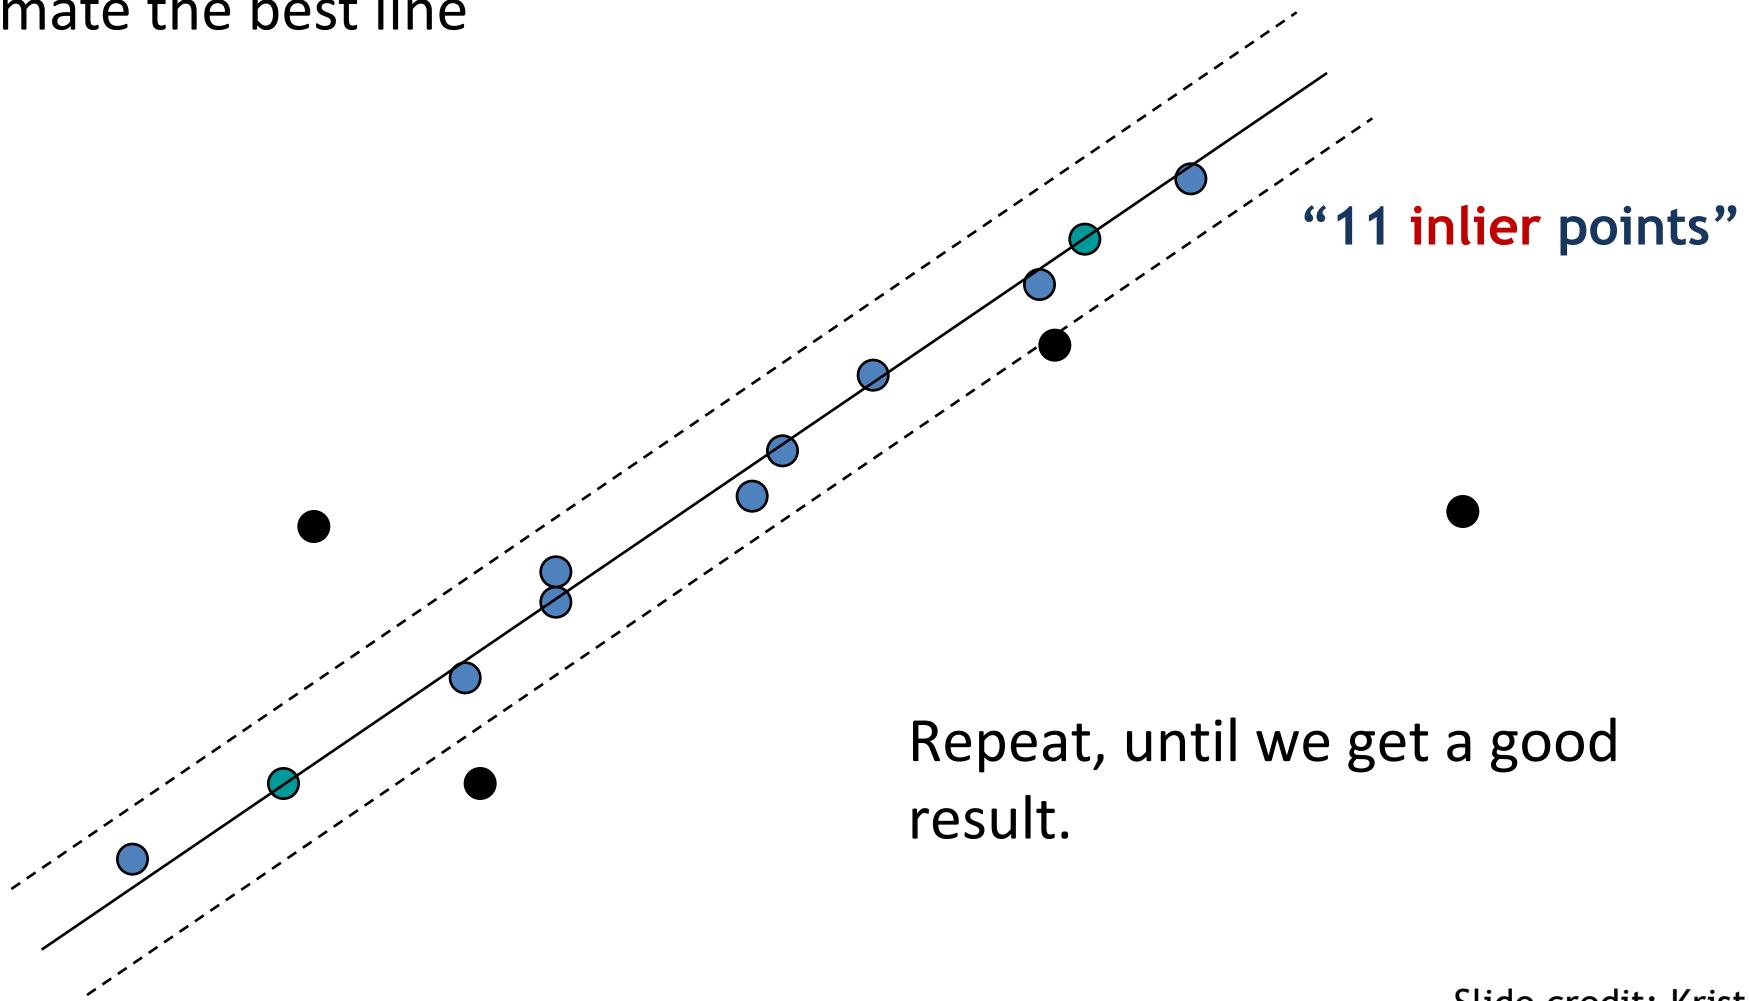
RANSAC Line Fitting Example

- Task: Estimate the best line



RANSAC Line Fitting Example

- Task: Estimate the best line





Algorithm 15.4: RANSAC: fitting lines using random sample consensus

Determine:

n — the smallest number of points required

k — the number of iterations required

t — the threshold used to identify a point that fits well

d — the number of nearby points required
to assert a model fits well

Until k iterations have occurred

 Draw a sample of n points from the data
 uniformly and at random

 Fit to that set of n points

 For each data point outside the sample

 Test the distance from the point to the line
 against t ; if the distance from the point to the line
 is less than t , the point is close

 end

 If there are d or more points close to the line
 then there is a good fit. Refit the line using all
 these points.

end

Use the best fit from this collection, using the
fitting error as a criterion



RANSAC: How many samples?

- How many samples are needed?
 - Suppose w is fraction of inliers (points from line).
 - n points needed to define hypothesis (2 for lines)
 - k samples chosen.
- Prob. that a single sample of n points is correct: w^n
- Prob. that all k samples fail is: $(1 - w^n)^k$

⇒ Choose k high enough to keep this below desired failure rate.



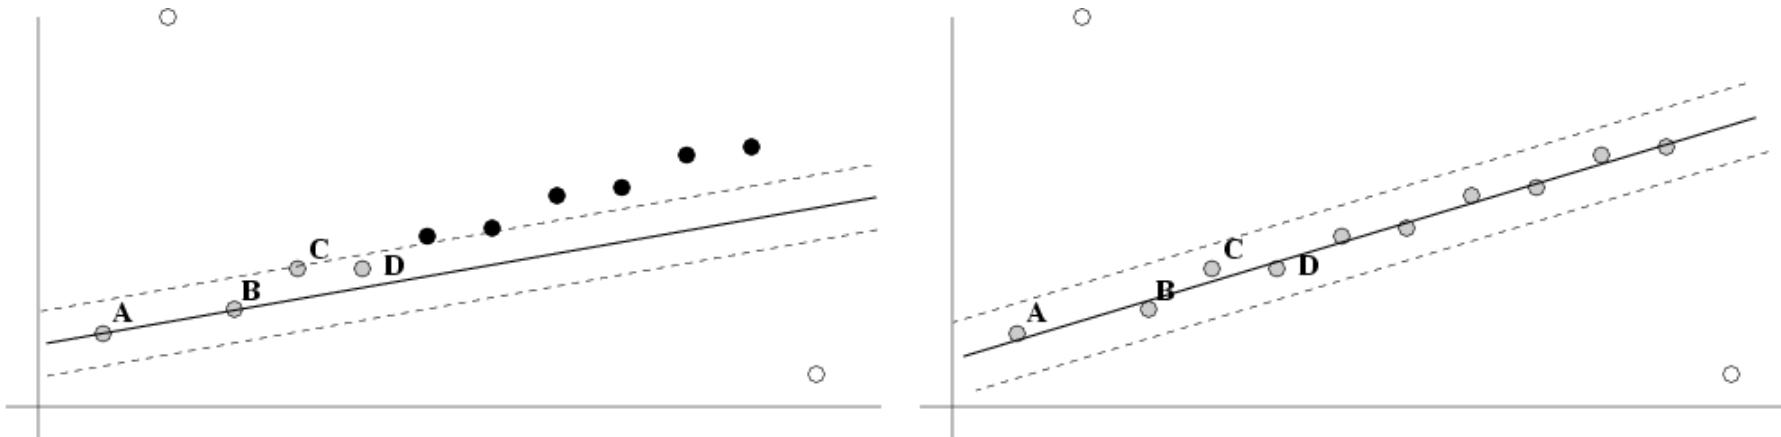
RANSAC: Computed k ($p=0.99$)

Sample size n	Proportion of outliers						
	5%	10%	20%	25%	30%	40%	50%
2	2	3	5	6	7	11	17
3	3	4	7	9	11	19	35
4	3	5	9	13	17	34	72
5	4	6	12	17	26	57	146
6	4	7	16	24	37	97	293
7	4	8	20	33	54	163	588
8	5	9	26	44	78	272	1177



After RANSAC

- RANSAC divides data into inliers and outliers and yields estimate computed from minimal set of inliers.
- Improve this initial estimate with estimation over all inliers (e.g. with standard least-squares minimization).
- But this may change inliers, so alternate fitting with re-classification as inlier/outlier.





RANSAC: Pros and Cons

- **Pros:**
 - General method suited for a wide range of model fitting problems
 - Easy to implement and easy to calculate its failure rate
- **Cons:**
 - Only handles a moderate percentage of outliers without cost blowing up
 - Many real problems have high rate of outliers (but sometimes selective choice of random subsets can help)
- A voting strategy, The Hough transform, can handle high percentage of outliers



What we will learn today?

- A model fitting method for edge detection
 - RANSAC
- Local invariant features
 - Motivation
 - Requirements, invariances
- Keypoint localization
 - Harris corner detector

Some background reading:

Rick Szeliski, Chapter 4.1.1; David Lowe, IJCV 2004



Image matching: a challenging problem





Image matching: a challenging problem



by [Diva Sian](#)



by [swashford](#)



Harder Case



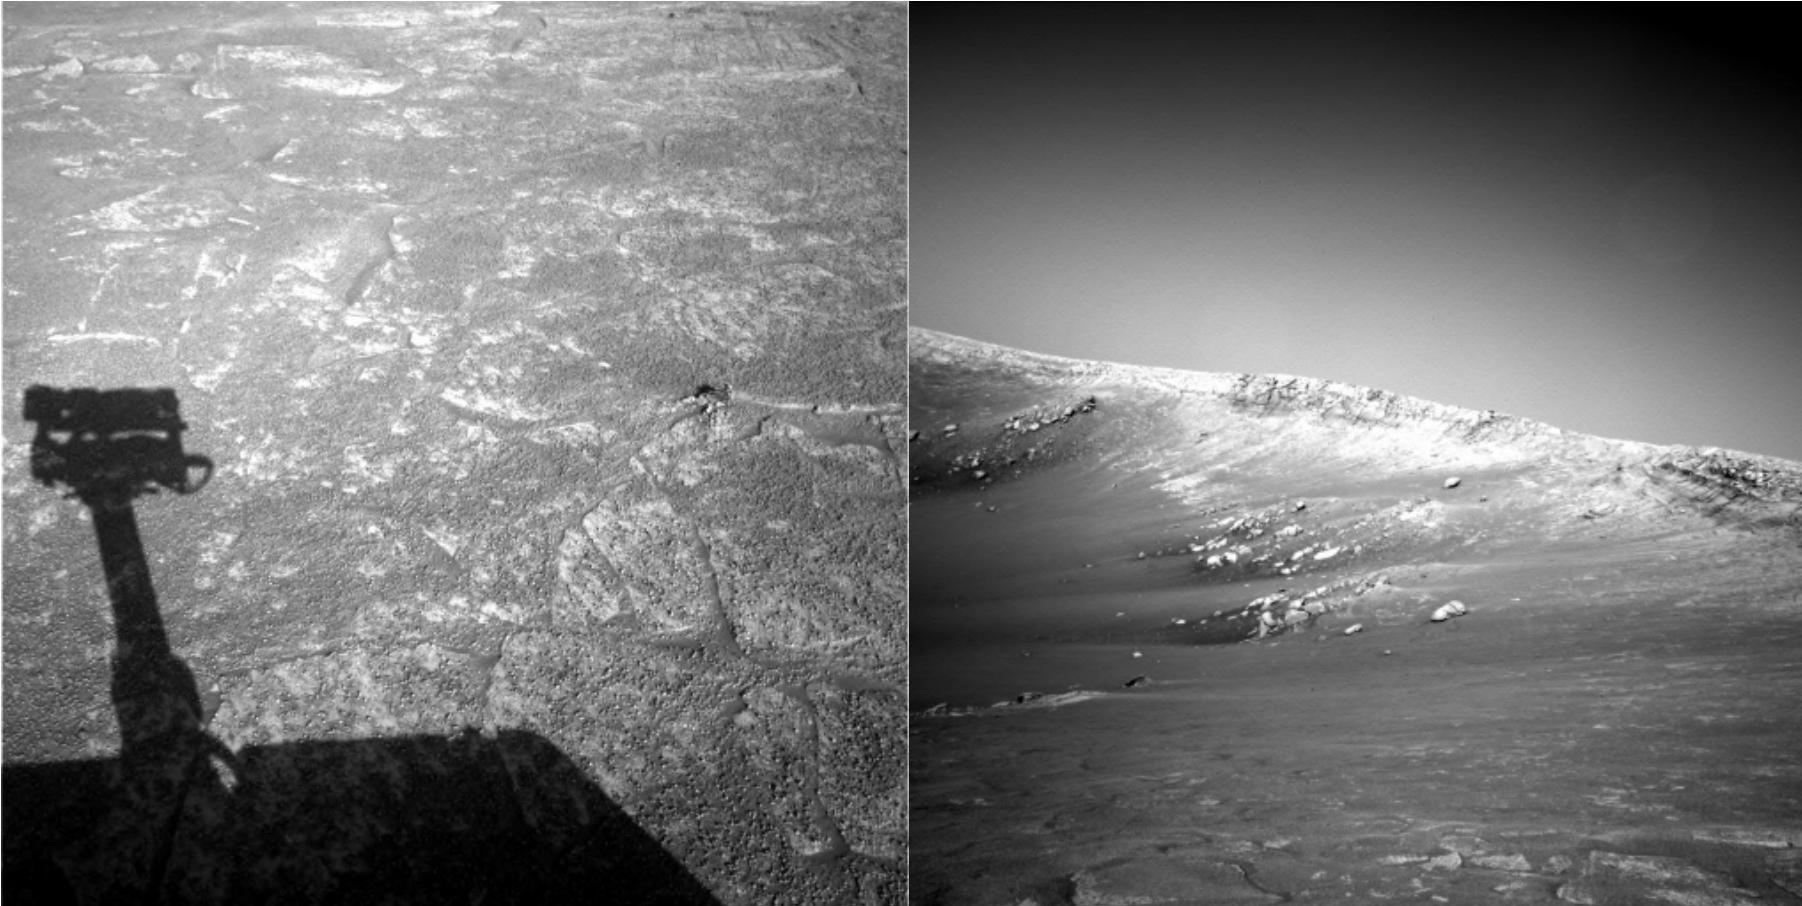
by [Diva Sian](#)



by [scgbt](#)



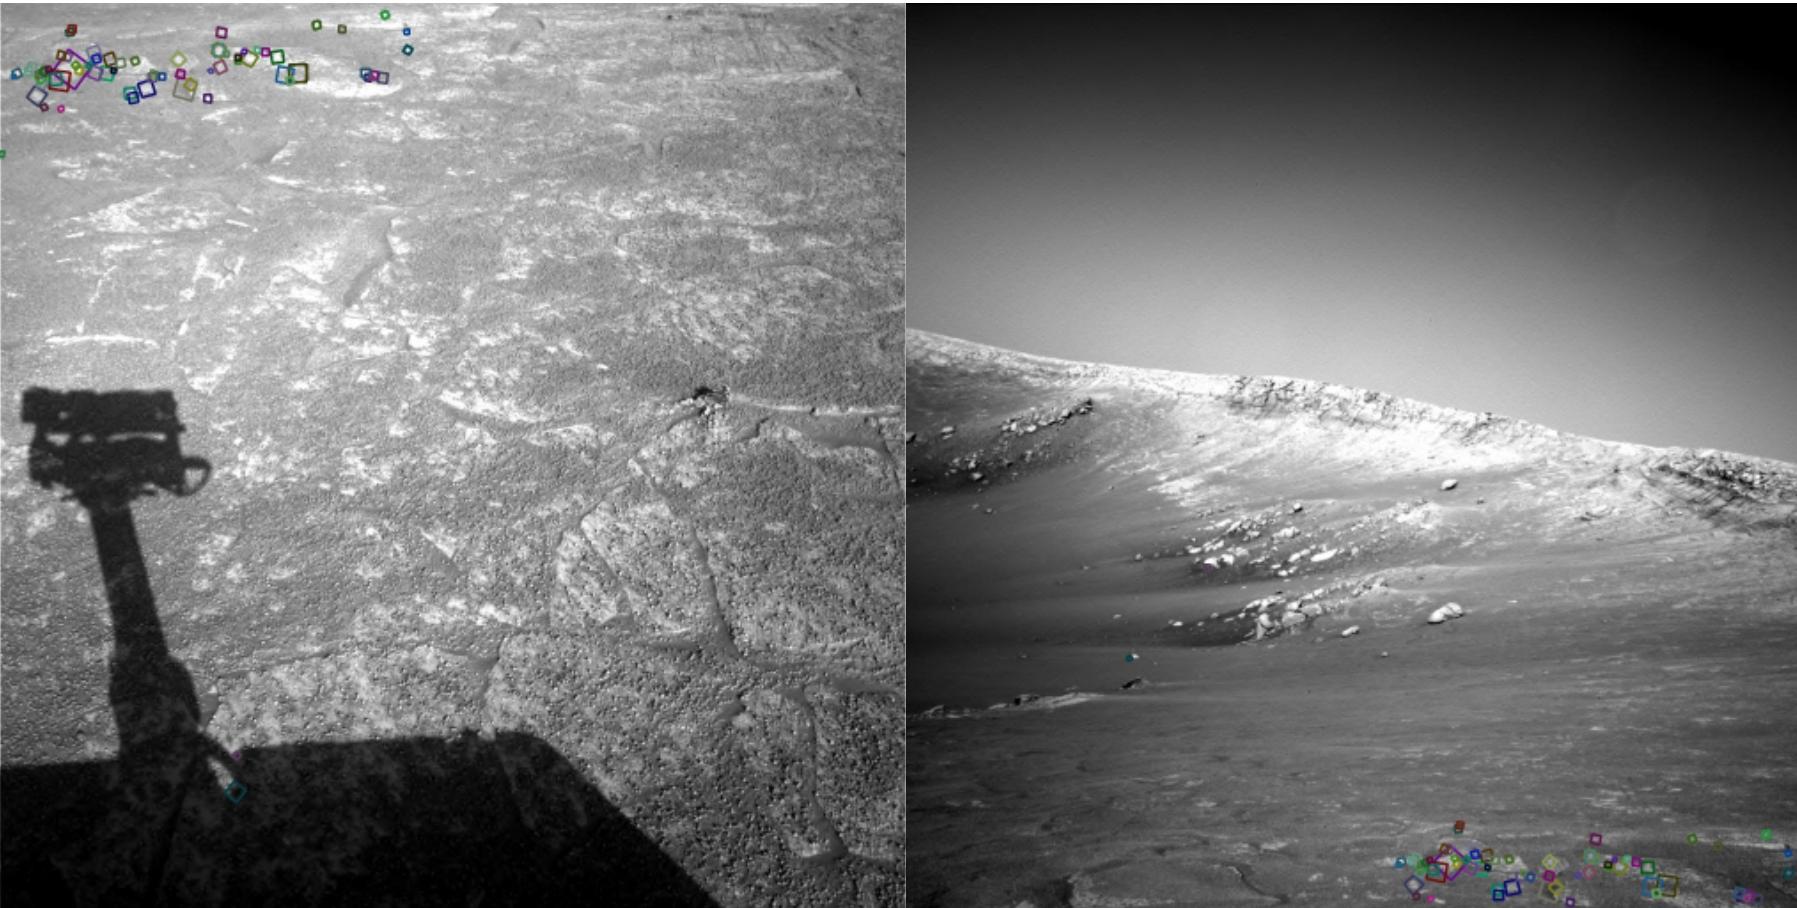
Harder Still?



NASA Mars Rover images



Answer Below (Look for tiny colored squares)



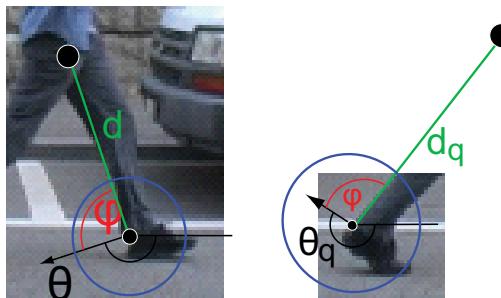
NASA Mars Rover images with SIFT feature matches
(Figure by Noah Snavely)

Motivation for using local features

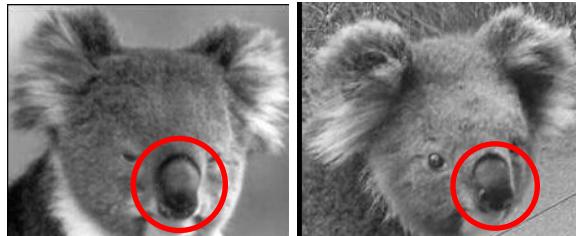
- Global representations have major limitations
- Instead, describe and match only local regions
- Increased robustness to
 - Occlusions



- Articulation

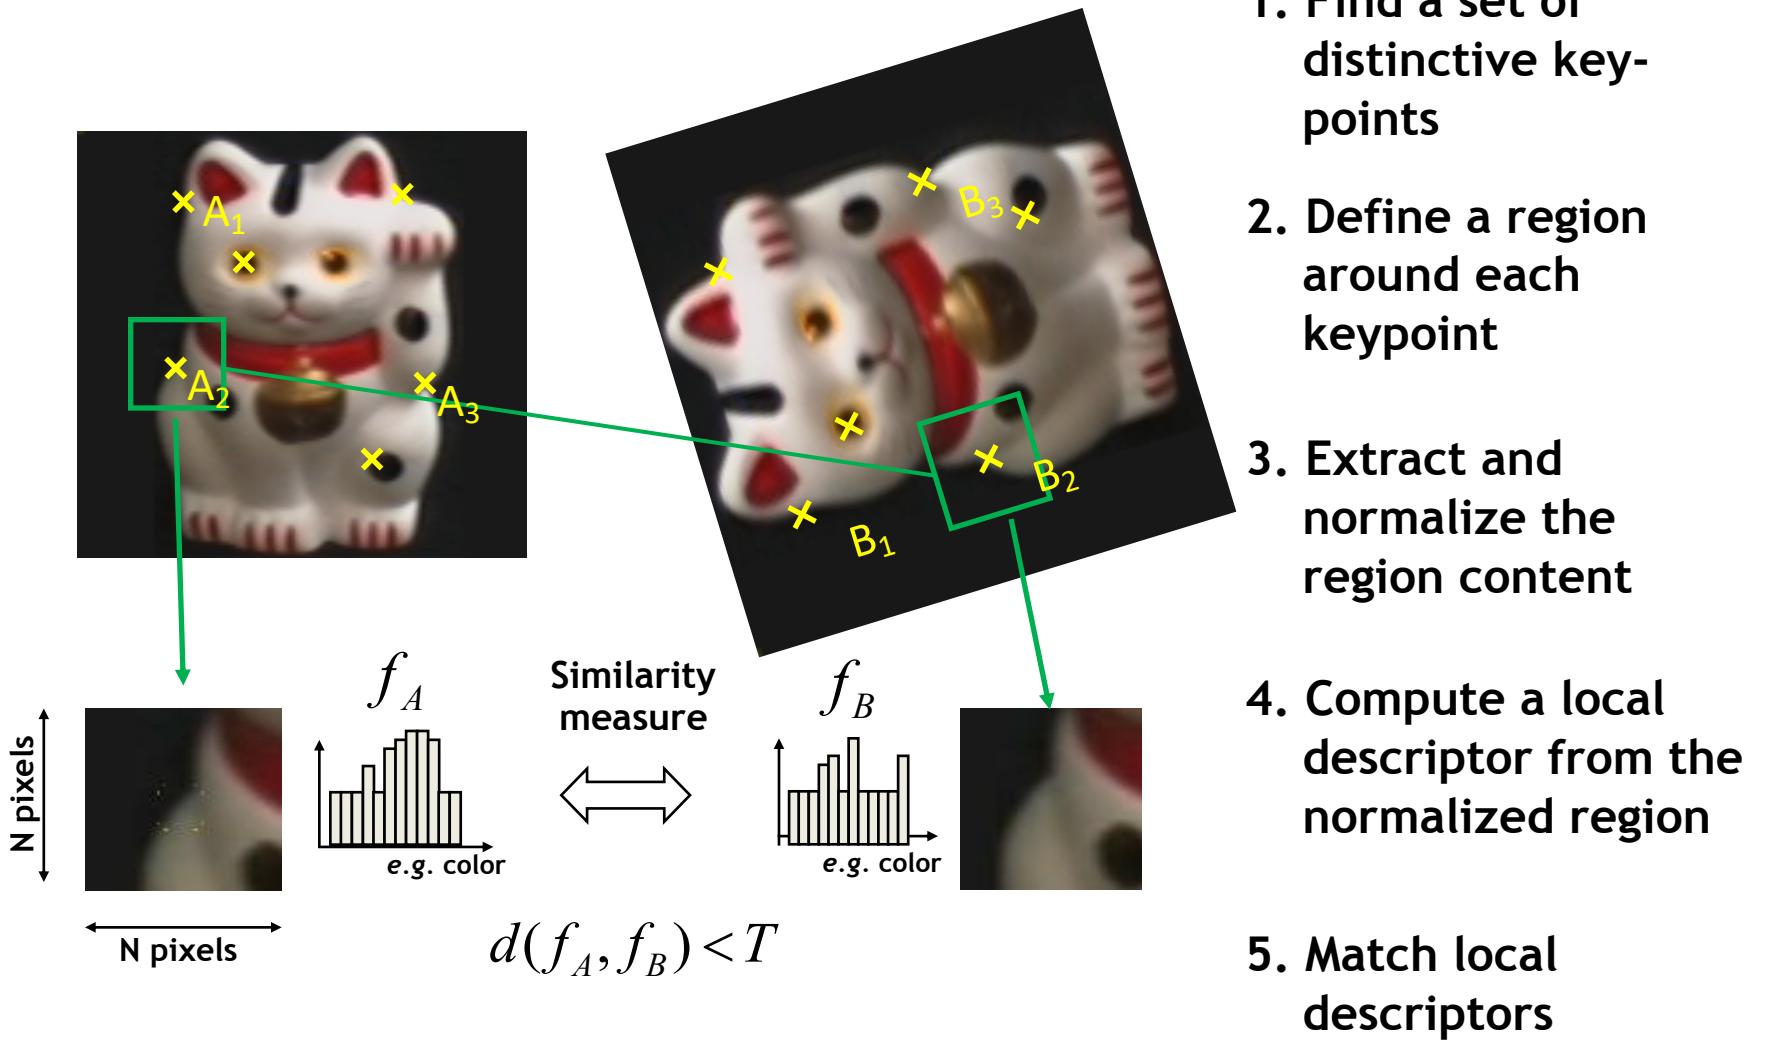


- Intra-category variations



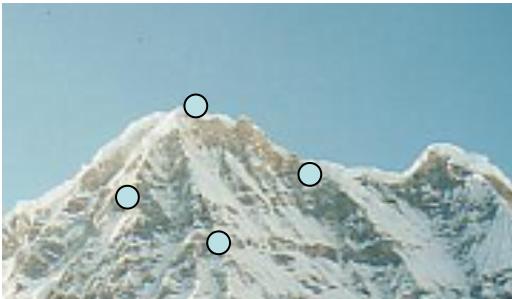


General Approach



Common Requirements

- Problem 1:
 - Detect the same point *independently* in both images

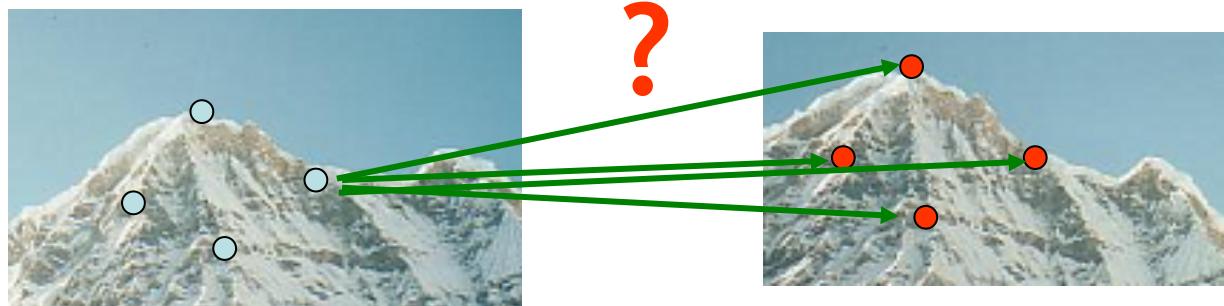


No chance to match!

We need a repeatable detector!

Common Requirements

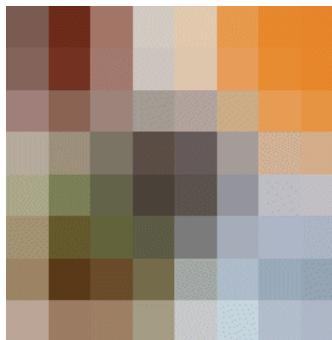
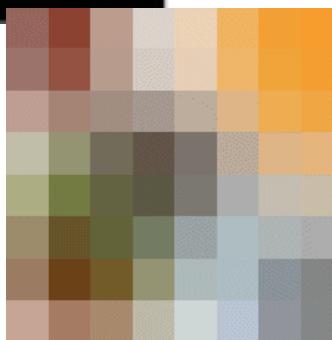
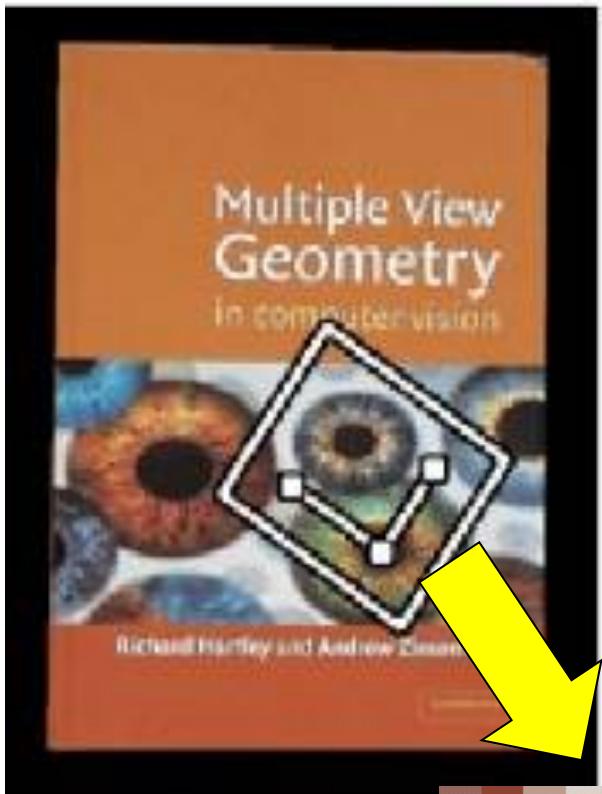
- Problem 1:
 - Detect the same point *independently* in both images
- Problem 2:
 - For each point correctly recognize the corresponding one



We need a reliable and distinctive descriptor!

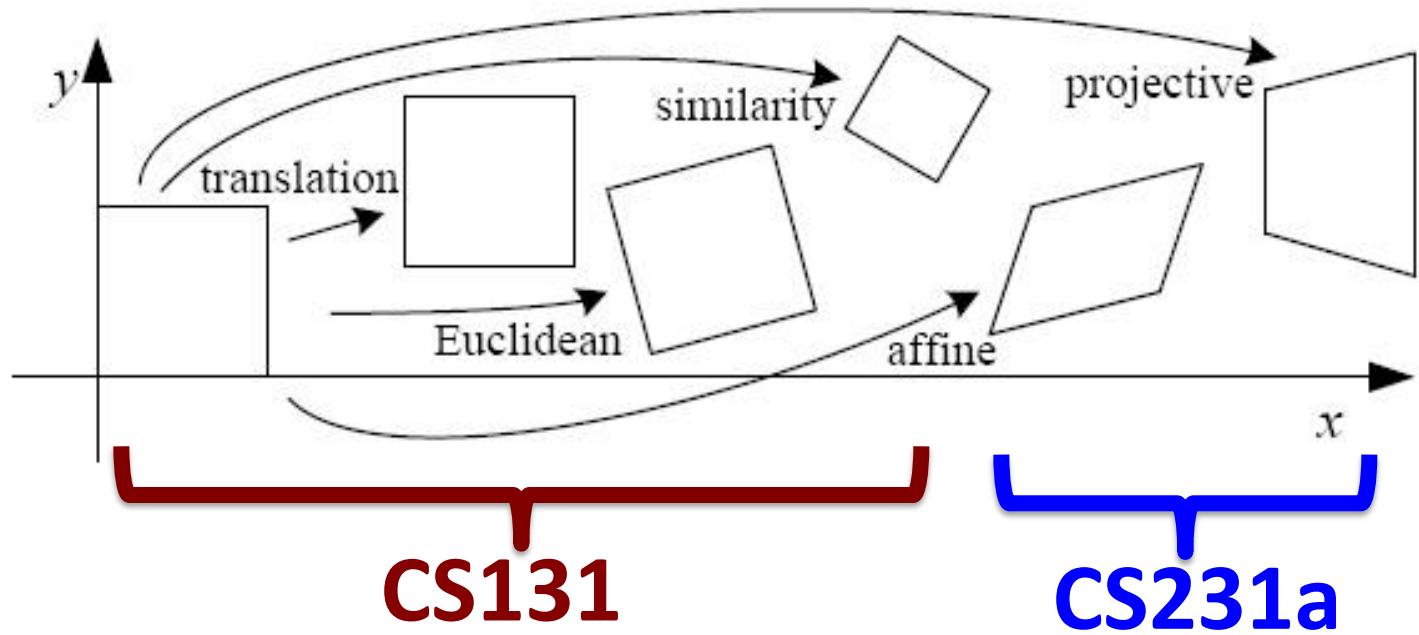


Invariance: Geometric Transformations



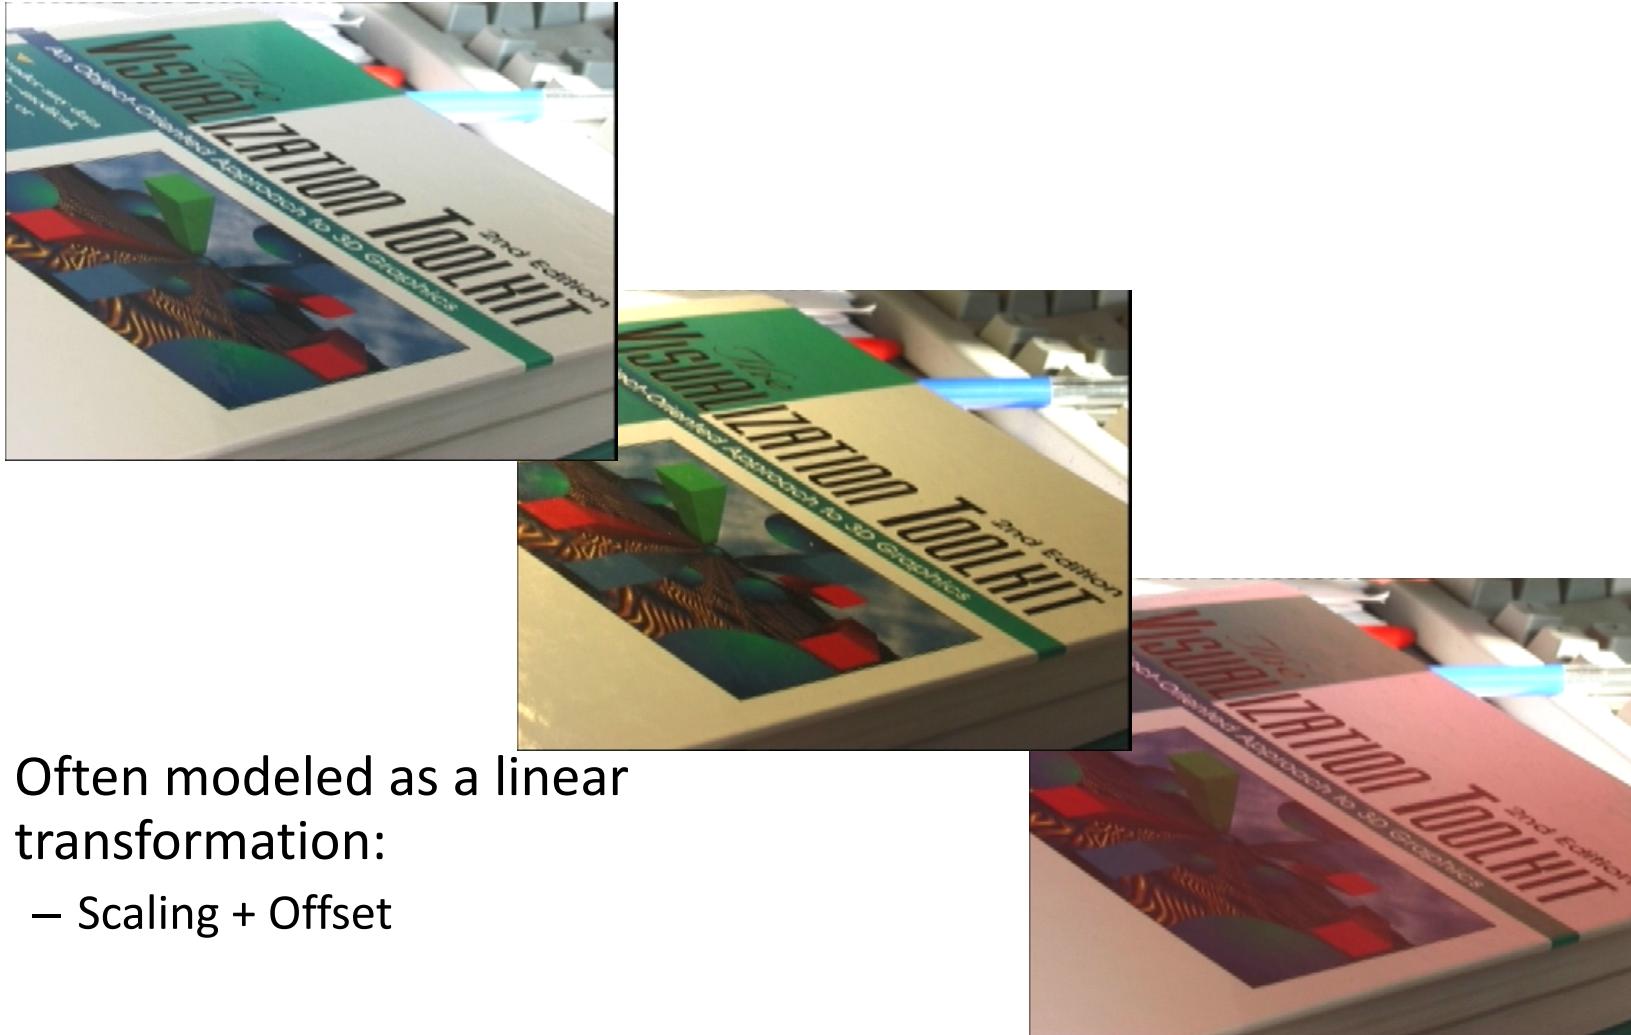


Levels of Geometric Invariance





Invariance: Photometric Transformations



- Often modeled as a linear transformation:
 - Scaling + Offset



Requirements

- Region extraction needs to be **repeatable** and **accurate**
 - **Invariant** to translation, rotation, scale changes
 - **Robust** or **covariant** to out-of-plane (\approx affine) transformations
 - **Robust** to lighting variations, noise, blur, quantization
- **Locality**: Features are local, therefore robust to occlusion and clutter.
- **Quantity**: We need a sufficient number of regions to cover the object.
- **Distinctiveness** : The regions should contain “interesting” structure.
- **Efficiency**: Close to real-time performance.



Many Existing Detectors Available

- Hessian & Harris [Beaudet '78], [Harris '88]
- Laplacian, DoG [Lindeberg '98], [Lowe '99]
- Harris-/Hessian-Laplace [Mikolajczyk & Schmid '01]
- Harris-/Hessian-Affine [Mikolajczyk & Schmid '04]
- EBR and IBR [Tuytelaars & Van Gool '04]
- MSER [Matas '02]
- Salient Regions [Kadir & Brady '01]
- Others...
- *Those detectors have become a basic building block for many applications in Computer Vision.*



What we will learn today?

- A model fitting method for edge detection
 - RANSAC
- Local invariant features
 - Motivation
 - Requirements, invariances
- Keypoint localization
 - Harris corner detector

Some background reading:

Rick Szeliski, Chapter 4.1.1; David Lowe, IJCV 2004

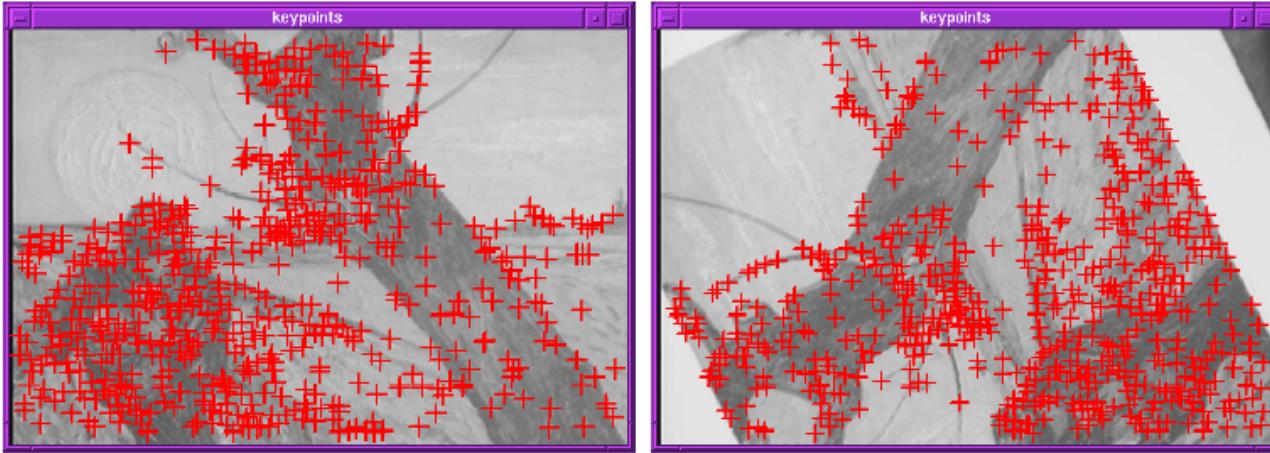
Keypoint Localization

- Goals:
 - Repeatable detection
 - Precise localization
 - Interesting content
- ⇒ *Look for two-dimensional signal changes*





Finding Corners



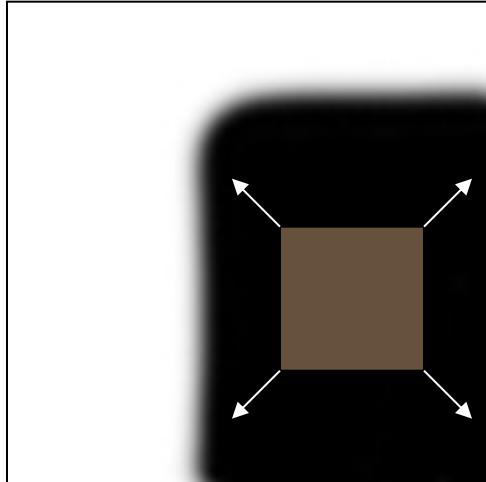
- Key property:
 - In the region around a corner, image gradient has two or more dominant directions
- Corners are *repeatable* and *distinctive*

C.Harris and M.Stephens. "[A Combined Corner and Edge Detector.](#)"
Proceedings of the 4th Alvey Vision Conference, 1988.

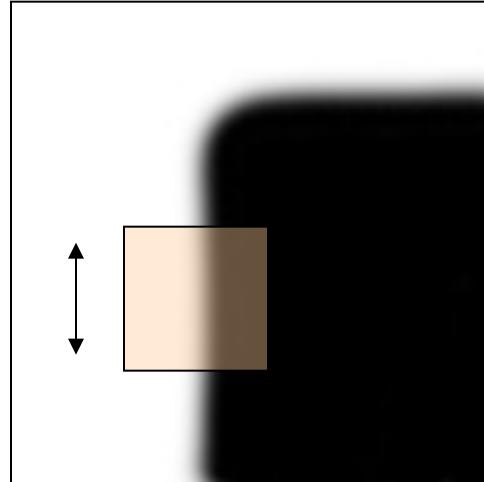


Corners as Distinctive Interest Points

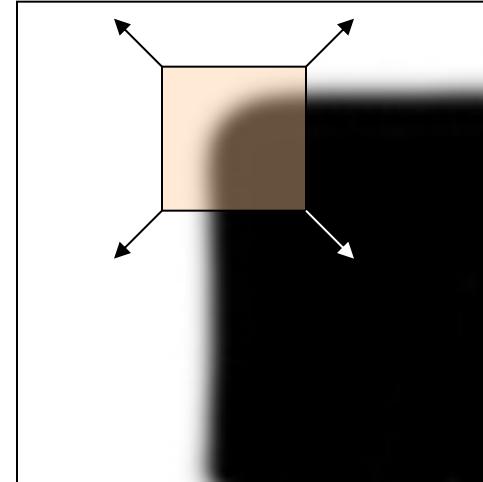
- Design criteria
 - We should easily recognize the point by looking through a small window (*locality*)
 - Shifting the window in *any direction* should give *a large change* in intensity (*good localization*)



“flat” region:
no change in all
directions



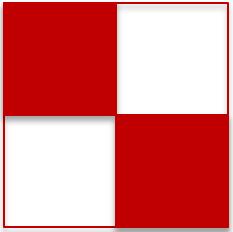
“edge”:
no change along
the edge direction



“corner”:
significant change
in all directions



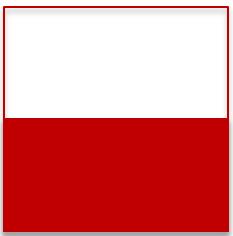
Corners versus edges



$$\sum I_x^2 \longrightarrow \text{Large}$$

$$\sum I_y^2 \longrightarrow \text{Large}$$

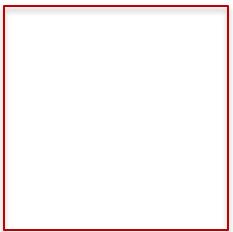
Corner



$$\sum I_x^2 \longrightarrow \text{Small}$$

$$\sum I_y^2 \longrightarrow \text{Large}$$

Edge



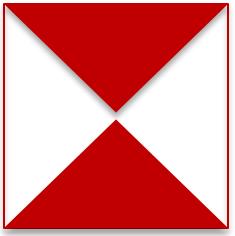
$$\sum I_x^2 \longrightarrow \text{Small}$$

$$\sum I_y^2 \longrightarrow \text{Small}$$

Nothing



Corners versus edges

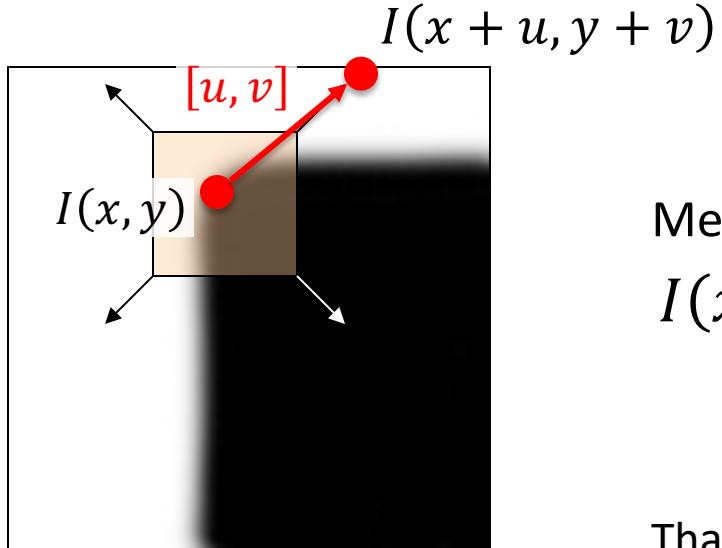


$$\sum I_x^2 \longrightarrow ??$$
$$\sum I_y^2 \longrightarrow ??$$

Corner



Measure change in all directions $[u, v]$



“corner”:
significant change
in all directions

Measure change as intensity difference:

$$I(x + u, y + v) - I(x, y)$$

That's for a single point, but we have to accumulate over a “small window” around that point...



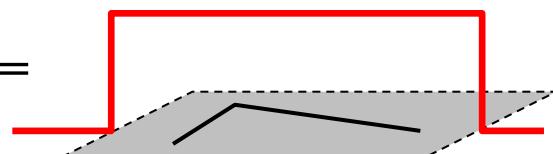
Harris Detector Formulation

- Change of intensity for the shift $[u,v]$:

$$E(u, v) = \sum_{x,y} w(x, y) [I(x + u, y + v) - I(x, y)]^2$$

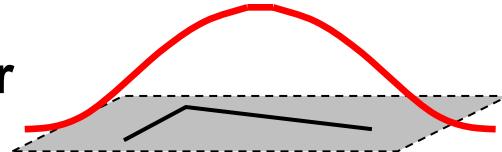
Window function **Shifted intensity** **Intensity**

Window function $w(x,y) =$



1 in window, 0 outside

or



Gaussian



Harris Detector Formulation

- This measure of change can be approximated by (Taylor expansion):

$$E(u, v) \approx [u \ v] M \begin{bmatrix} u \\ v \end{bmatrix}$$

where M is a 2×2 matrix computed from image derivatives:

$$M = \sum_{x,y} w(x, y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

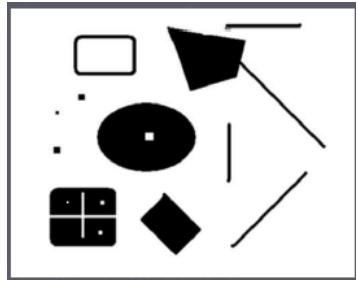
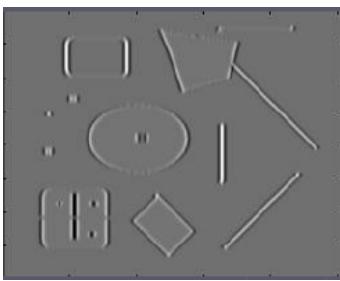
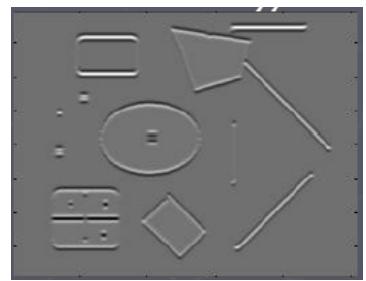
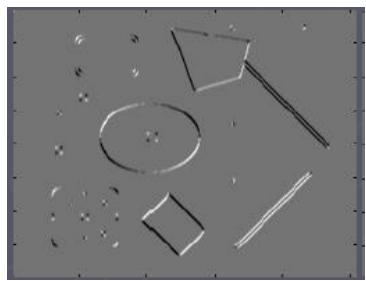
Sum over image region – the area we are checking for corner

Gradient with respect to x , times gradient with respect to y

$$M = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \ I_y]$$



Harris Detector Formulation

Image I  I_x  I_y  $I_x I_y$

where M is a 2×2 matrix computed from image derivatives:

$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

↑
Sum over image region – the area we are
checking for corner

**Gradient with
respect to x ,
times gradient
with respect to y**

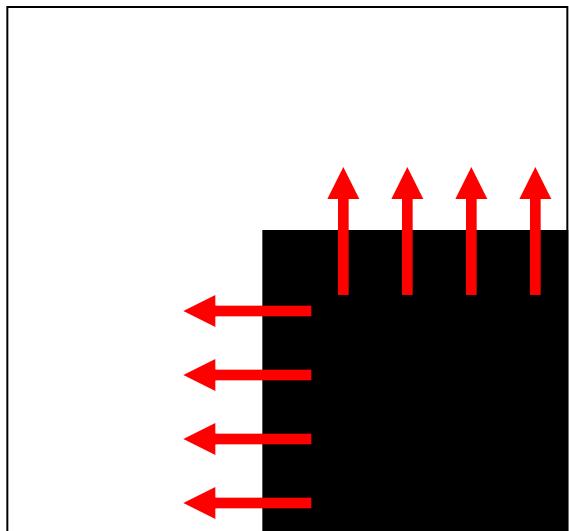
$$M = \begin{bmatrix} \sum I_x I_x & \sum I_x I_y \\ \sum I_x I_y & \sum I_y I_y \end{bmatrix} = \sum \begin{bmatrix} I_x \\ I_y \end{bmatrix} [I_x \ I_y]$$



What Does This Matrix Reveal?

- First, let's consider an axis-aligned corner:

$$M = \begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{bmatrix} = \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix}$$

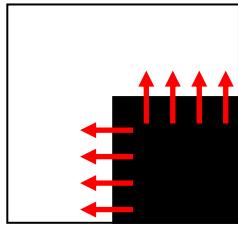




What Does This Matrix Reveal?

- First, let's consider an axis-aligned corner:

$$M = \begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_x I_y & \sum I_y^2 \end{bmatrix} = \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix}$$



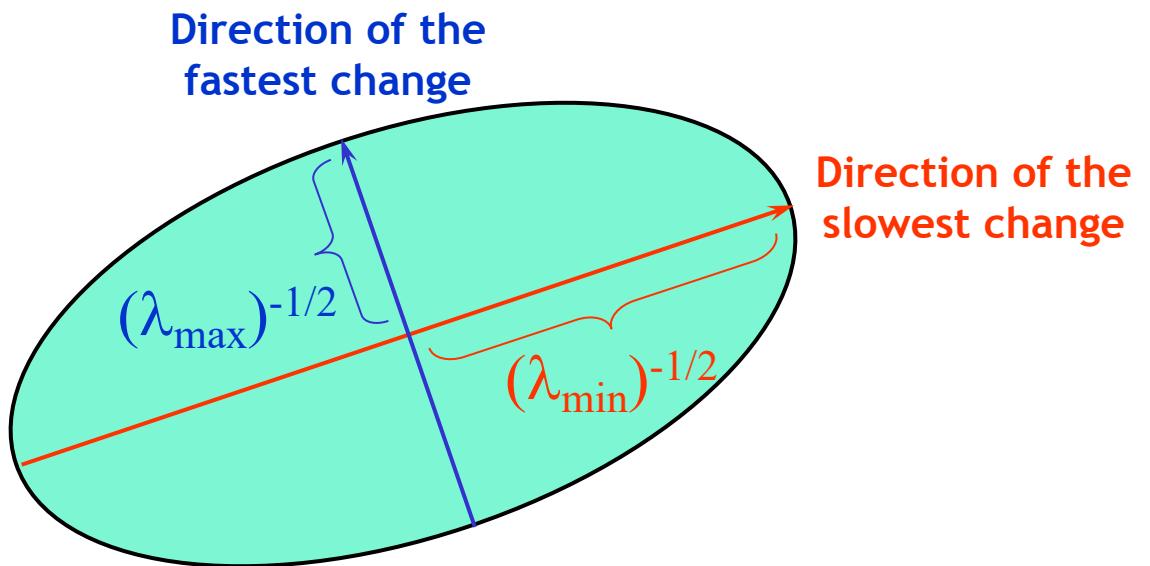
- This means:
 - Dominant gradient directions align with x or y axis
 - If either λ is close to 0, then this is not a corner, so look for locations where both are large.
- What if we have a corner that is not aligned with the image axes?



General Case

- Since M is symmetric, we have
$$M = R^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} R$$

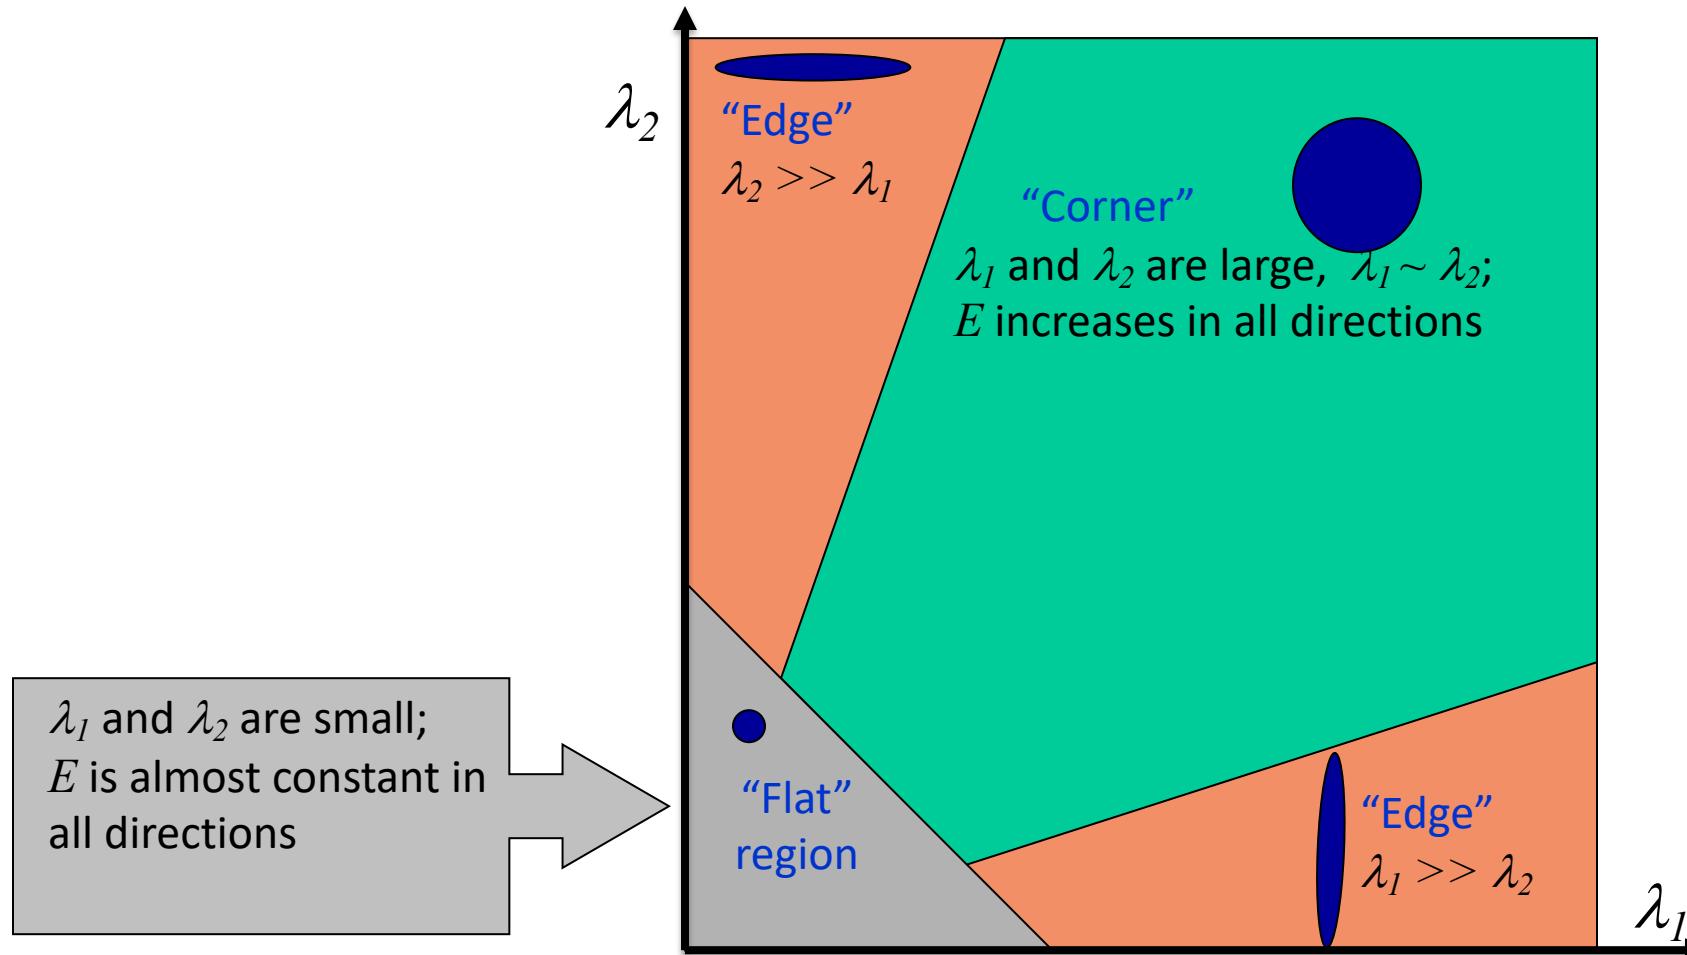
(Eigenvalue decomposition)
- We can think of M as an ellipse with axis lengths determined by the eigenvalues and orientation determined by R





Interpreting the Eigenvalues

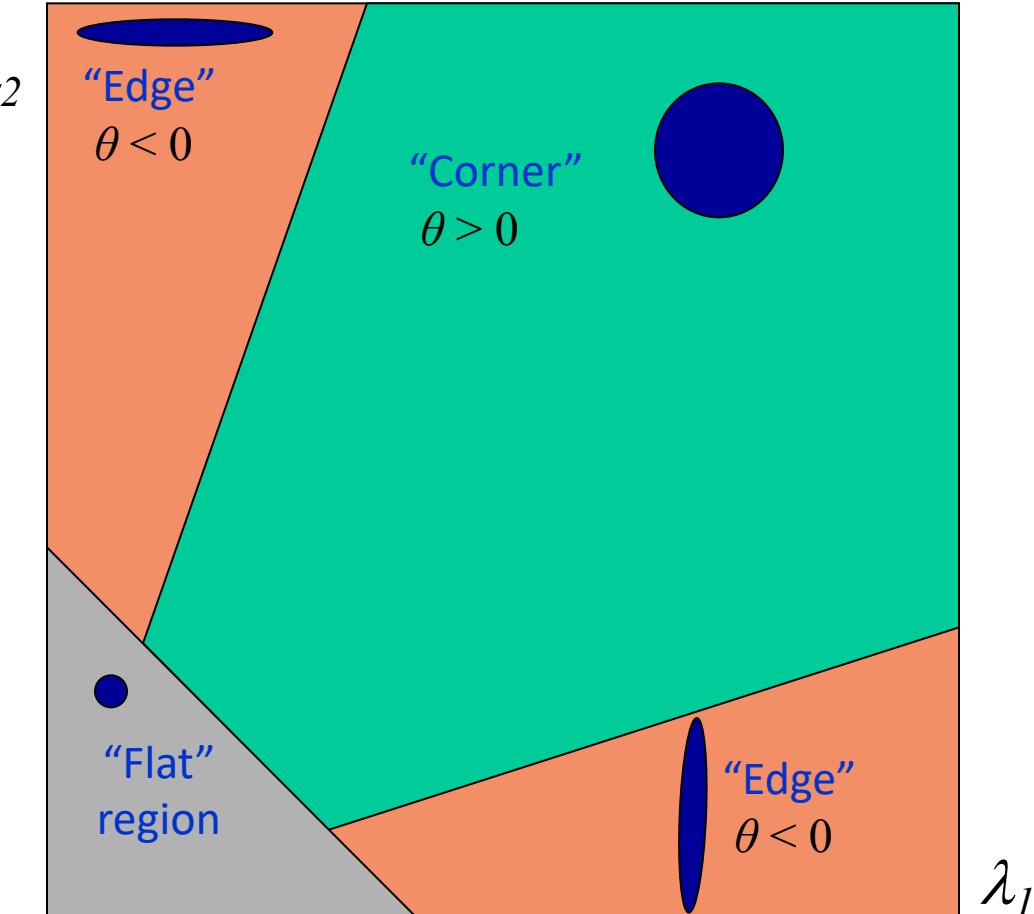
- Classification of image points using eigenvalues of M :





Corner Response Function

$$\theta = \det(M) - \alpha \operatorname{trace}(M)^2 = \lambda_1 \lambda_2 - \alpha(\lambda_1 + \lambda_2)^2$$



- Fast approximation
 - Avoid computing the eigenvalues
 - α : constant (0.04 to 0.06)

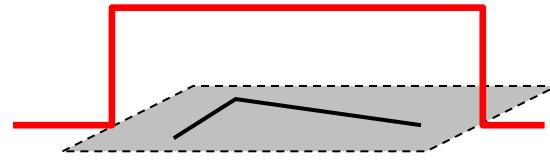
Window Function $w(x,y)$

$$M = \sum_{x,y} w(x,y) \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- Option 1: uniform window
 - Sum over square window

$$M = \sum_{x,y} \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- Problem: not rotation invariant

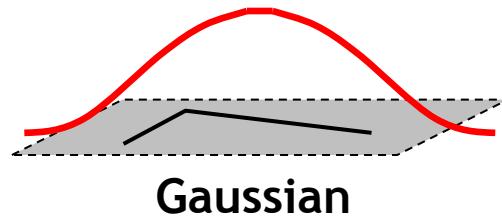


1 in window, 0 outside

- Option 2: Smooth with Gaussian
 - Gaussian already performs weighted sum

$$M = g(\sigma) * \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix}$$

- Result is rotation invariant



Gaussian



Summary: Harris Detector [Harris88]

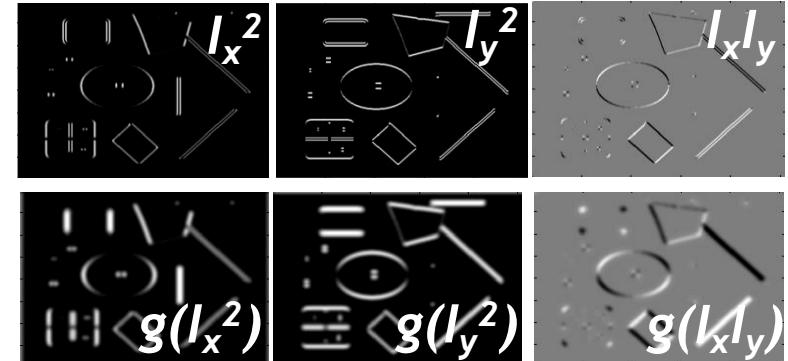
- Compute second moment matrix (autocorrelation matrix)

$$M(\sigma_I, \sigma_D) = g(\sigma_I) * \begin{bmatrix} I_x^2(\sigma_D) & I_x I_y(\sigma_D) \\ I_x I_y(\sigma_D) & I_y^2(\sigma_D) \end{bmatrix}$$

σ_D : for Gaussian in the derivative calculation
 σ_I : for Gaussian in the windowing function

2. Square of derivatives

1. Image derivatives



3. Gaussian filter $g(\sigma_I)$



4. Cornerness function - two strong eigenvalues

$$\begin{aligned}\theta &= \det[M(\sigma_I, \sigma_D)] - \alpha[\text{trace}(M(\sigma_I, \sigma_D))]^2 \\ &= g(I_x^2)g(I_y^2) - [g(I_x I_y)]^2 - \alpha[g(I_x^2) + g(I_y^2)]^2\end{aligned}$$

5. Perform non-maximum suppression



Harris Detector: Workflow

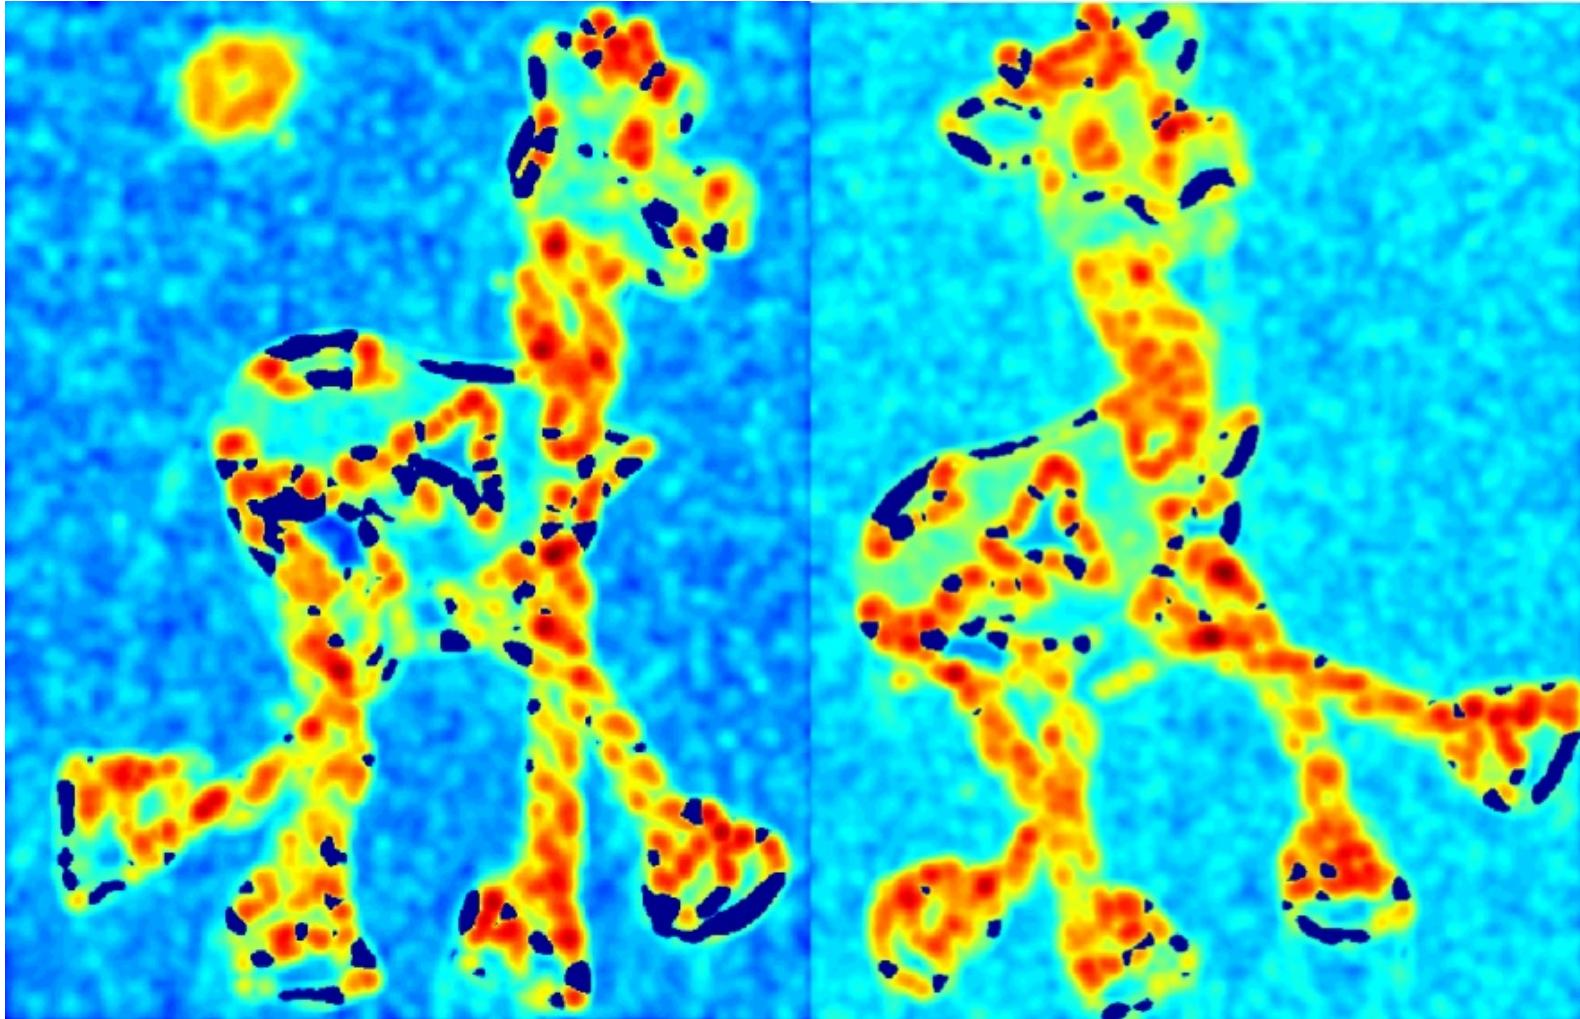


Slide adapted from Darya Frolova, Denis Simakov



Harris Detector: Workflow

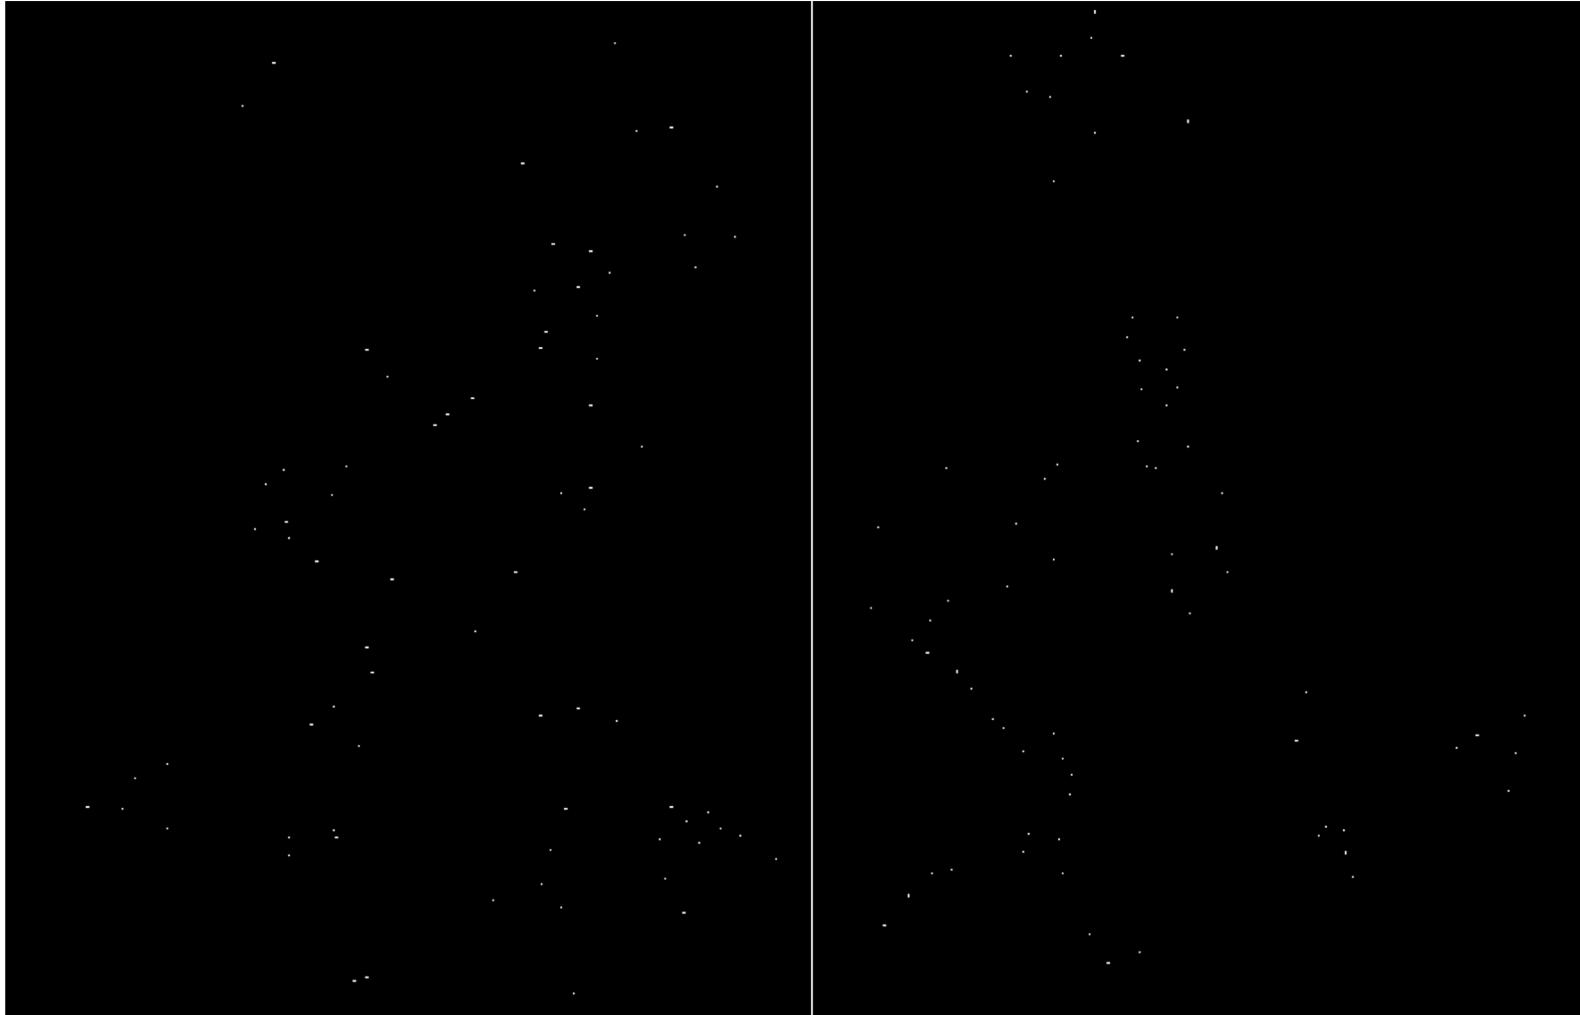
- computer corner responses θ





Harris Detector: Workflow

- Take only the local maxima of θ , where $\theta > \text{threshold}$





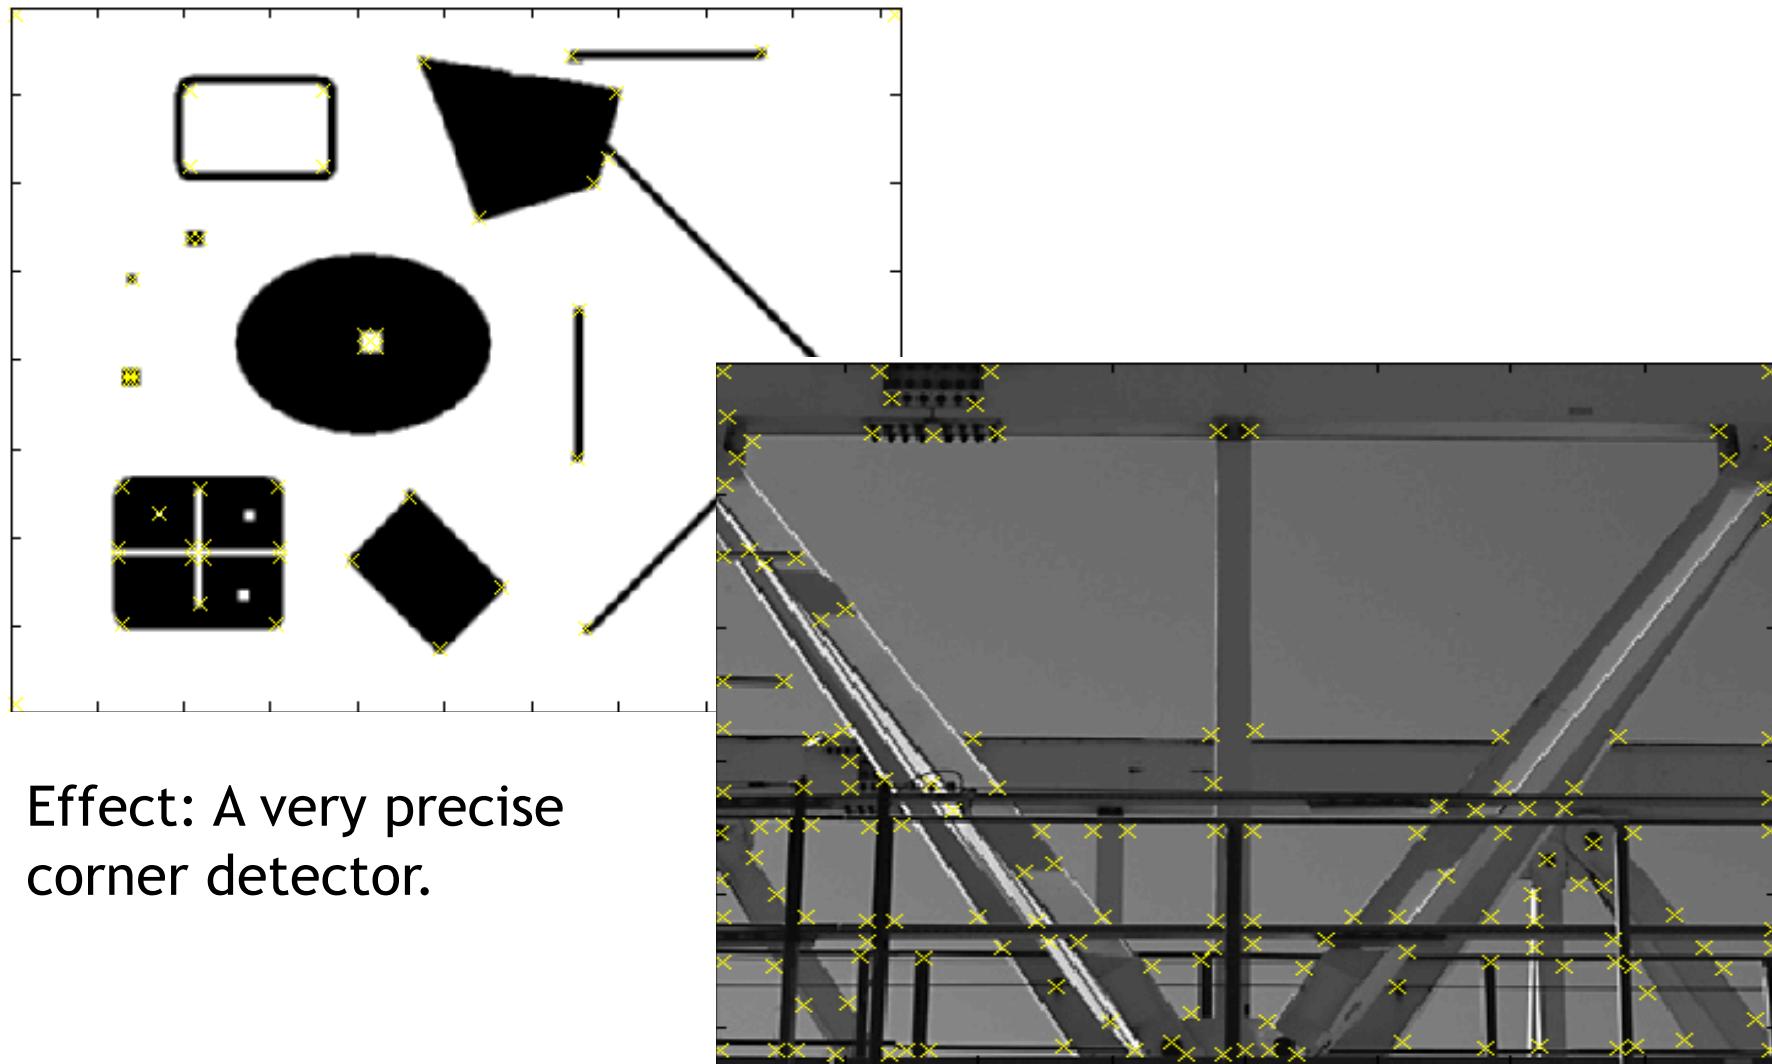
Harris Detector: Workflow - Resulting Harris points



Slide adapted from Darya Frolova, Denis Simakov



Harris Detector – Responses [Harris88]





Harris Detector – Responses [Harris88]





Harris Detector – Responses [Harris88]



- Results are well suited for finding stereo correspondences



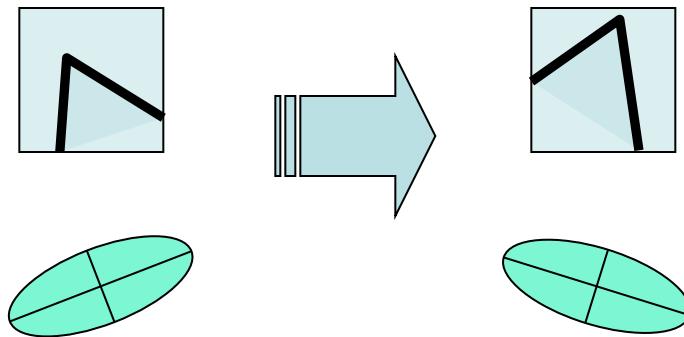
Harris Detector: Properties

- Translation invariance?



Harris Detector: Properties

- Translation invariance
- Rotation invariance?



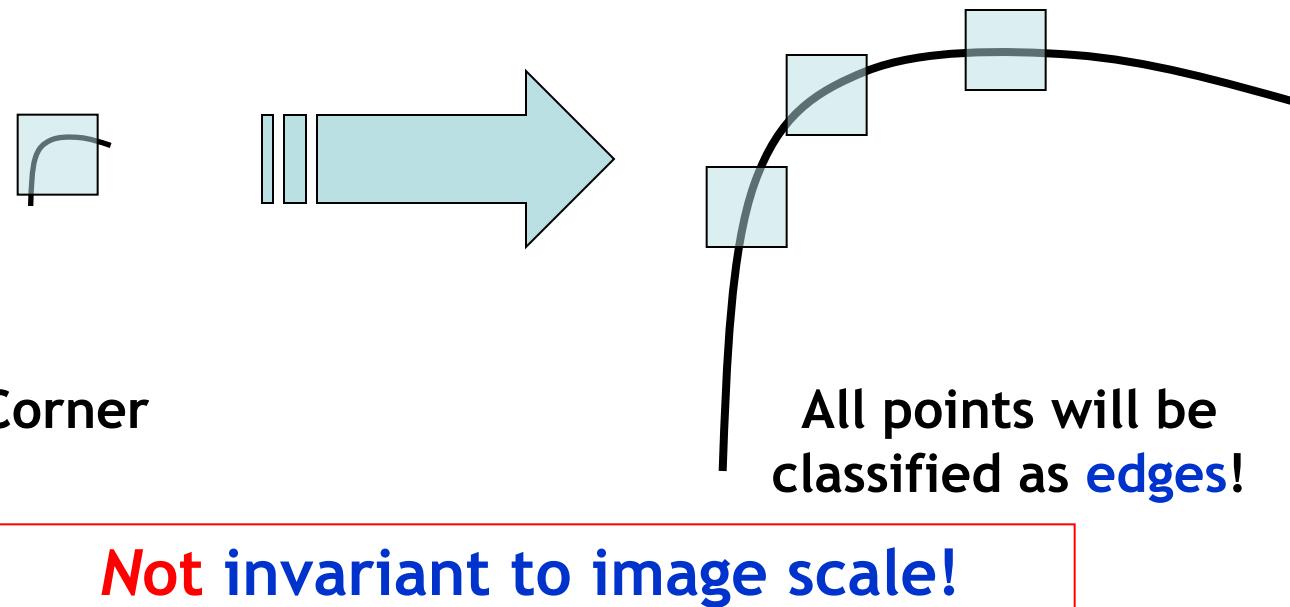
Ellipse rotates but its shape (i.e. eigenvalues) remains the same

Corner response θ is invariant to image rotation



Harris Detector: Properties

- Translation invariance
- Rotation invariance
- Scale invariance?





What we are learned today?

- A model fitting method for edge detection
 - RANSAC
- Local invariant features
 - Motivation
 - Requirements, invariances
- Keypoint localization
 - Harris corner detector