



# Lecture: Face Recognition

Juan Carlos Niebles and Ranjay Krishna  
Stanford Vision and Learning Lab



# CS 131 Roadmap





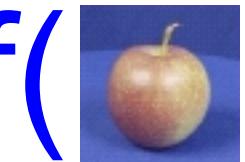
# Let's recap

- A simple object recognition pipeline with kNN
- PCA



# Object recognition: a classification framework

- Apply a prediction function to a feature representation of the image to get the desired output:

$$f(\text{apple}) = \text{"apple"}$$
$$f(\text{tomato}) = \text{"tomato"}$$
$$f(\text{cow}) = \text{"cow"}$$




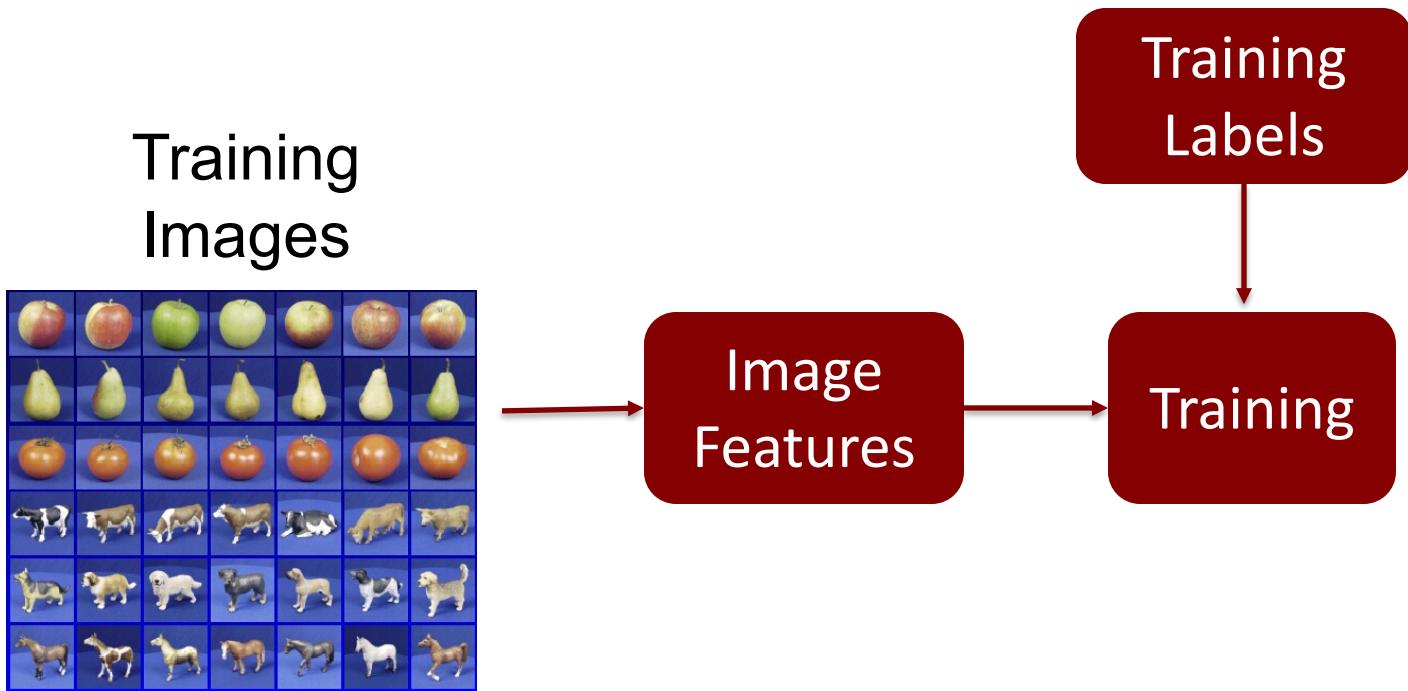
# A simple pipeline - Training

Training  
Images



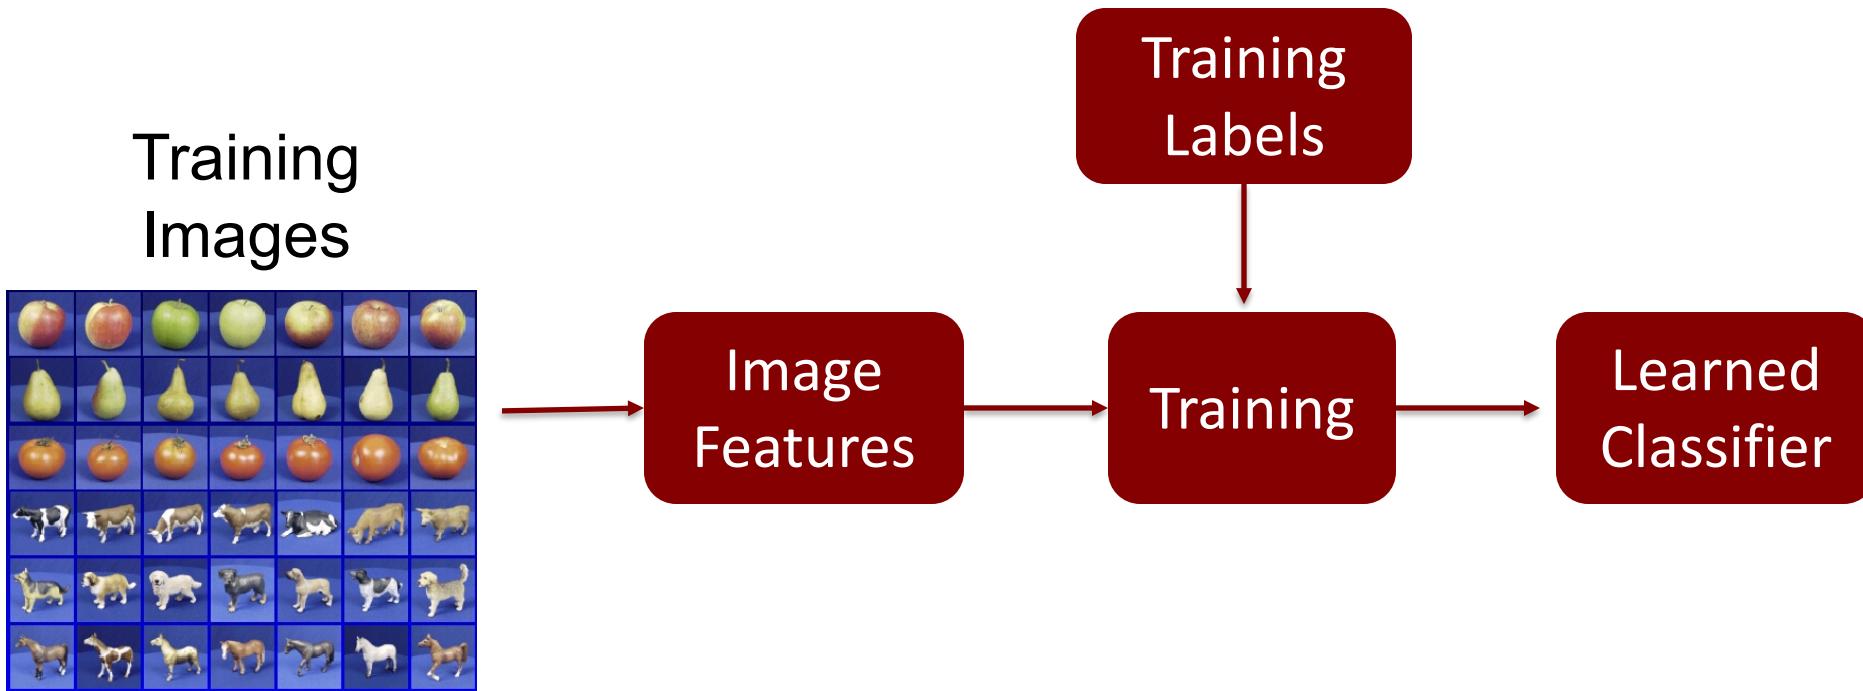


# A simple pipeline - Training



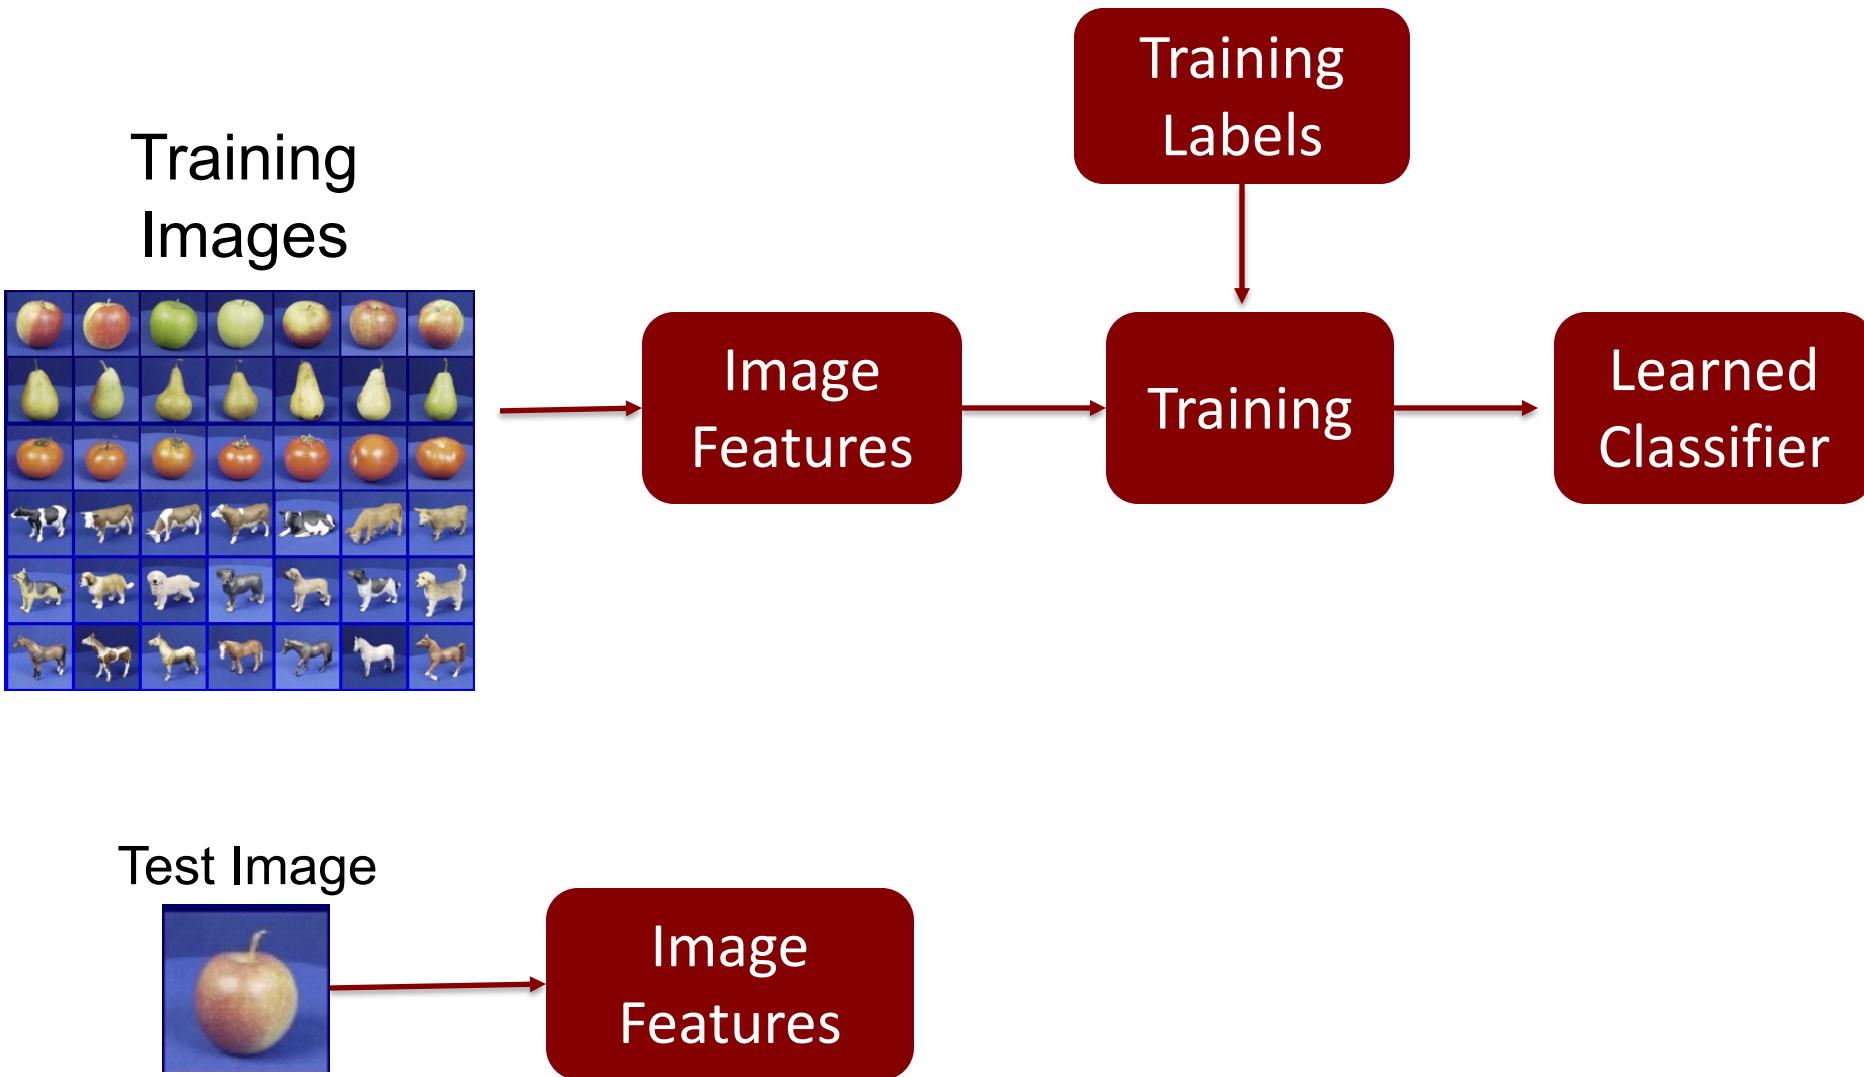


# A simple pipeline - Training



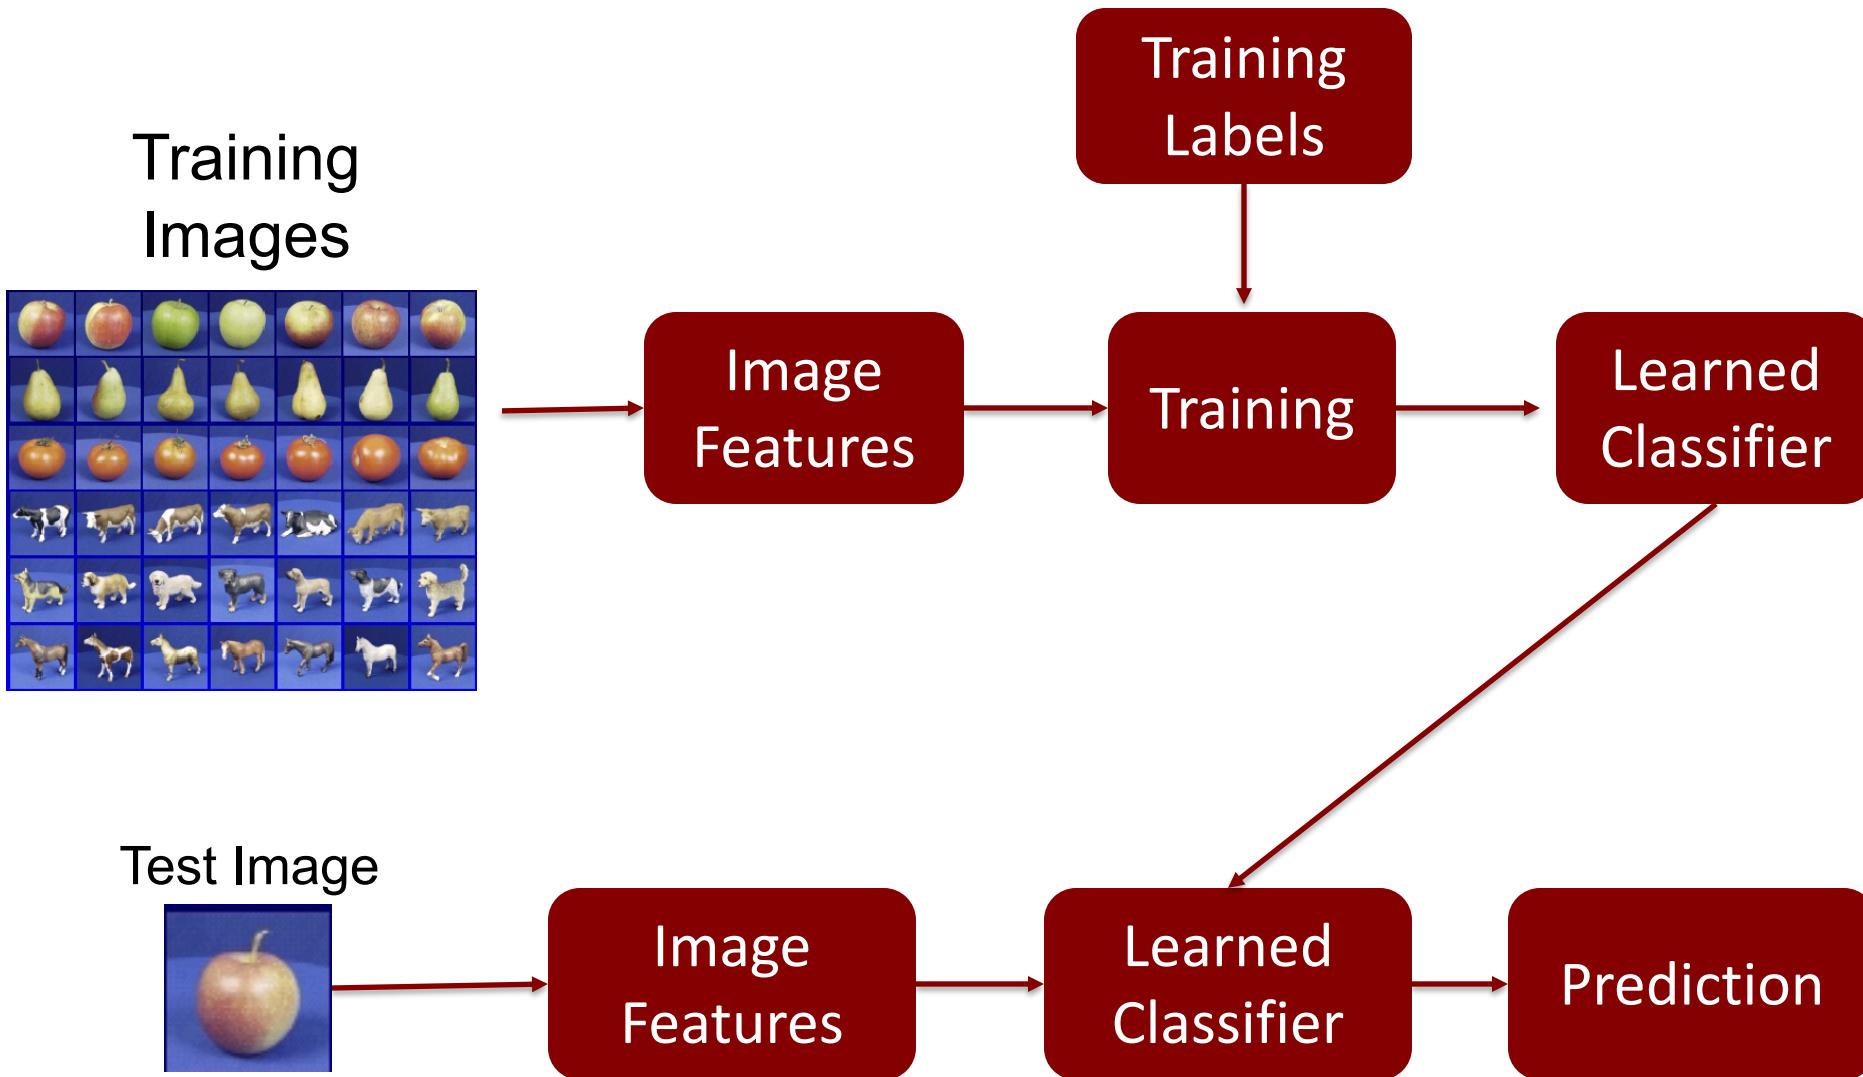


# A simple pipeline - Training



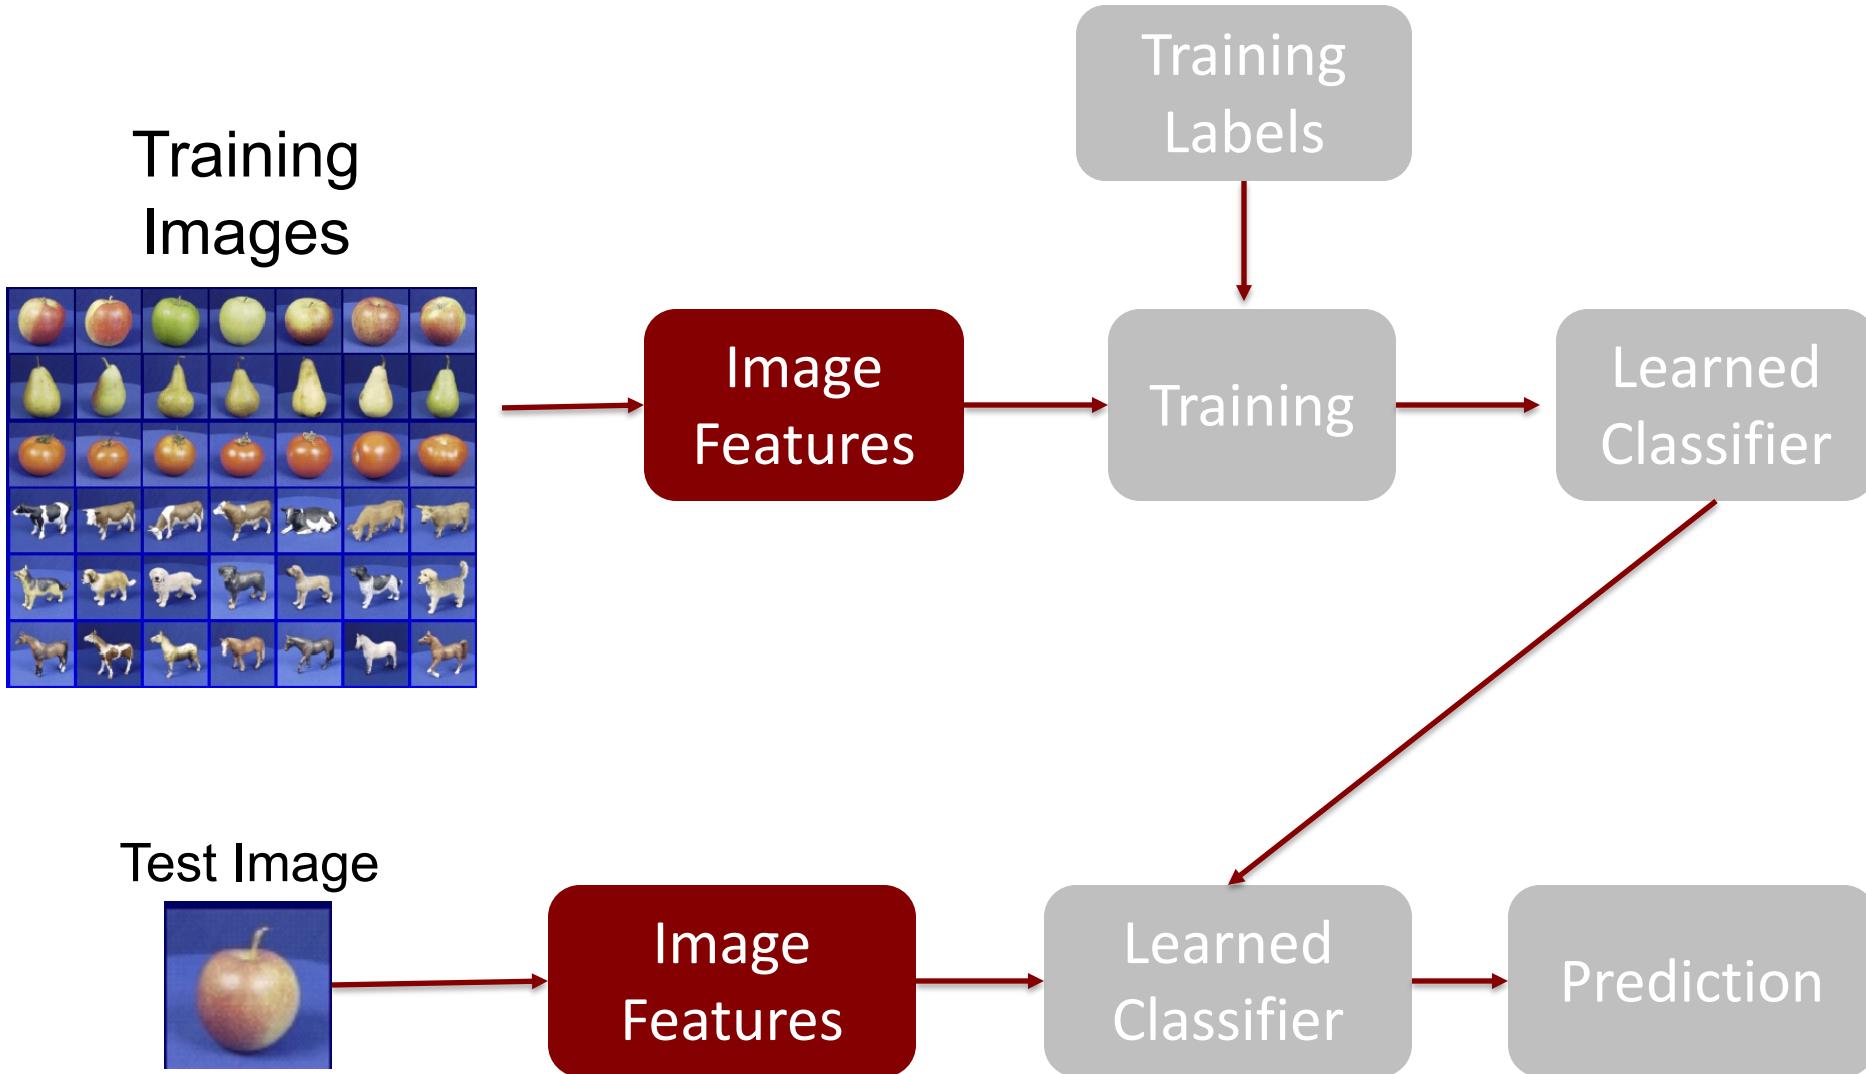


# A simple pipeline - Training





# A simple pipeline - Training



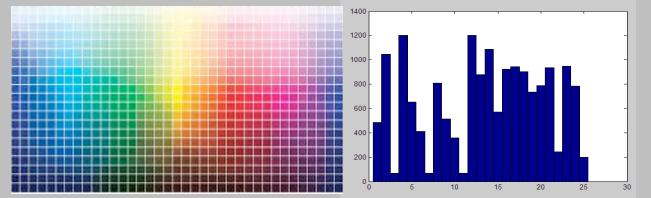


# Image features

Input image



Color: Quantize RGB values



Invariance?

- ? Translation
- ? Scale
- ? Rotation
- ? Occlusion

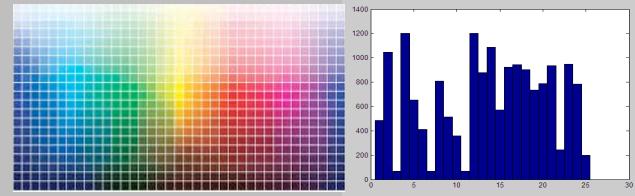


# Image features

Input image



Color: Quantize RGB values



Invariance?

- 😊 Translation
- 😢 Scale
- 😊 Rotation (in-planar)
- 😢 Occlusion

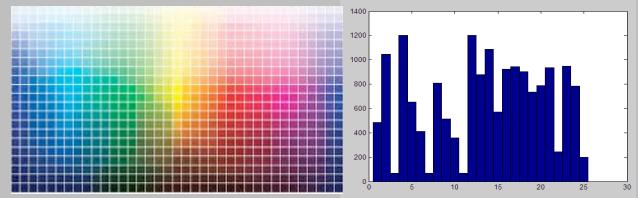


# Image features

Input image



Color: Quantize RGB values



Invariance?

- 😊 Translation
- 😢 Scale
- 😊 Rotation (in-planar)
- 😢 Occlusion

Global shape: PCA space



Invariance?

- ? Translation
- ? Scale
- ? Rotation (in-planar)
- ? Occlusion

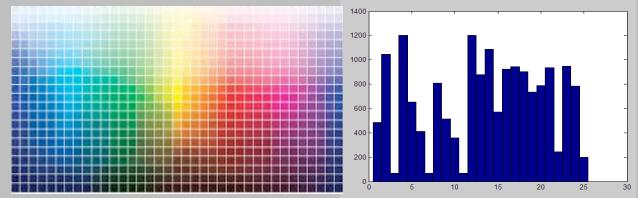


# Image features

Input image



Color: Quantize RGB values



Invariance?

- 😊 Translation
- 😢 Scale
- 😊 Rotation (in-planar)
- 😢 Occlusion

Global shape: PCA space



Invariance?

- 😊 Translation
- 😢 Scale
- 😊 Rotation (in-planar)
- 😢 Occlusion

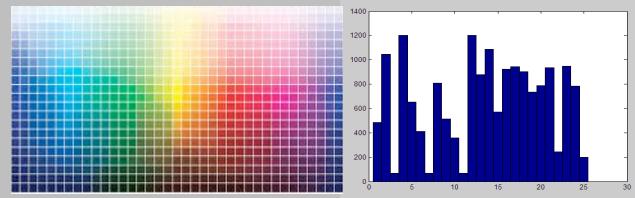


# Image features

Input image



Color: Quantize RGB values



Invariance?

- 😊 Translation
- 😢 Scale
- 😊 Rotation
- 😢 Occlusion

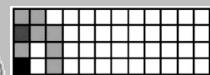
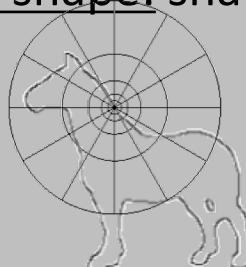
Global shape: PCA space



Invariance?

- 😊 Translation
- 😢 Scale
- 😊 Rotation
- 😢 Occlusion

Local shape: shape context



Invariance?

- ? Translation
- ? Scale
- ? Rotation (in-planar)
- ? Occlusion

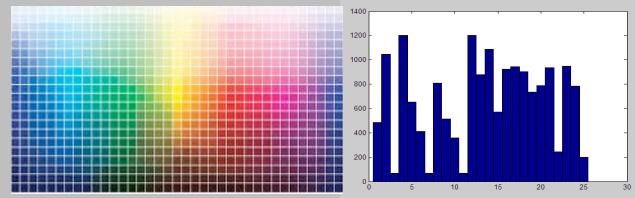


# Image features

Input image



Color: Quantize RGB values



Invariance?

- 😊 Translation
- 😊 Scale
- 😊 Rotation
- 😢 Occlusion

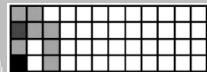
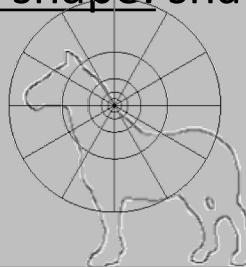
Global shape: PCA space



Invariance?

- 😊 Translation
- 😢 Scale
- 😊 Rotation
- 😢 Occlusion

Local shape: shape context



Invariance?

- 😊 Translation
- 😊 Scale
- 😊 Rotation (in-planar)
- 😢 Occlusion

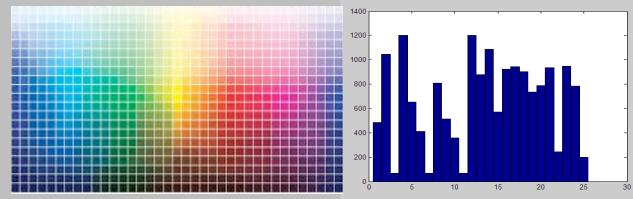


# Image features

Input image



Color: Quantize RGB values



Invariance?

- 😊 Translation
- 😊 Scale
- 😊 Rotation
- 😢 Occlusion

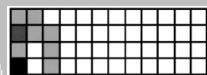
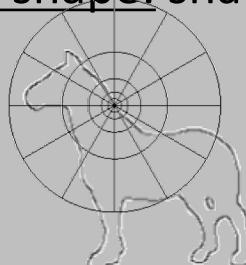
Global shape: PCA space



Invariance?

- 😊 Translation
- 😢 Scale
- 😊 Rotation
- 😢 Occlusion

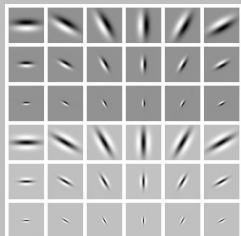
Local shape: shape context



Invariance?

- 😊 Translation
- 😊 Scale
- 😊 Rotation (in-planar)
- 😢 Occlusion

Texture: Filter banks



Invariance?

- ? Translation
- ? Scale
- ? Rotation (in-planar)
- ? Occlusion

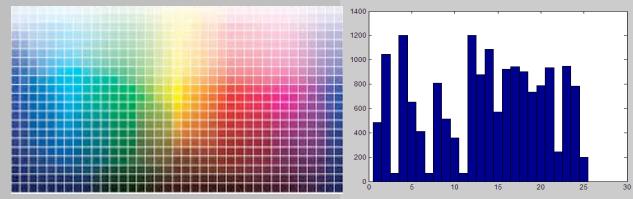


# Image features

Input image



Color: Quantize RGB values



Invariance?

- 😊 Translation
- 😊 Scale
- 😊 Rotation
- 😢 Occlusion

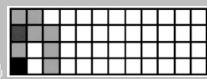
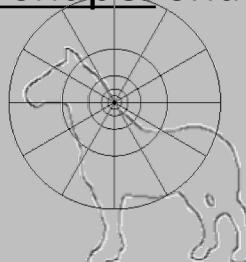
Global shape: PCA space



Invariance?

- 😊 Translation
- 😢 Scale
- 😊 Rotation
- 😢 Occlusion

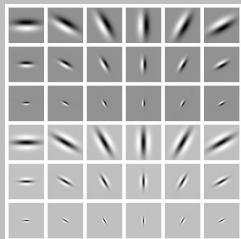
Local shape: shape context



Invariance?

- 😊 Translation
- 😊 Scale
- 😊 Rotation (in-planar)
- 😢 Occlusion

Texture: Filter banks

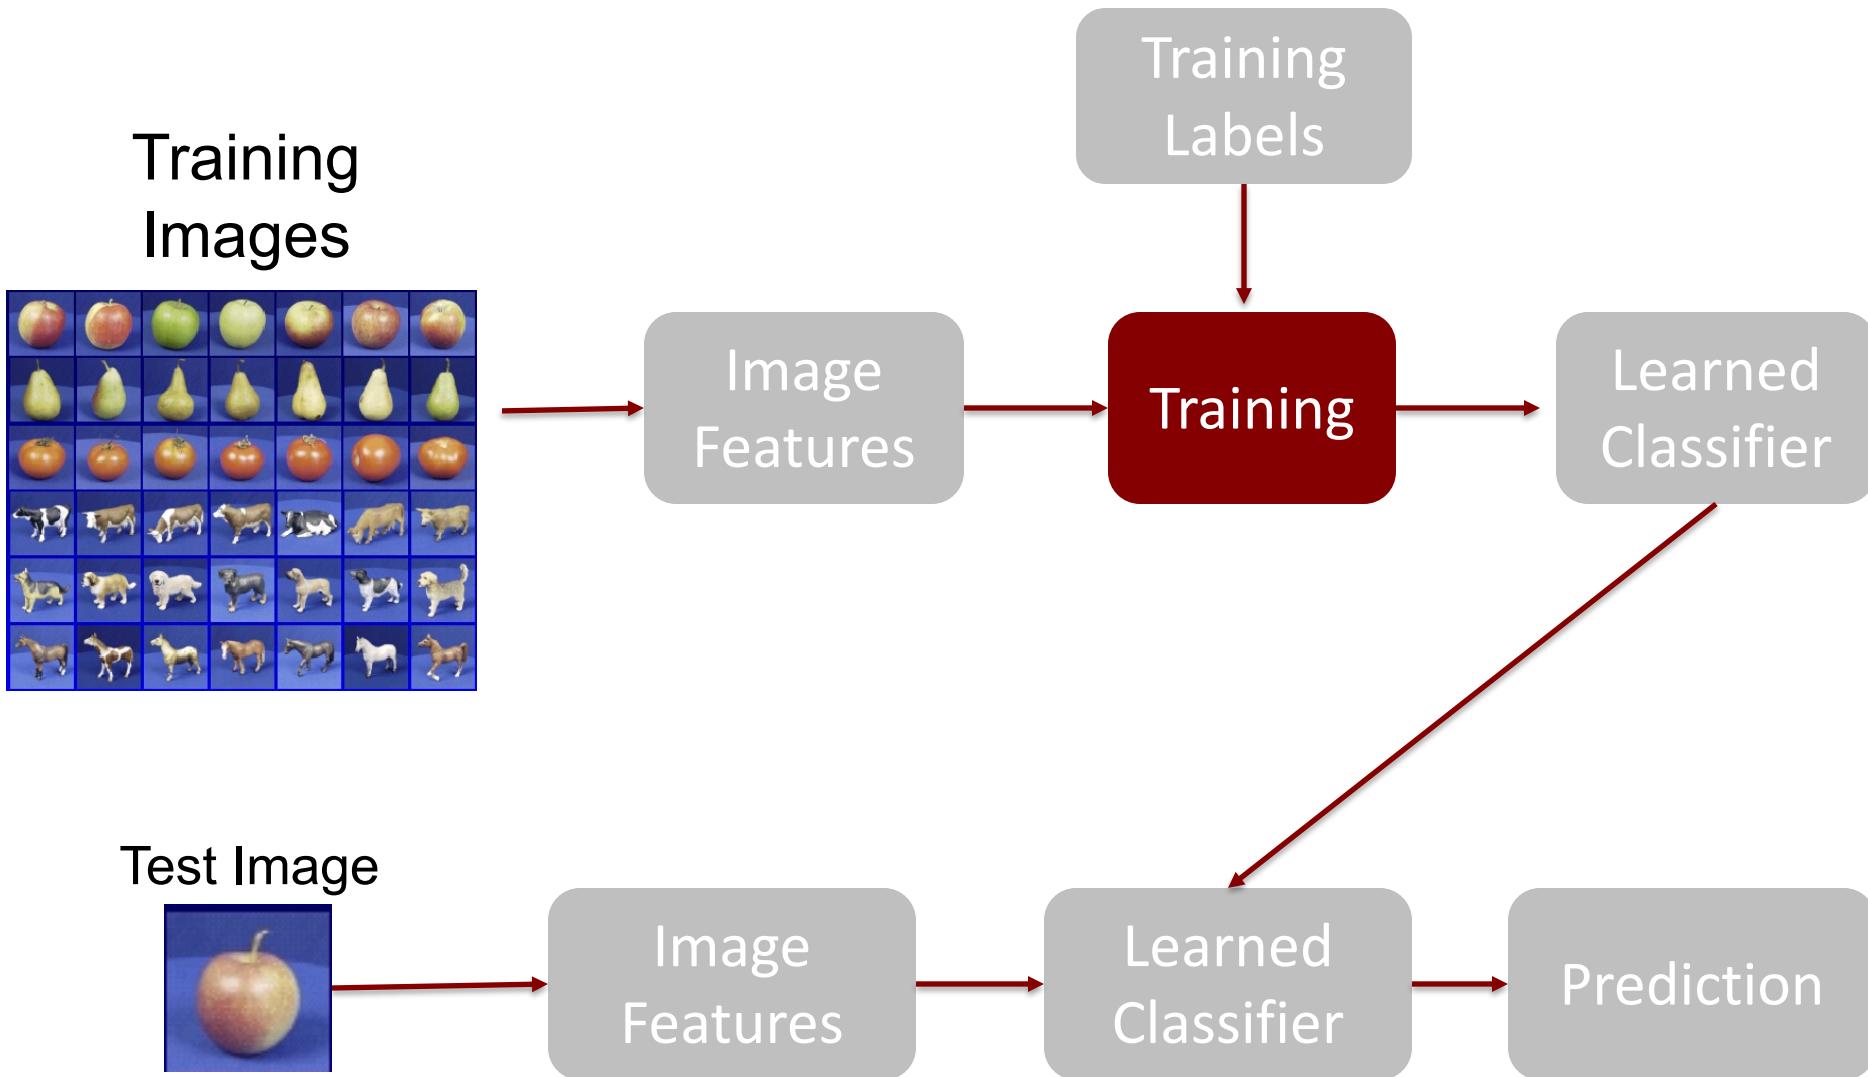


Invariance?

- 😊 Translation
- ? Scale
- ? Rotation (in-planar)
- 😢 Occlusion

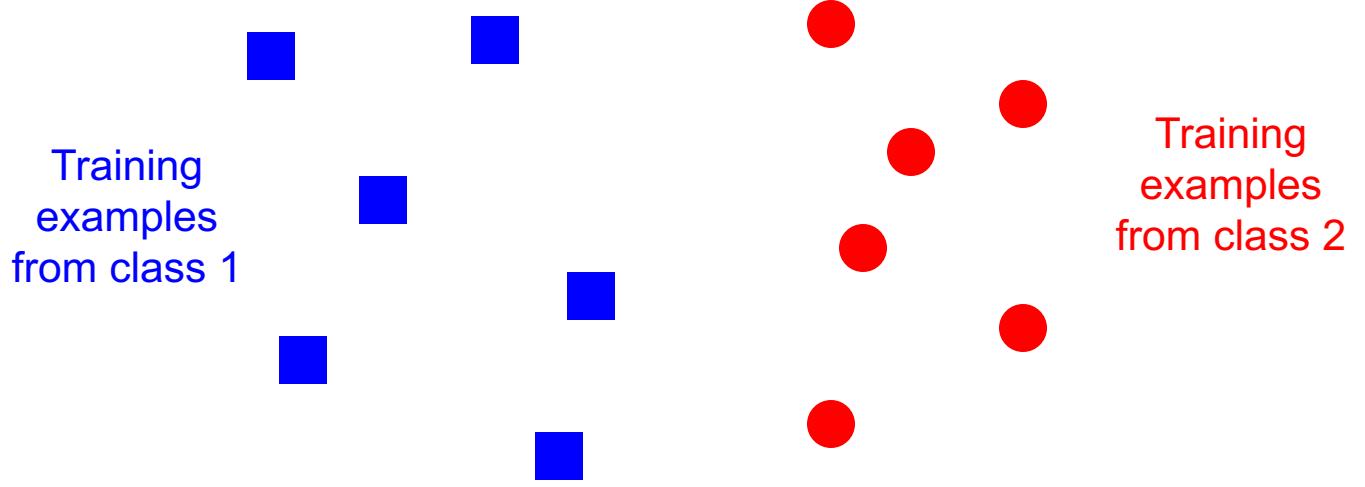


# A simple pipeline - Training



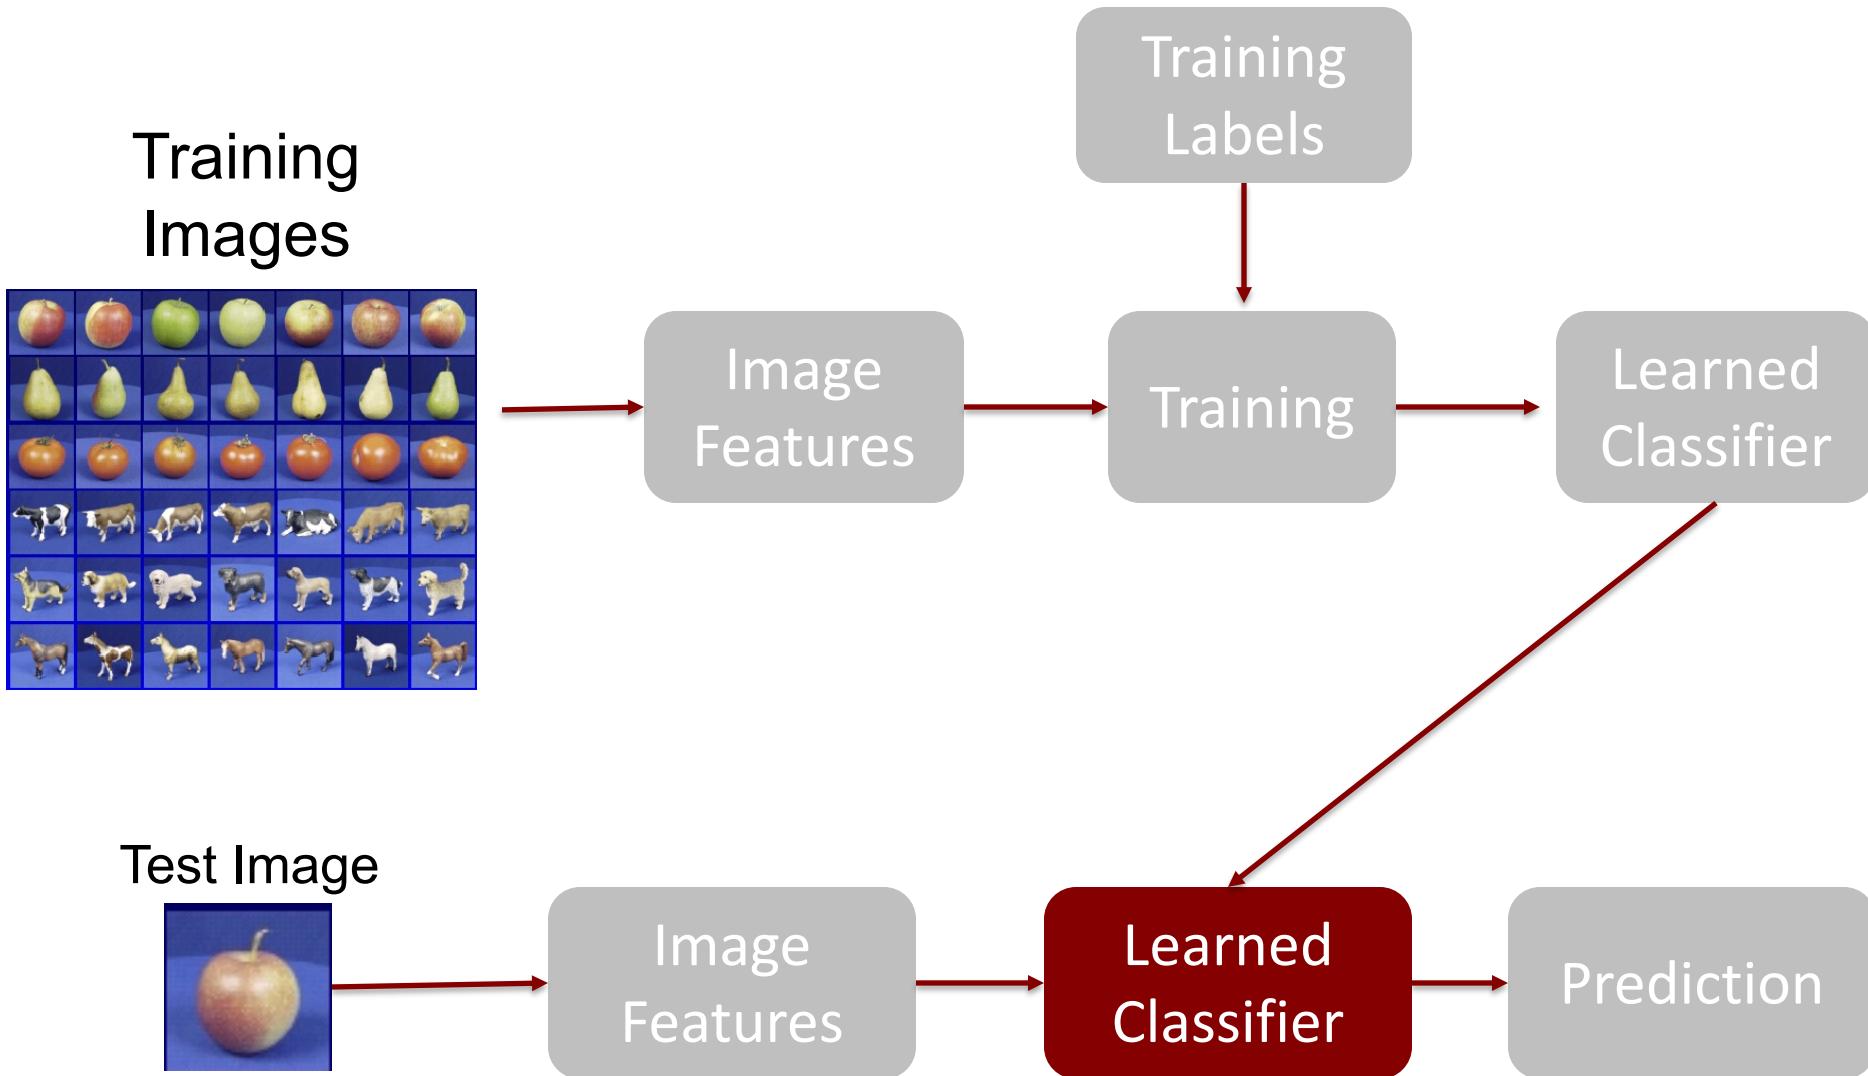


# Classifiers: Nearest neighbor



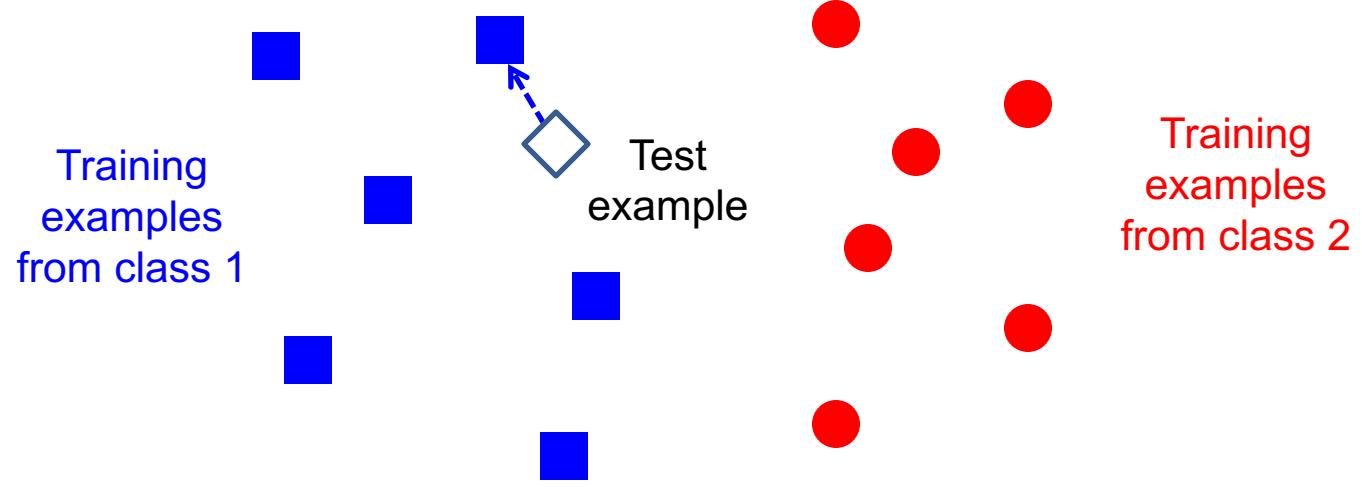


# A simple pipeline - Training





# Classifiers: Nearest neighbor





# Let's recap

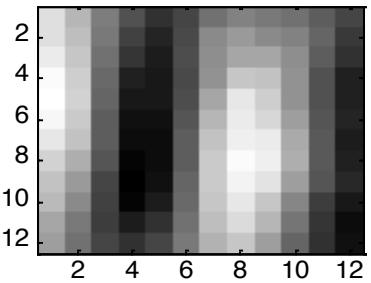
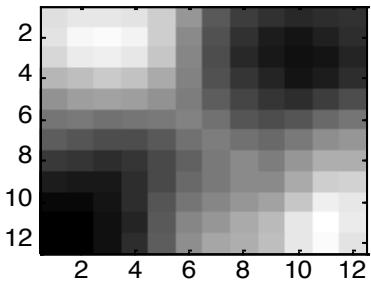
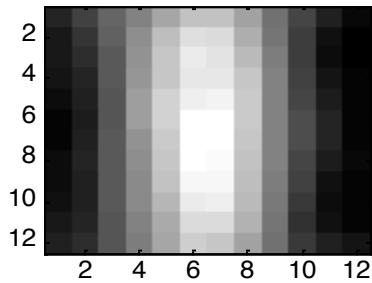
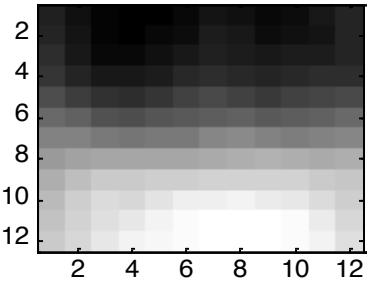
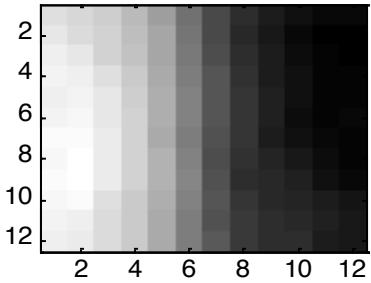
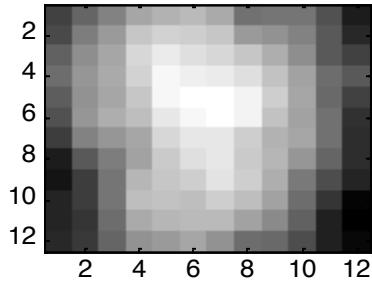
- A simple object recognition pipeline with kNN
- PCA



PCA compression: 144D  $\rightarrow$  6D



# 6 most important eigenvectors

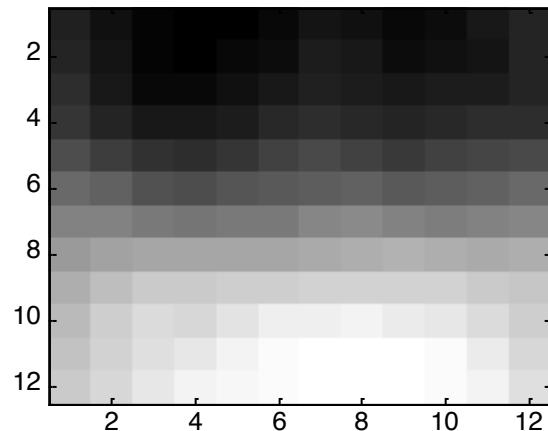
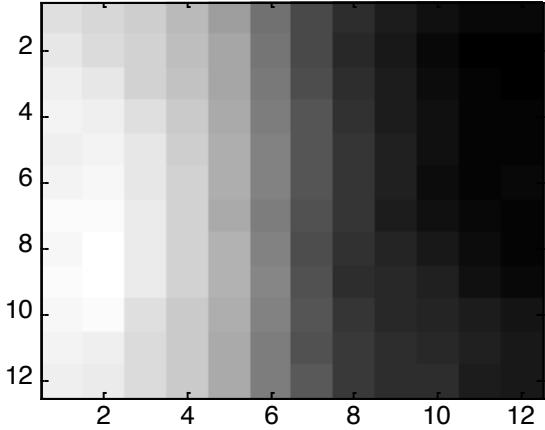
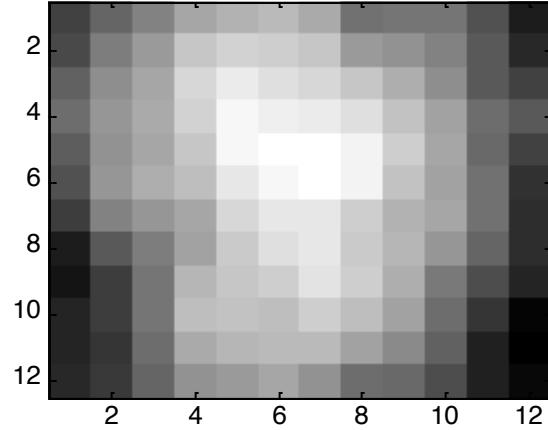




PCA compression: 144D → 3D



# 3 most important eigenvectors





# What we will learn today

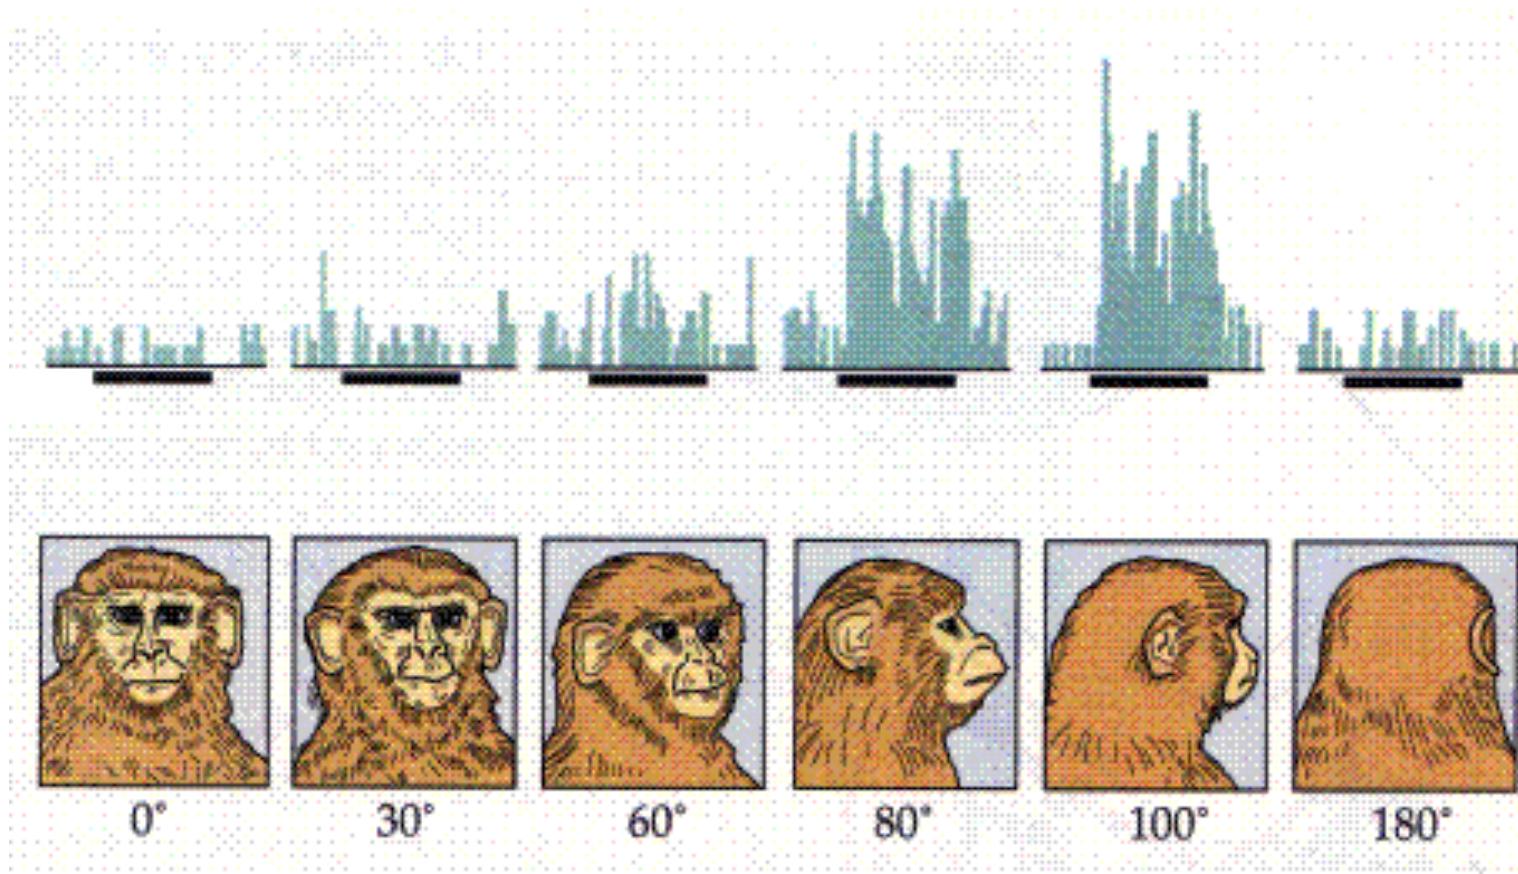
- Introduction to face recognition
- The Eigenfaces Algorithm
- Linear Discriminant Analysis (LDA)

Turk and Pentland, Eigenfaces for Recognition, *Journal of Cognitive Neuroscience* **3** (1): 71–86.

P. Belhumeur, J. Hespanha, and D. Kriegman. "Eigenfaces vs. Fisherfaces: Recognition Using Class Specific Linear Projection". *IEEE Transactions on pattern analysis and machine intelligence* **19** (7): 711. 1997.



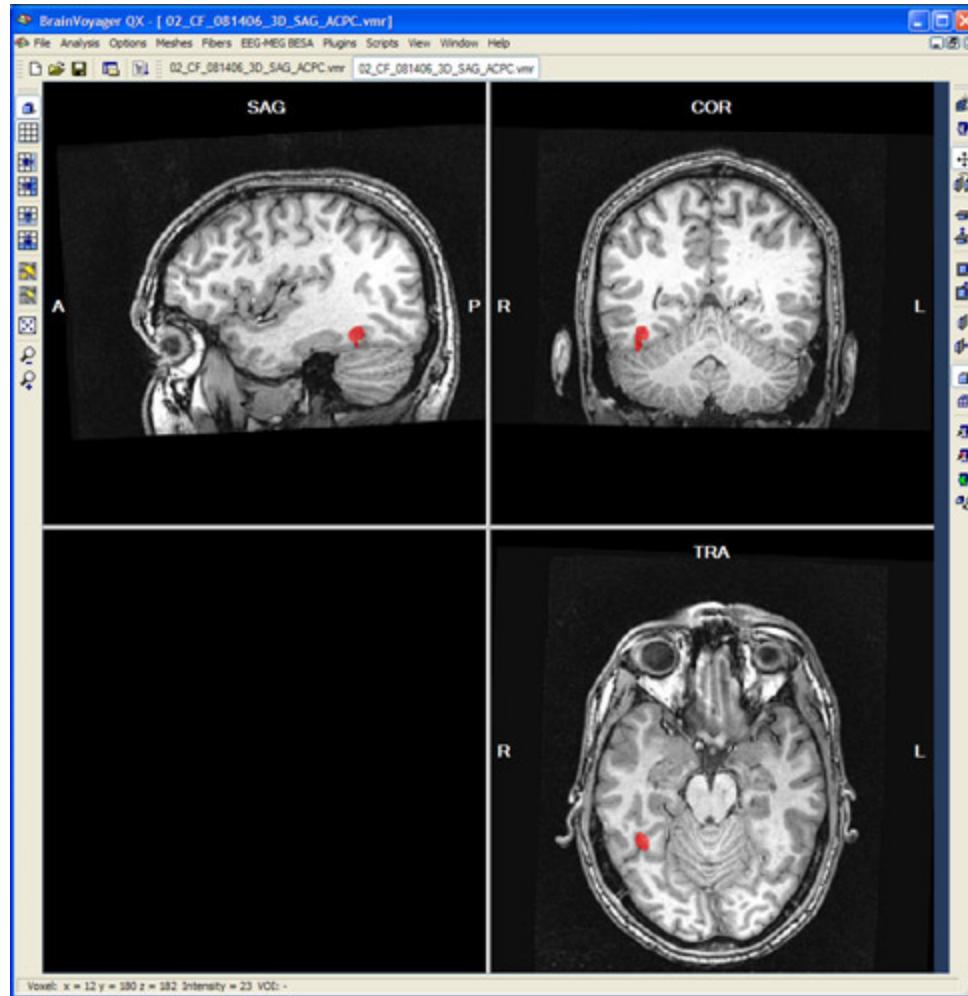
# “Faces” in the brain



Courtesy of Johannes M. Zanker



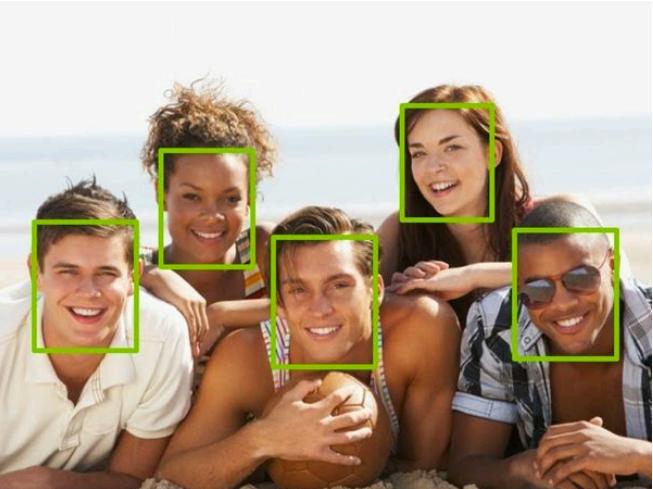
# “Faces” in the brain fusiform face area



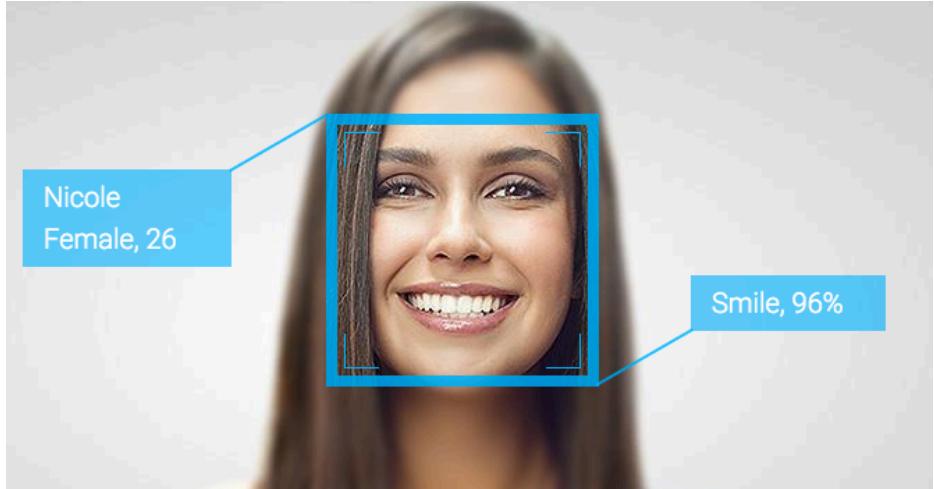
Kanwisher, et al. 1997



# Detection versus Recognition



Detection finds the faces in images



Recognition recognizes WHO the person is



# Face Recognition

- Digital photography





# Face Recognition

- Digital photography
- Surveillance

■ Recording

Detecting....

**Matching with Database**

Name: Alireza,  
Date: 25 My 2007 15:45  
Place: Main corridor

Name: Unknown  
Date: 25 My 2007 15:45  
Place: Main corridor

Report



# Face Recognition

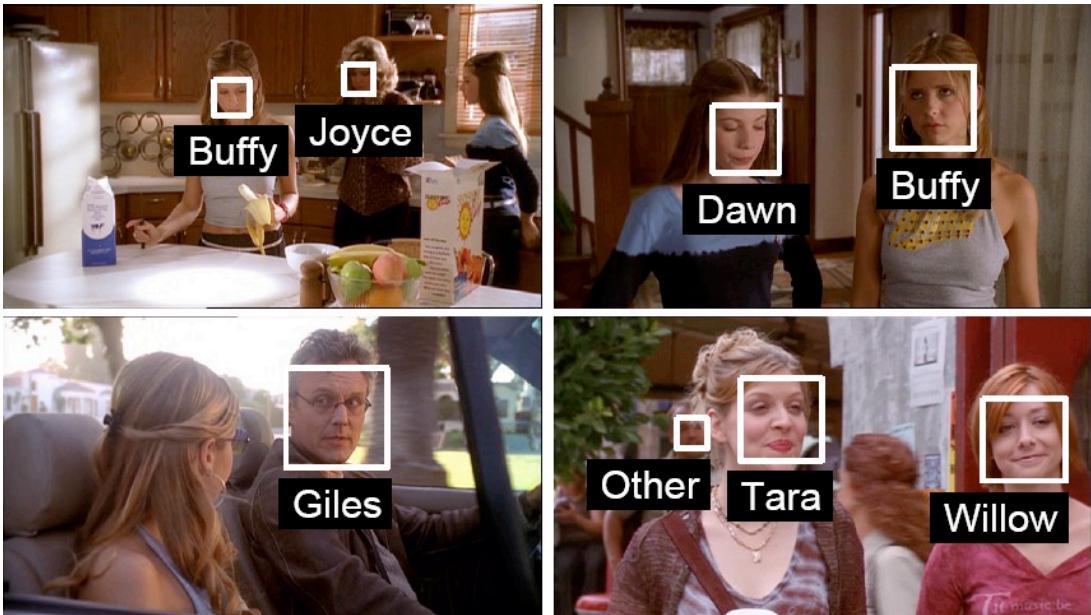
- Digital photography
- Surveillance
- Album organization





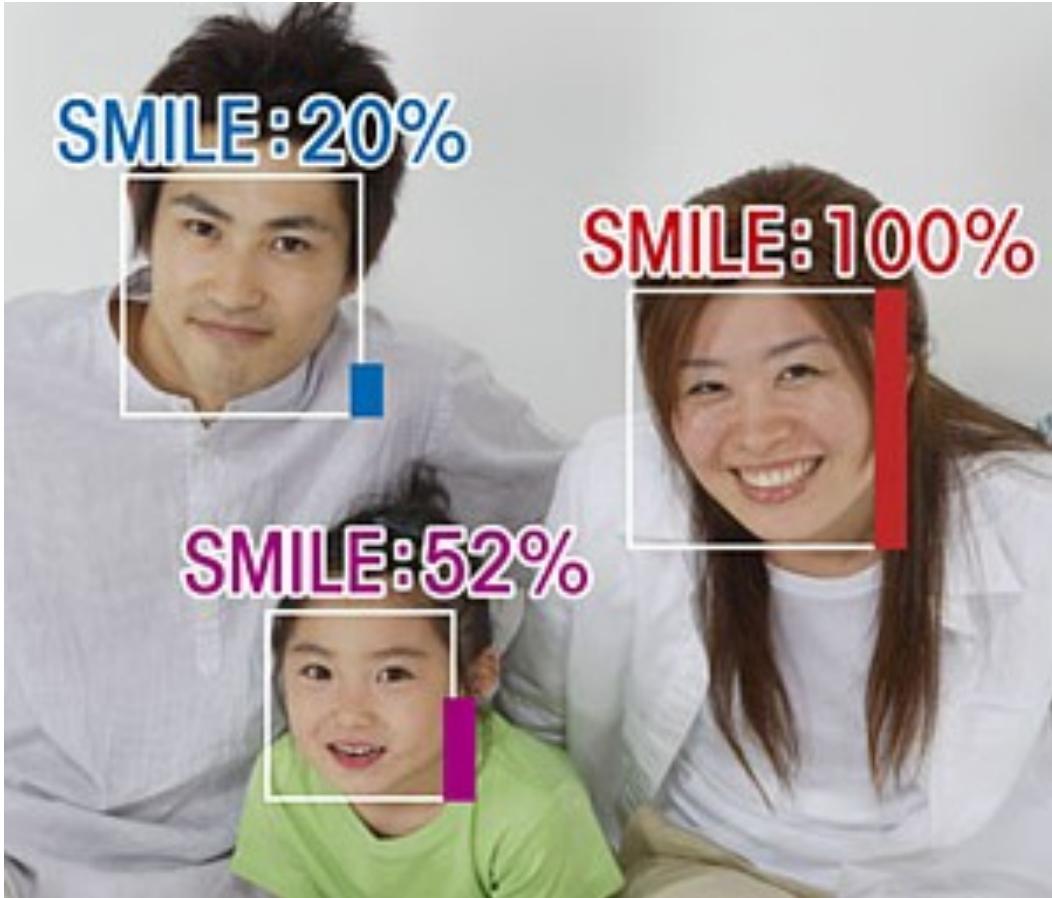
# Face Recognition

- Digital photography
- Surveillance
- Album organization
- Person tracking/id.



# Face Recognition

- Digital photography
- Surveillance
- Album organization
- Person tracking/id.
- Emotions and expressions





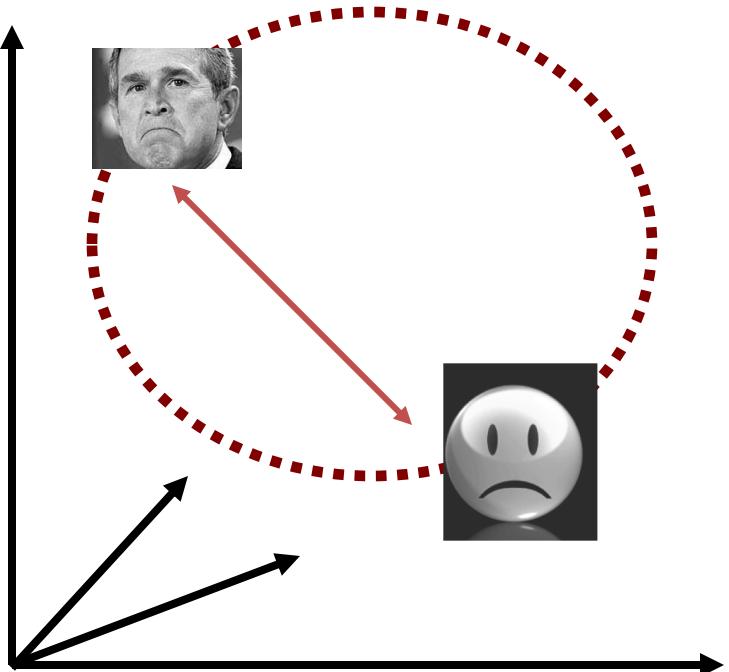
# Face Recognition

- Digital photography
- Surveillance
- Album organization
- Person tracking/id.
- Emotions and expressions
- Security/warfare
- Tele-conferencing
- Etc.



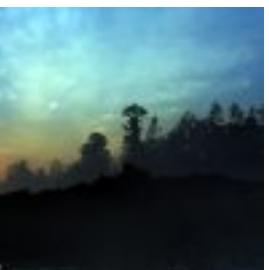
# The Space of Faces

- An image is a point in a high dimensional space
  - If represented in grayscale intensity, an  $N \times M$  image is a point in  $R^{NM}$
  - E.g. 100x100 image = 10,000 dim



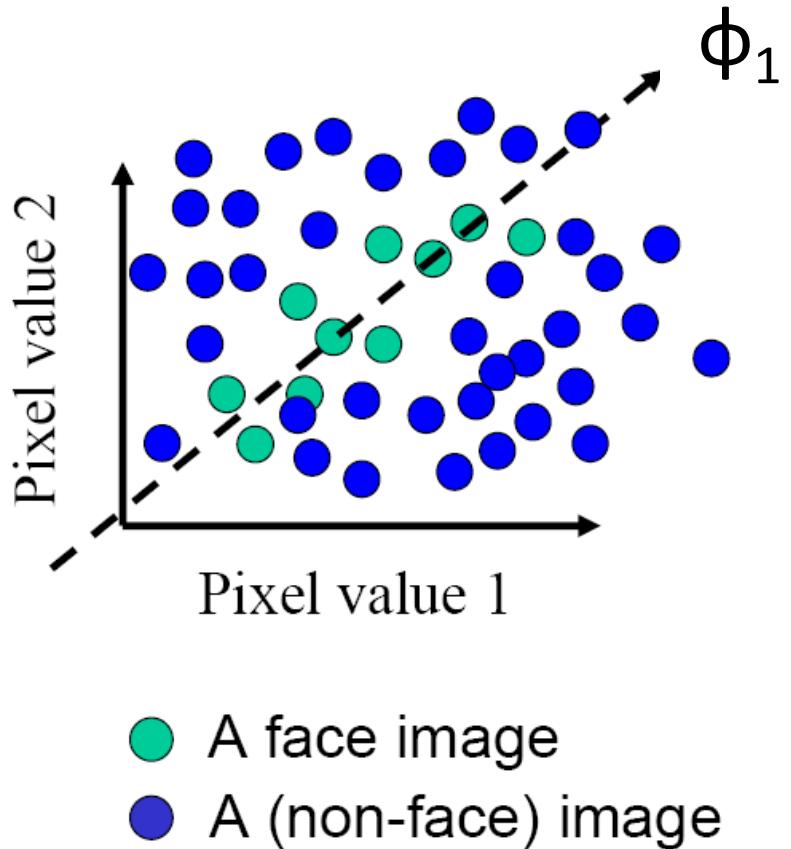
Slide credit: Chuck Dyer, Steve Seitz, Nishino

# 100x100 images can contain many things other than faces!





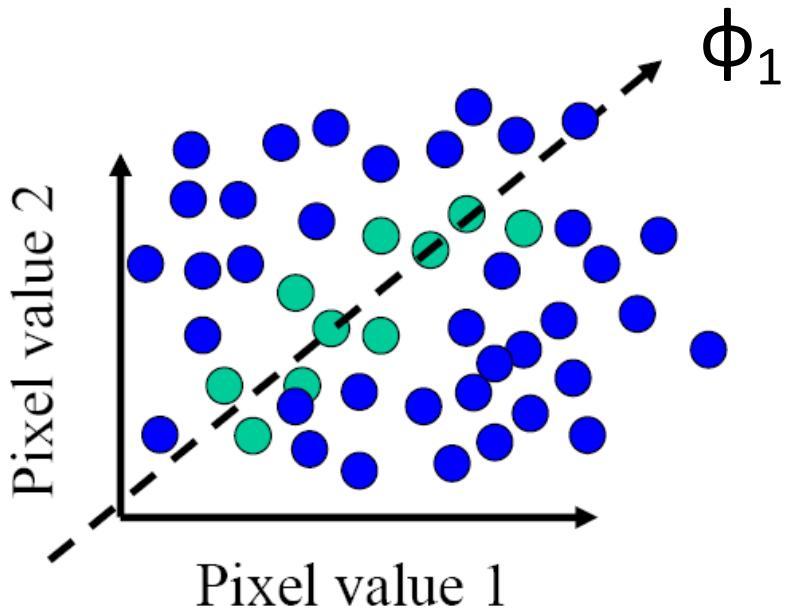
# The Space of Faces



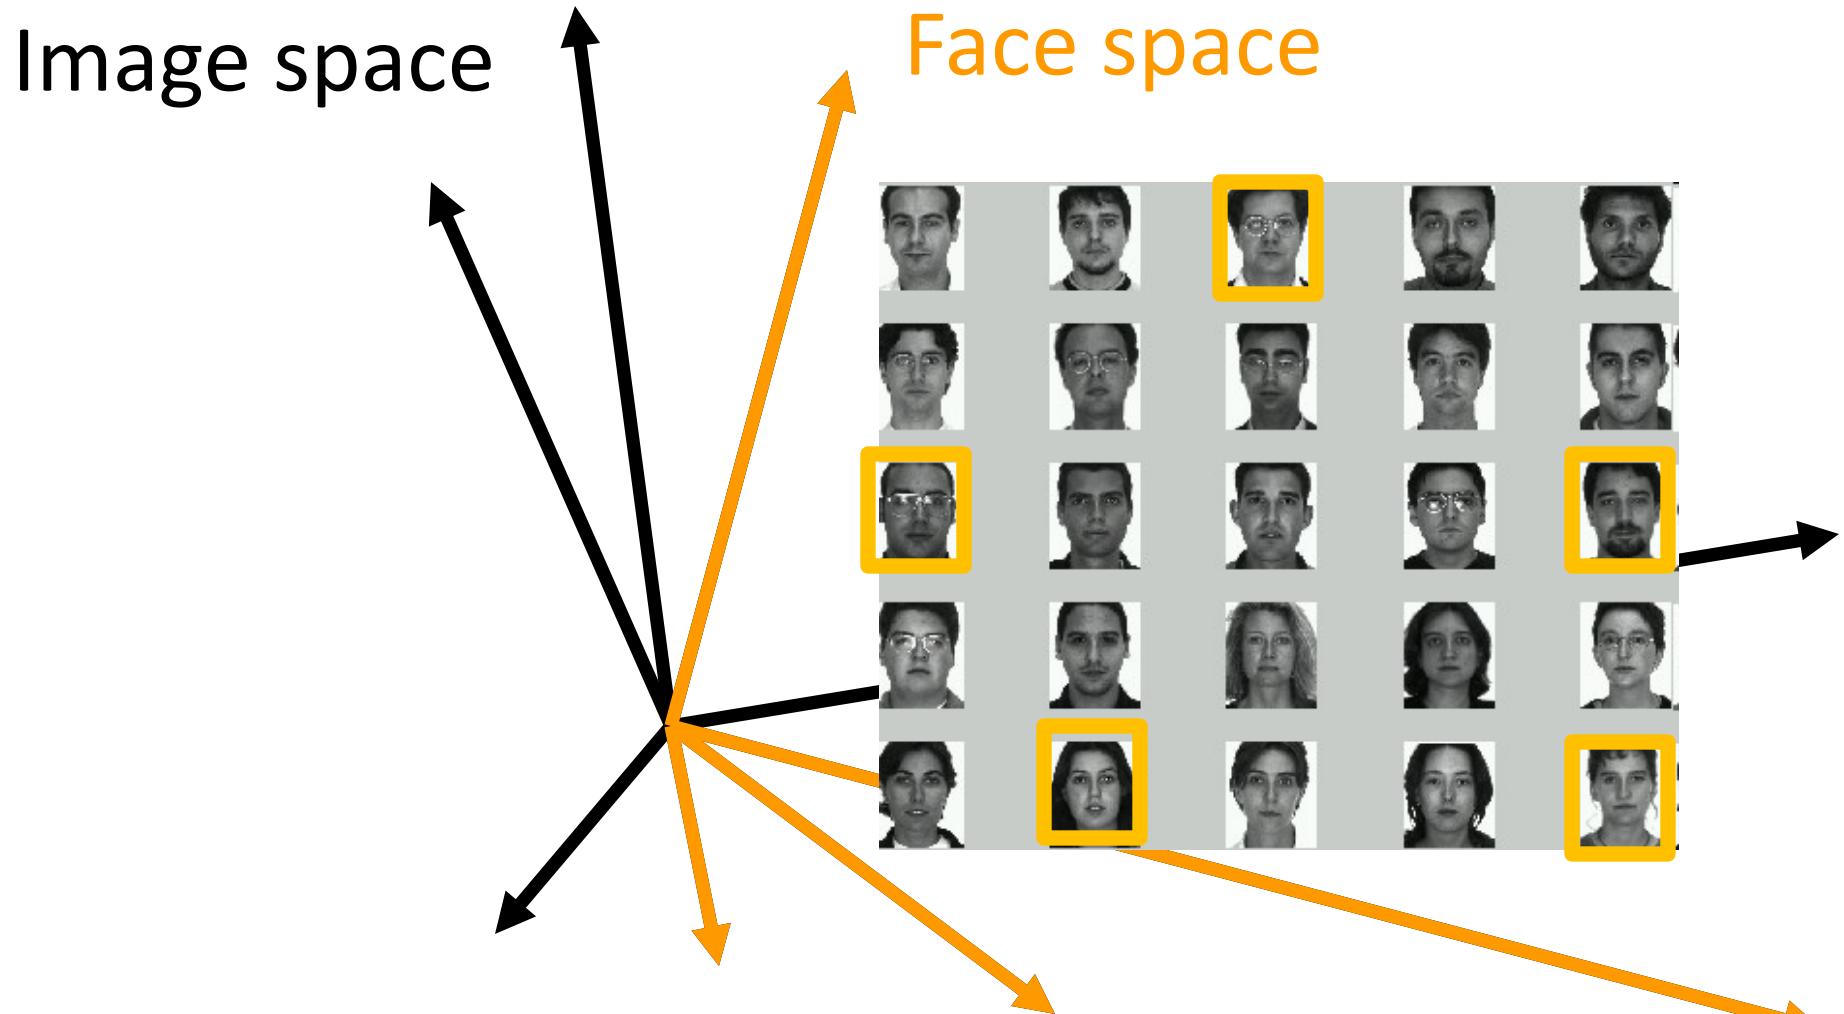
- An image is a point in a high dimensional space
  - If represented in grayscale intensity, an  $N \times M$  image is a point in  $R^{NM}$
  - E.g.  $100 \times 100$  image = 10,000 dim
- However, relatively few high dimensional vectors correspond to valid face images
- We want to effectively model the subspace of face images



Where have we seen something like this before?



- A face image
- A (non-face) image



- Compute n-dim subspace such that the projection of the data points onto the subspace has **the largest variance** among all n-dim subspaces.
- **Maximize the scatter** of the training images in face space



# Key Idea

- So, compress them to a low-dimensional subspace that captures key appearance characteristics of the visual DOFs.
- USE PCA for estimating the sub-space (dimensionality reduction)
- Compare two faces by projecting the images into the subspace and measuring the EUCLIDEAN distance between them.



# What we will learn today

- Introduction to face recognition
- The Eigenfaces Algorithm
- Linear Discriminant Analysis (LDA)

Turk and Pentland, Eigenfaces for Recognition, *Journal of Cognitive Neuroscience* **3** (1): 71–86.

P. Belhumeur, J. Hespanha, and D. Kriegman. "Eigenfaces vs. Fisherfaces: Recognition Using Class Specific Linear Projection". *IEEE Transactions on pattern analysis and machine intelligence* **19** (7): 711. 1997.



# Eigenfaces: key idea

- Assume that most face images lie on a low-dimensional subspace determined by the first  $k$  ( $k \ll d$ ) directions of maximum variance
- Use PCA to determine the vectors or “eigenfaces” that span that subspace
- Represent all face images in the dataset as linear combinations of eigenfaces

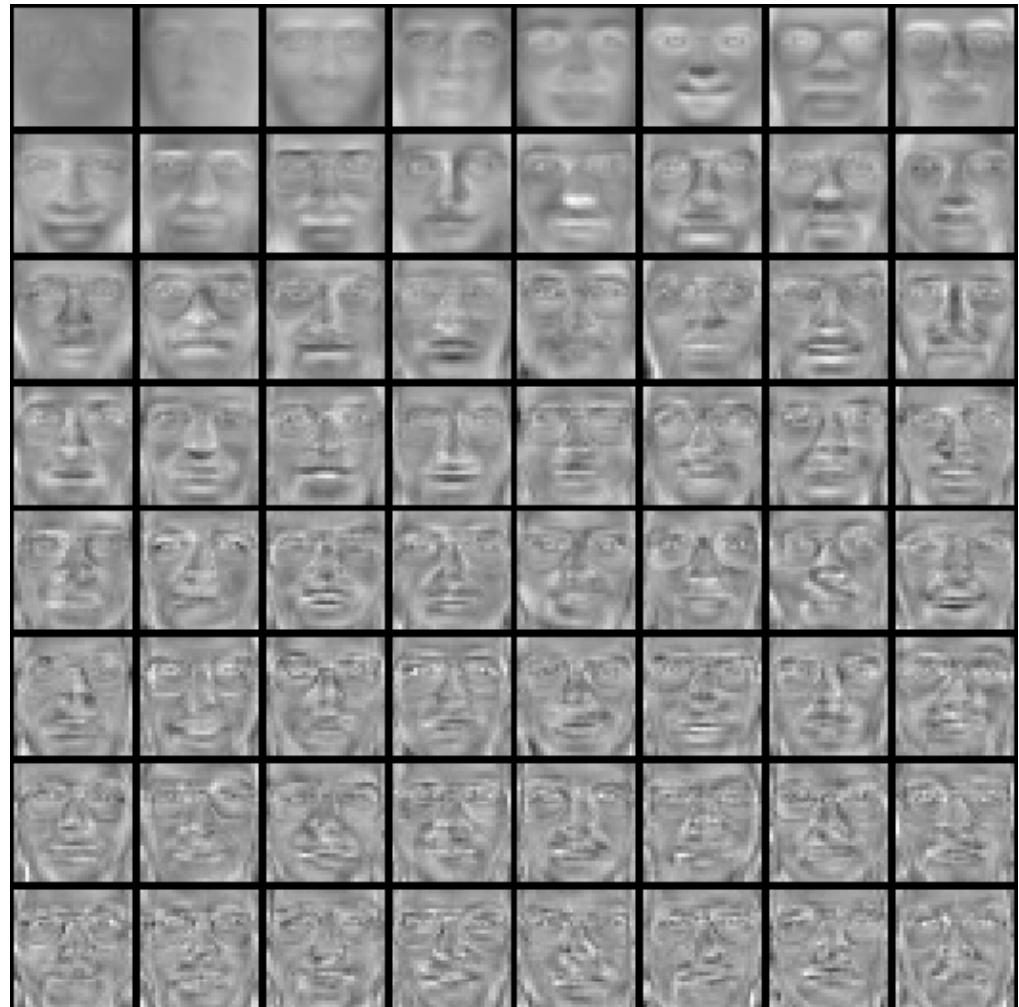
M. Turk and A. Pentland, [Face Recognition using Eigenfaces](#), CVPR 1991

Training images:  $\mathbf{x}_1, \dots, \mathbf{x}_N$



Top eigenvectors:  $\phi_1, \dots, \phi_k$

Mean:  $\mu$





# Visualization of eigenfaces

Principal component (eigenvector)  $\phi_k$



$\mu + 3\sigma_k \phi_k$



$\mu - 3\sigma_k \phi_k$





# Eigenface algorithm

- Training

1. Align training images  $x_1, x_2, \dots, x_N$



Note that each image is formulated into a long vector!

2. Compute average face  $\mu = \frac{1}{N} \sum x_i$
3. Compute the difference image (the centered data matrix)

$$\begin{aligned} X_c &= \begin{bmatrix} | & | \\ x_1 & \dots & x_n \\ | & | \end{bmatrix} - \begin{bmatrix} | & | \\ \mu & \dots & \mu \\ | & | \end{bmatrix} \\ &= X - \mu \mathbf{1}^T = X - \frac{1}{n} X \mathbf{1} \mathbf{1}^T = X \left( I - \frac{1}{n} \mathbf{1} \mathbf{1}^T \right) \end{aligned}$$



# Eigenface algorithm

4. Compute the covariance matrix

$$\Sigma = \frac{1}{n} \begin{bmatrix} | & & | \\ x_1^c & \dots & x_n^c \\ | & & | \end{bmatrix} \begin{bmatrix} - & x_1^c & - \\ \vdots & \vdots & \vdots \\ - & x_n^c & - \end{bmatrix} = \frac{1}{n} X_c X_c^T$$

5. Compute the eigenvectors of the covariance matrix  $\Sigma$

6. Compute each training image  $x_i$  's projections as

$$x_i \rightarrow (x_i^c \cdot \phi_1, x_i^c \cdot \phi_2, \dots, x_i^c \cdot \phi_K) \equiv (a_1, a_2, \dots, a_K)$$

7. Visualize the estimated training face  $x_i$

$$x_i \approx \mu + a_1 \phi_1 + a_2 \phi_2 + \dots + a_K \phi_K$$



# Eigenface algorithm



6. Compute each training image  $x_i$  's projections as

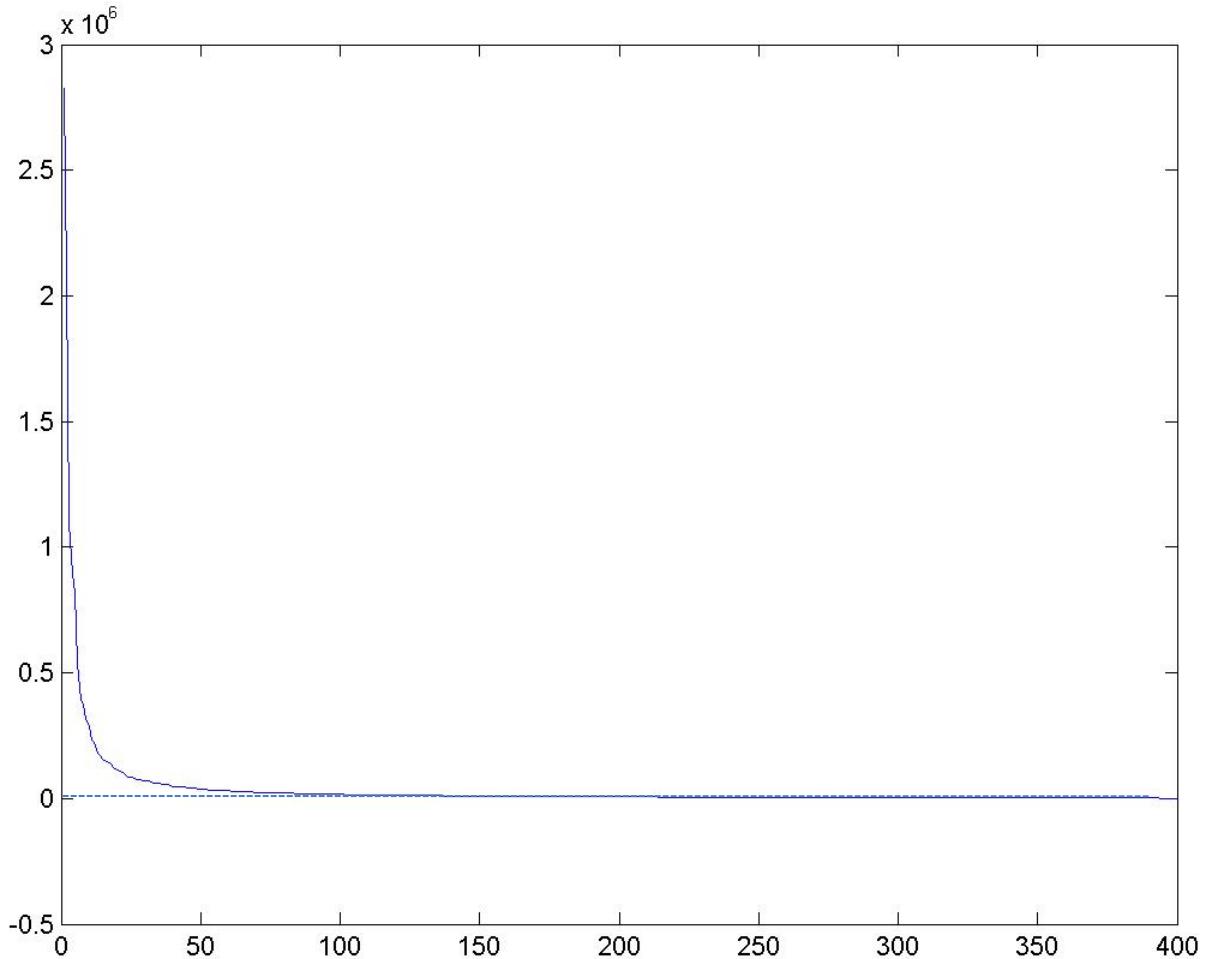
$$x_i \rightarrow (x_i^c \cdot \phi_1, x_i^c \cdot \phi_2, \dots, x_i^c \cdot \phi_K) \equiv (a_1, a_2, \dots, a_K)$$

7. Visualize the reconstructed training face  $x_i$

$$x_i \approx \mu + a_1\phi_1 + a_2\phi_2 + \dots + a_K\phi_K$$

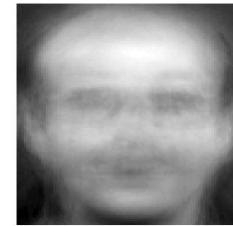


# Eigenvalues (variance along eigenvectors)



# Reconstruction and Errors

$K = 4$



$K = 200$



$K = 400$



- Only selecting the top  $K$  eigenfaces → reduces the dimensionality.
- Fewer eigenfaces result in more information loss, and hence less discrimination between faces.



# Eigenface algorithm

- Testing

1. Take query image  $t$
2. Project into eigenface space and compute projection

$$t \rightarrow ((t - \mu) \cdot \phi_1, (t - \mu) \cdot \phi_2, \dots, (t - \mu) \cdot \phi_K) \equiv (w_1, w_2, \dots, w_K)$$

3. Compare projection  $w$  with all  $N$  training projections

- Simple comparison metric: Euclidean
- Simple decision: K-Nearest Neighbor

(note: this “K” refers to the k-NN algorithm, is different from the previous K’s referring to the # of principal components)



# Shortcomings

- Requires carefully controlled data:
  - All faces centered in frame
  - Same size
  - Some sensitivity to angle
- Alternative:
  - “Learn” one set of PCA vectors for each angle
  - Use the one with lowest error
- Method is completely knowledge free
  - (sometimes this is good!)
  - Doesn’t know that faces are wrapped around 3D objects (heads)
  - Makes no effort to preserve class distinctions



# Summary for Eigenface

## Pros

- Non-iterative, globally optimal solution

## Limitations

- PCA projection is **optimal for reconstruction** from a low dimensional basis, but **may NOT be optimal for discrimination...** Is there a better dimensionality reduction?



Besides face recognitions, we can also do  
Facial expression recognition

# Happiness subspace (method A)

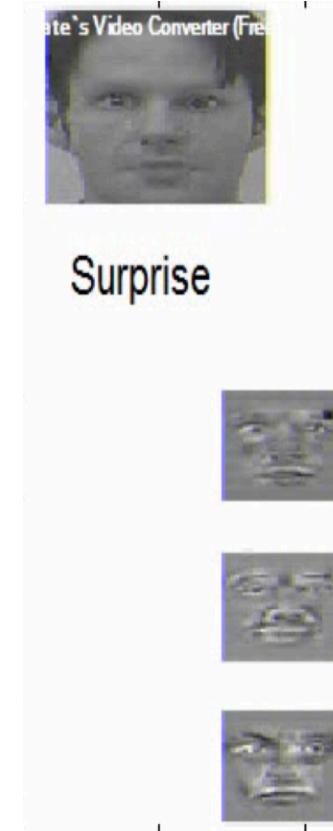
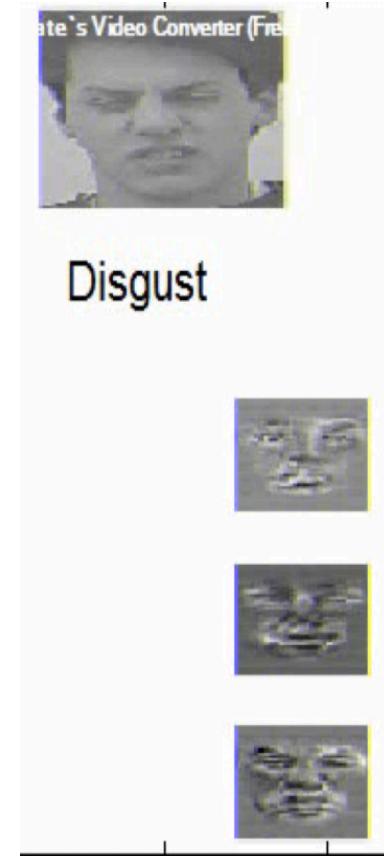
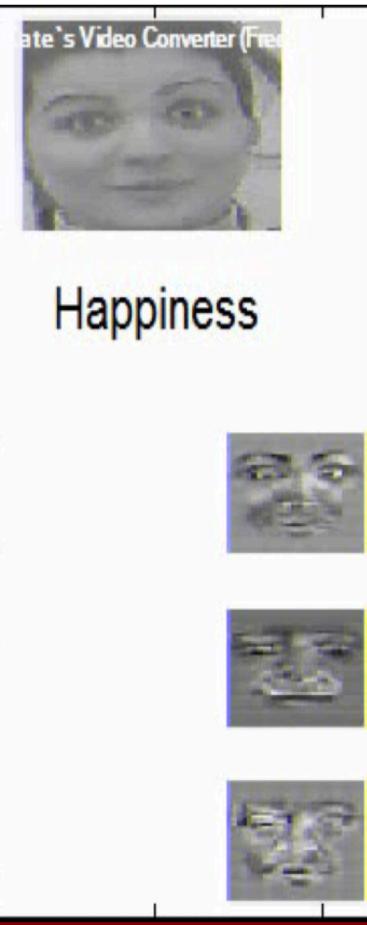




# Disgust subspace (method A)



# Facial Expression Recognition Movies (method A)





# What we will learn today

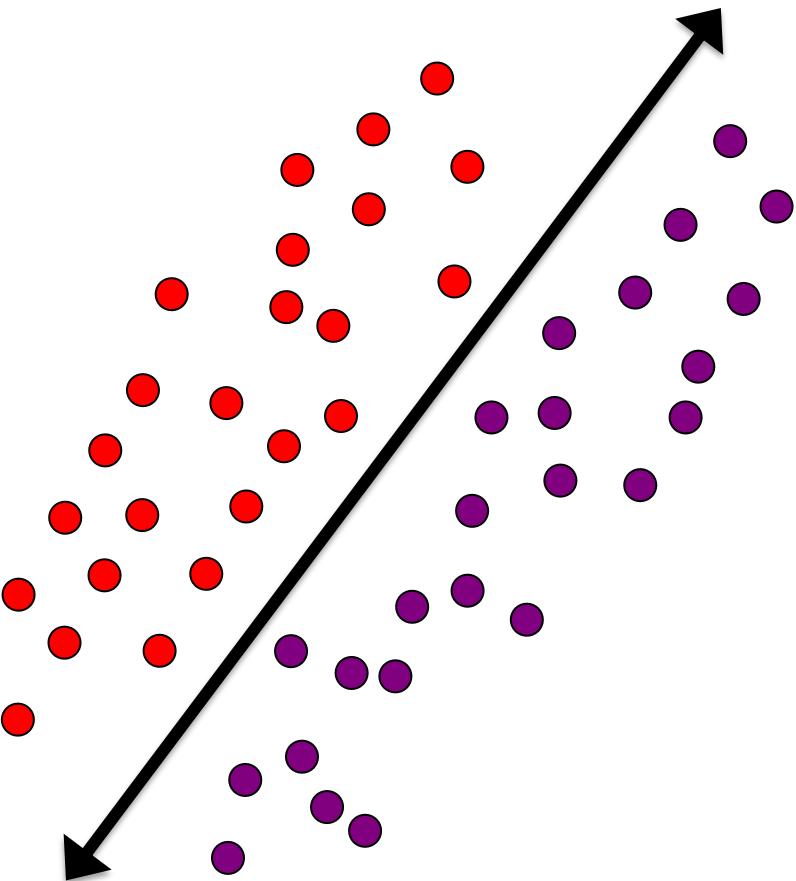
- Introduction to face recognition
- The Eigenfaces Algorithm
- Linear Discriminant Analysis (LDA)

Turk and Pentland, Eigenfaces for Recognition, *Journal of Cognitive Neuroscience* **3** (1): 71–86.

P. Belhumeur, J. Hespanha, and D. Kriegman. "Eigenfaces vs. Fisherfaces: Recognition Using Class Specific Linear Projection". *IEEE Transactions on pattern analysis and machine intelligence* **19** (7): 711. 1997.

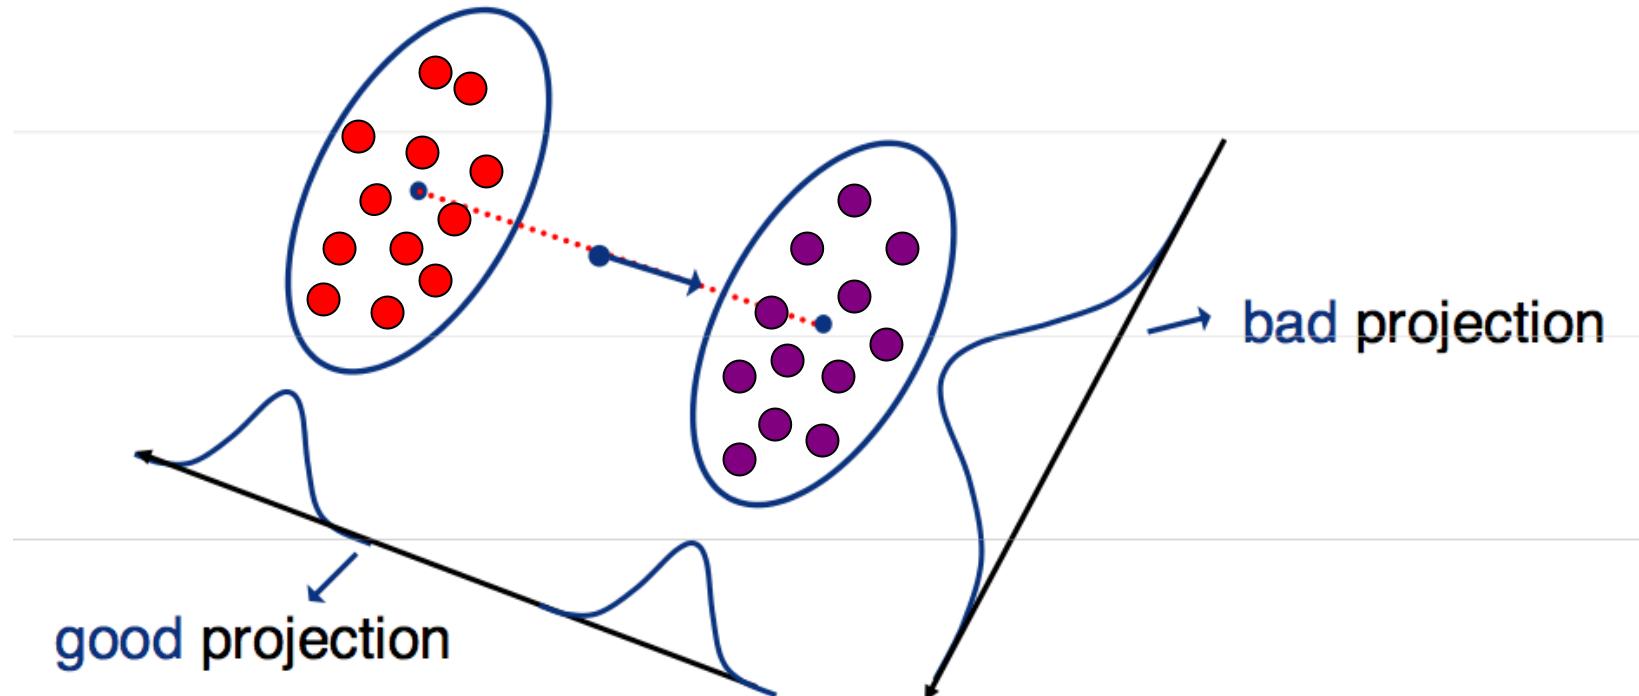


Which direction will is the first principle component?



# Fischer's Linear Discriminant Analysis

- Goal: find the best separation between two classes



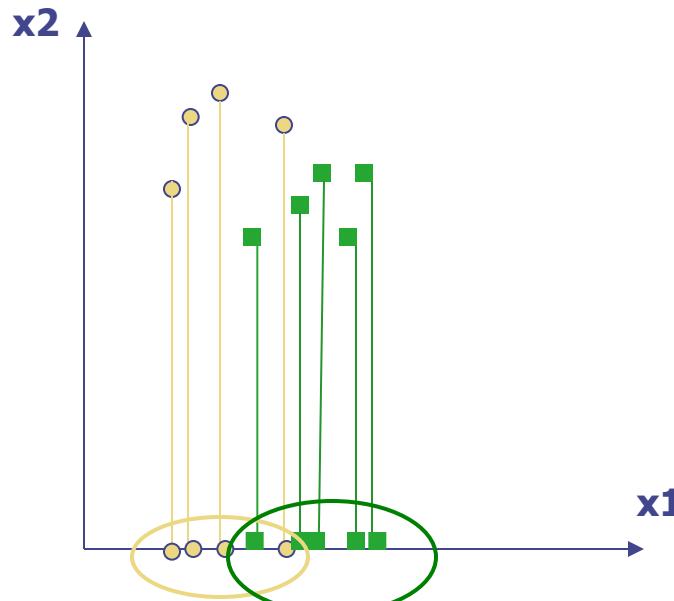


# Difference between PCA and LDA

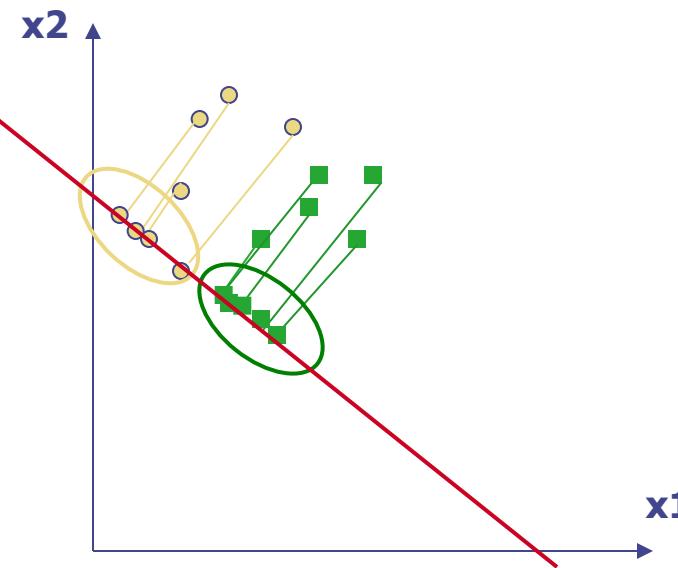
- PCA preserves maximum variance
- LDA preserves discrimination
  - Find projection that maximizes scatter between classes and minimizes scatter within classes

# Illustration of the Projection

- Using two classes as example:

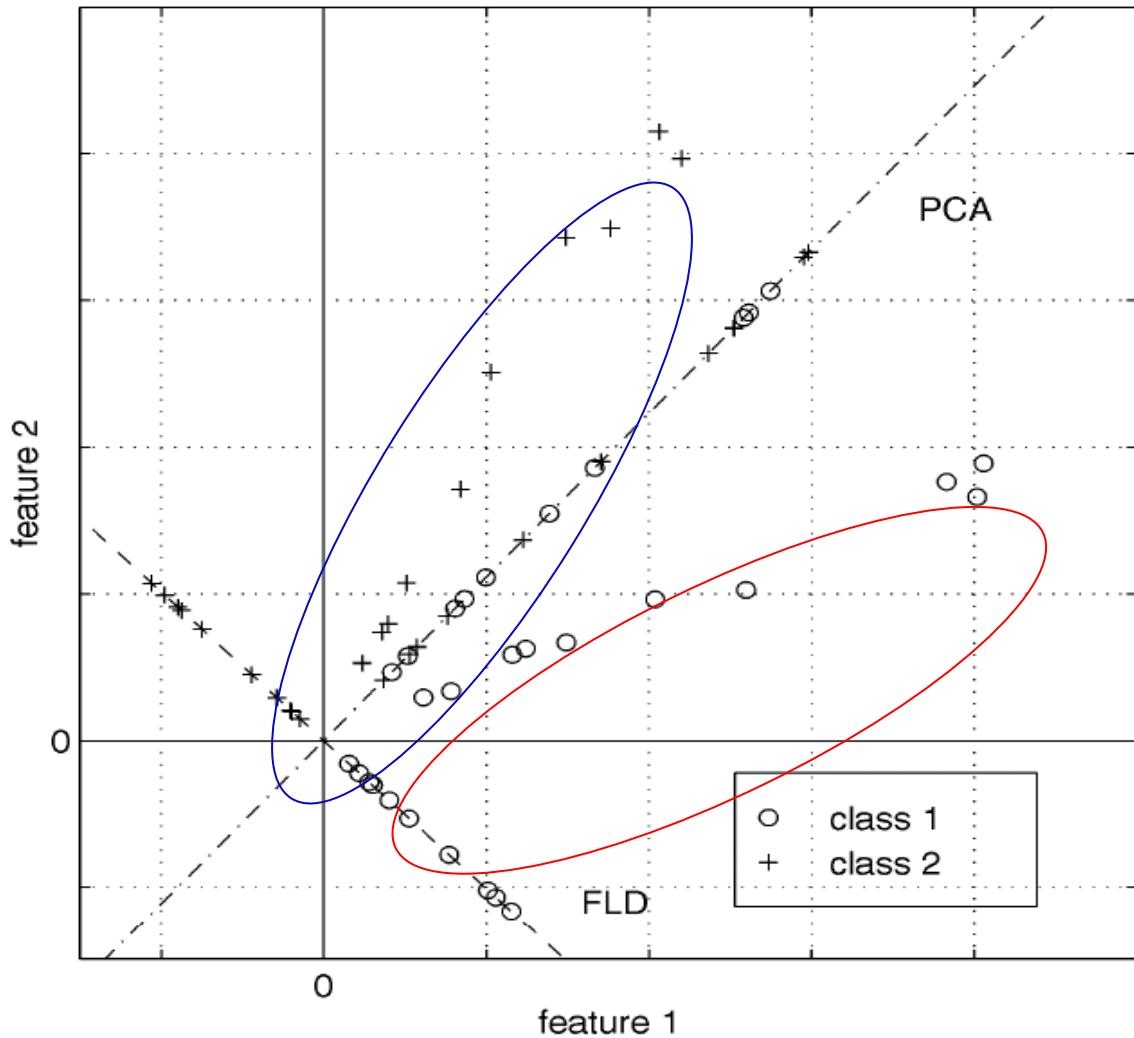


Poor Projection



Good

# Basic intuition: PCA vs. LDA





# LDA with 2 variables

- We want to learn a projection  $W$  such that the projection converts all the points from  $x$  to a new space (For this example, assume  $m == 1$ ):

$$z = w^T x \quad z \in \mathbf{R}^m \quad x \in \mathbf{R}^n$$

- Let the **per class** means be:

$$E_{X|Y}[X | Y = i] = \mu_i$$

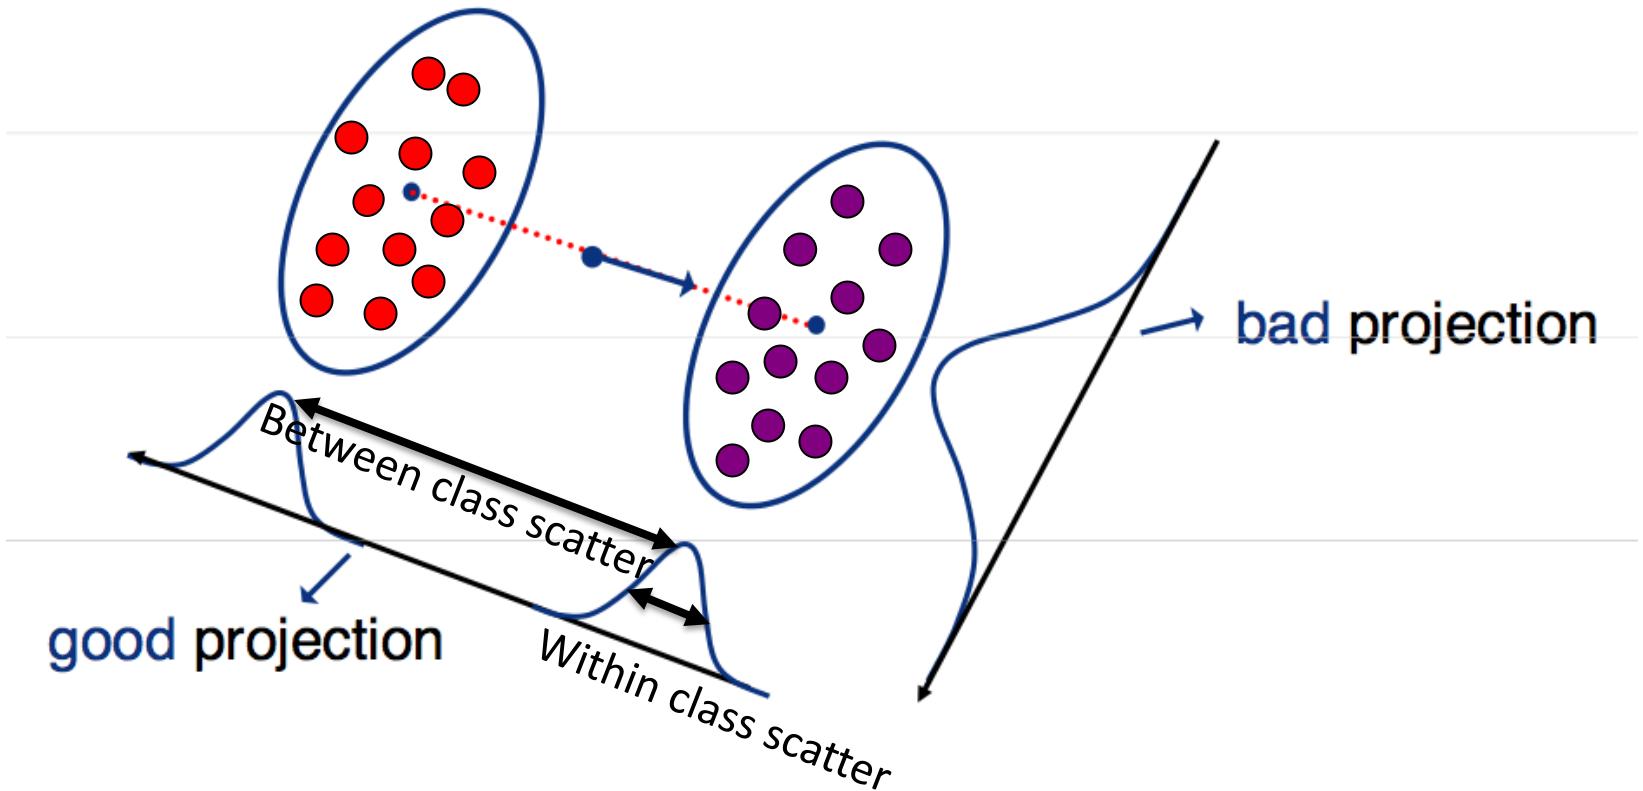
- And the **per class** covariance matrices be:

$$[(X - \mu_i)(X - \mu_i)^T | Y = i] = \Sigma_i$$

- We want a projection that maximizes:

$$J(w) = \max \frac{\text{between class scatter}}{\text{within class scatter}}$$

# Fischer's Linear Discriminant Analysis



Slide inspired by N. Vasconcelos



# LDA with 2 variables

The following objective function:

$$J(w) = \max \frac{\text{between class scatter}}{\text{within class scatter}}$$

Can be written as

$$J(w) = \frac{(E_{Z|Y}[Z|Y=1] - E_{Z|Y}[Z|Y=0])^2}{\text{var}[Z|Y=1] + \text{var}[Z|Y=0]}$$



# LDA with 2 variables

- We can write the between class scatter as:

$$\begin{aligned} \left( E_{Z|Y}[Z | Y = 1] - E_{Z|Y}[Z | Y = 0] \right)^2 &= \left( w^T [\mu_1 - \mu_0] \right)^2 \\ &= w^T [\mu_1 - \mu_0] [\mu_1 - \mu_0]^T w \end{aligned}$$

- Also, the within class scatter becomes:

$$\begin{aligned} \text{var}[Z | Y = i] &= E_{Z|Y} \left\{ (z - E_{Z|Y}[Z | Y = i])^2 | Y = i \right\} \\ &= E_{Z|Y} \left\{ (w^T [x - \mu_i])^2 | Y = i \right\} \\ &= E_{Z|Y} \left\{ w^T [x - \mu_i] [x - \mu_i]^T w | Y = i \right\} \\ &= w^T \Sigma_i w \end{aligned}$$



# LDA with 2 variables

- We can plug in these scatter values to our objective function:

$$J(w) = \frac{w^T S_B w}{w^T S_W w}$$

$$S_B = (\mu_1 - \mu_0)(\mu_1 - \mu_0)^T$$
$$S_W = (\Sigma_1 + \Sigma_0)$$

between class scatter

within class scatter

- And our objective becomes:

$$J(w) = \frac{(E_{Z|Y=1}[Z] - E_{Z|Y=0}[Z])^2}{\text{var}[Z|Y=1] + \text{var}[Z|Y=0]}$$

$$= \frac{w^T (\mu_1 - \mu_0)(\mu_1 - \mu_0)^T w}{w^T (\Sigma_1 + \Sigma_0) w}$$



# LDA with 2 variables

- The scatter variables

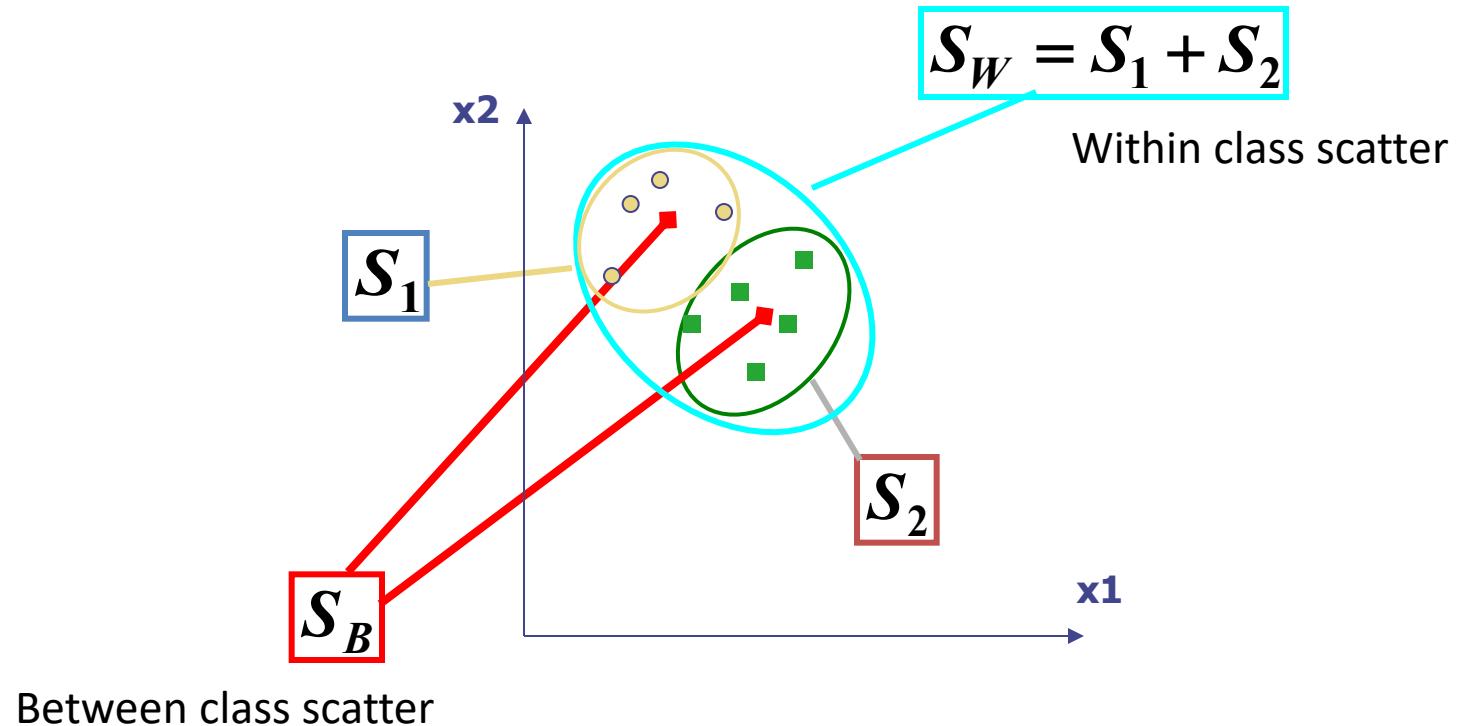
$$\boxed{S_B = (\mu_1 - \mu_0)(\mu_1 - \mu_0)^T}$$
$$S_W = (\Sigma_1 + \Sigma_0)$$

between class scatter

within class scatter



# Visualization





# Linear Discriminant Analysis (LDA)

- Maximizing the ratio

$$J(w) = \frac{w^T S_B w}{w^T S_W w}$$

- Is equivalent to maximizing the numerator while keeping the denominator constant, i.e.

$$\max_w w^T S_B w \quad \text{subject to} \quad w^T S_W w = K$$

- And can be accomplished using Lagrange multipliers, where we define the Lagrangian as

$$L = w^T S_B w - \lambda(w^T S_W w - K)$$

- And maximize with respect to both  $w$  and  $\lambda$



# Linear Discriminant Analysis (LDA)

- Setting the gradient of

$$L = w^T (S_B - \lambda S_W)w + \lambda K$$

With respect to  $w$  to zeros we get

$$\nabla_w L = 2(S_B - \lambda S_W)w = 0$$

or

$$S_B w = \lambda S_W w$$

- This is a generalized eigenvalue problem
- The solution is easy when  $S_w^{-1} = (\Sigma_1 + \Sigma_0)^{-1}$



# Linear Discriminant Analysis (LDA)

- In this case

$$S_w^{-1} S_B w = \lambda w$$

- And using the definition of  $S_B$

$$S_w^{-1} (\mu_1 - \mu_0) (\mu_1 - \mu_0)^T w = \lambda w$$

- Assuming that  $(\mu_1 - \mu_0)^T w = \alpha$  is a scalar, this can be written as

$$S_w^{-1} (\mu_1 - \mu_0) = \frac{\lambda}{\alpha} w$$

- and since we don't care about the magnitude of  $w$

$$w^* = S_w^{-1} (\mu_1 - \mu_0) = (\Sigma_1 + \Sigma_0)^{-1} (\mu_1 - \mu_0)$$



# LDA with N variables and C classes



# Variables

- N Sample images:  $\{x_1, \dots, x_N\}$
- C classes:  $\{Y_1, Y_2, \dots, Y_c\}$
- Average of each class:  $\mu_i = \frac{1}{N_i} \sum_{x_k \in Y_i} x_k$
- Average of all data:  $\mu = \frac{1}{N} \sum_{k=1}^N x_k$



# Scatter Matrices

- Scatter of class i:  
$$S_i = \sum_{x_k \in Y_i} (x_k - \mu_i)(x_k - \mu_i)^T$$
- Within class scatter:  
$$S_W = \sum_{i=1}^c S_i$$
- Between class scatter:  
$$S_B = \sum_{i=1}^c \sum_{j \neq i} (\mu_i - \mu_j)(\mu_i - \mu_j)^T$$



# Mathematical Formulation

- Recall that we want to learn a projection  $W$  such that the projection converts all the points from  $x$  to a new space  $z$ :

$$z = w^T x \quad z \in \mathbf{R}^m \quad x \in \mathbf{R}^n$$

- After projection:
  - Between class scatter  $\tilde{S}_B = W^T S_B W$
  - Within class scatter  $\tilde{S}_W = W^T S_W W$
- So, the objective becomes:

$$W_{opt} = \arg \max_w \frac{|\tilde{S}_B|}{|\tilde{S}_W|} = \arg \max_w \frac{|W^T S_B W|}{|W^T S_W W|}$$



# Mathematical Formulation

$$W_{opt} = \arg \max_w \frac{|W^T S_B W|}{|W^T S_W W|}$$

- Solve generalized eigenvector problem:

$$S_B w_i = \lambda_i S_W w_i \quad i = 1, \dots, m$$



# Mathematical Formulation

- Solution: Generalized Eigenvectors

$$S_B w_i = \lambda_i S_W w_i \quad i = 1, \dots, m$$

- Rank of  $S_{opt}$  is limited
  - $\text{Rank}(S_B) \leq |C|-1$
  - $\text{Rank}(S_W) \leq N-C$

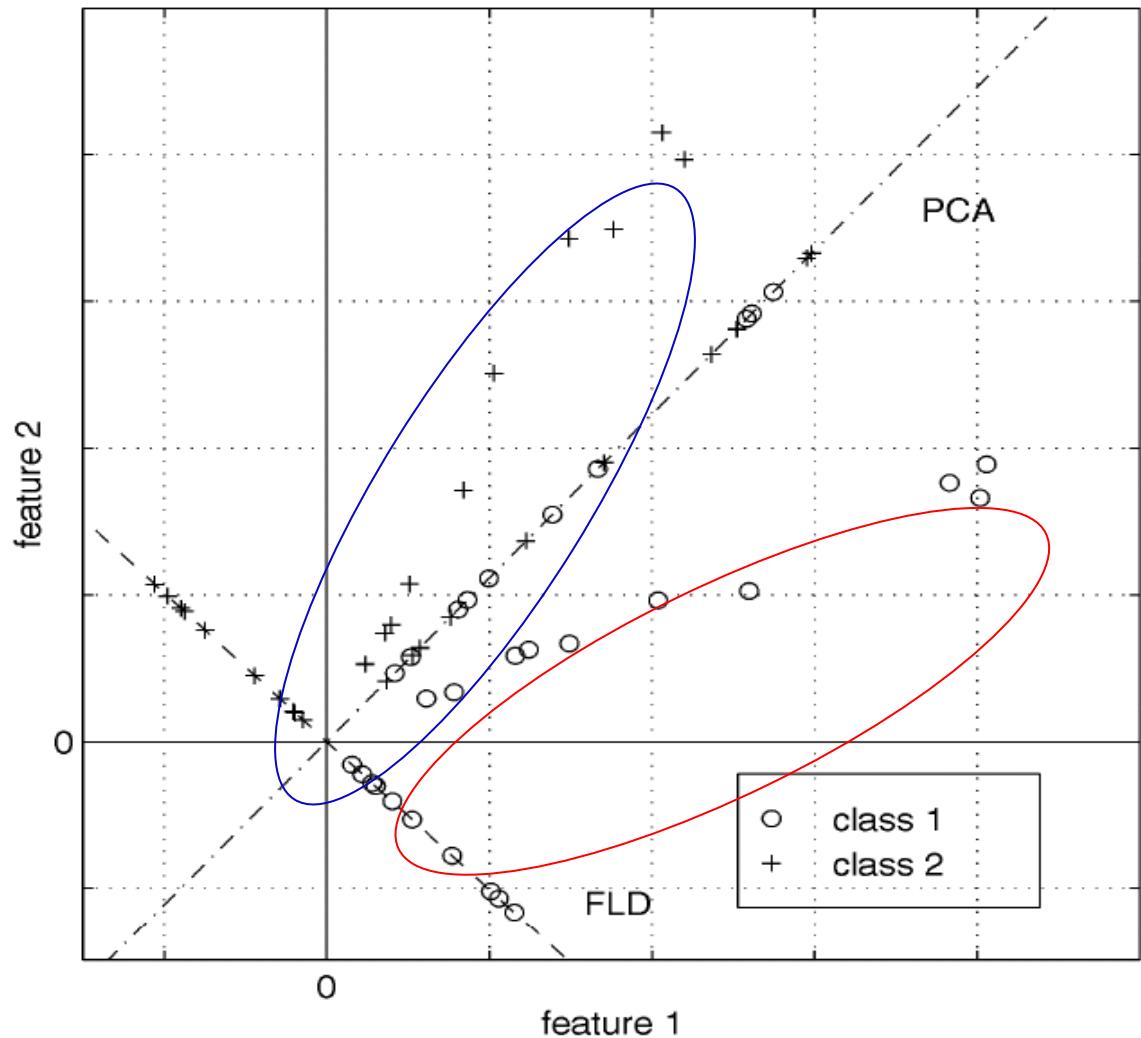


# PCA vs. LDA

- Eigenfaces exploit the max scatter of the training images in face space
- Fisherfaces attempt to maximise the **between class scatter**, while minimising the **within class scatter**.

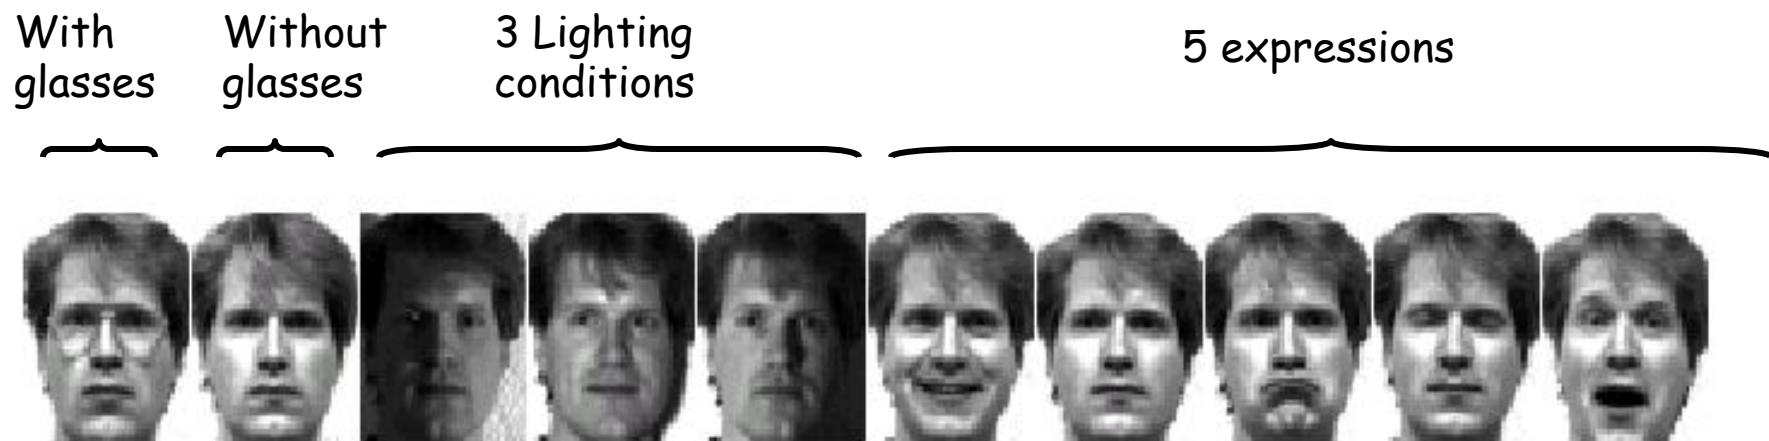


# Basic intuition: PCA vs. LDA

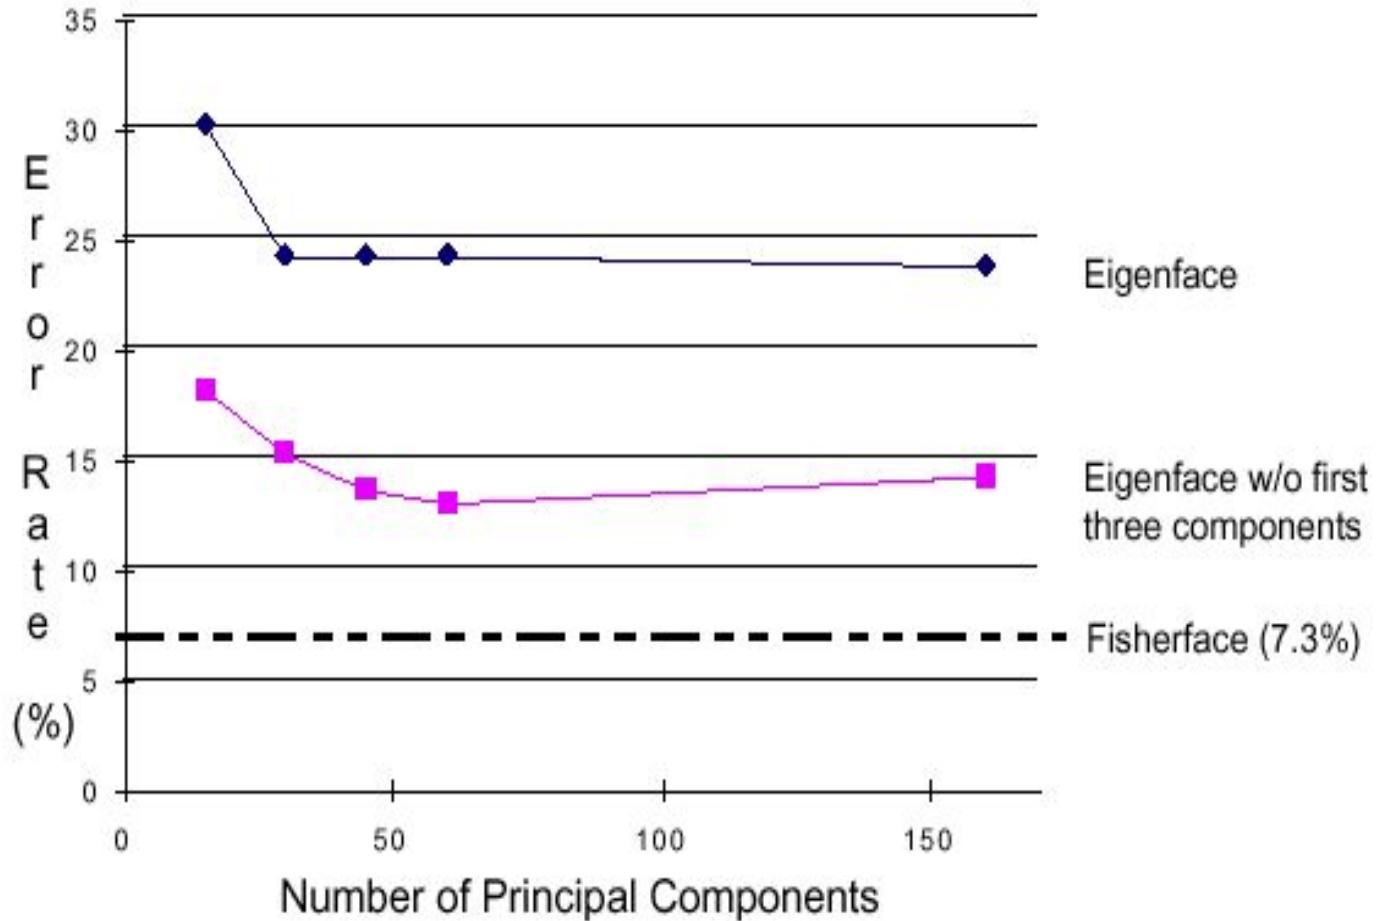


# Results: Eigenface vs. Fisherface

- Input: 160 images of 16 people
- Train: 159 images
- Test: 1 image
- Variation in Facial Expression, Eyewear, and Lighting



# Eigenface vs. Fisherface





# What we have learned today

- Introduction to face recognition
- The Eigenfaces Algorithm
- Linear Discriminant Analysis (LDA)

Turk and Pentland, Eigenfaces for Recognition, *Journal of Cognitive Neuroscience* **3** (1): 71–86.

P. Belhumeur, J. Hespanha, and D. Kriegman. "Eigenfaces vs. Fisherfaces: Recognition Using Class Specific Linear Projection". *IEEE Transactions on pattern analysis and machine intelligence* **19** (7): 711. 1997.